

# CS 561: Data Systems Architectures

class 2

## Data Systems 101

Prof. Manos Athanassoulis

<https://bu-disc.github.io/CS561/>

# some reminders



no smartphones




no laptop

# What do we do in this class?



reading papers



presentations



quizzes + exam



projects

# class summary

2 classes per week & OH/Labs 4 multiple per week

**each student**

1 paper discussion (as presenter or critic or proponent) + Quizzes + exam

project 0 (individual project) + project 1 (group project)

systems or research project (group project)

proposal + mid-semester report + final report/presentation

# Projects

# AND

## project 0

A small implementation project  
to sharpen dev skills

independent project



Due on Jan 31, 2025

## project 1

A medium project to give you a flavor of  
large-scale production system

groups of 3




# Projects

**AND**

## **project 0**

A small implementation project  
to sharpen dev skills

independent project



Due on Jan 31, 2025

## **project 1**

A medium project to give you a flavor of  
large-scale production system

groups of 3



Start forming groups NOW!  
Due on Feb 14, 2025

# Projects

## systems project

## groups of 3

# implementation-heavy C/C++ project

OR

# research project

## groups of 3

**pick a subject (list available on the website)**

# design & analysis

# experimentation



# Projects

# systems project

## groups of 3

# implementation-heavy C/C++ project

**OR**

# research project

## groups of 3

**pick a subject (list available on the website)**

# design & analysis

# experimentation




1. Proposal
  2. Mid-semester report
  3. Final report + Presentation

# class timeline



*discussions  
interaction in OH & Lab  
questions*



# Piazza



2 classes per week & OH/Labs multiple times per week

all discussions & announcements

<http://piazza.com/bu/spring2026/cs561/>

also available on class website

We have added everyone who already registered!  
Please double-check!

size (volume)

rate (velocity)

sources (variety)

veracity & value

*big data*

*(it's not only about size)*

The 3 V's

size (volume)

rate (velocity)

sources (variety)

veracity & value

*big data*

(it's not only about size)

The 3 V's

+ our ability to collect *machine-generated* data



scientific experiments




sensors


social A black icon of two user profiles, representing social data.

Internet-of-things A black icon of a network graph with nodes and connections, representing the Internet of Things.


a **data system** is a large software system  
that **stores data**, and provides the **interface** to  
**update** and **access** them **efficiently**



a **data system** is a large software system  
that **stores data**, and provides the **interface** to  
**update** and **access** them **efficiently**




a **data system** is a large software system  
that **stores data**, and provides the **interface** to  
**update** and **access** them **efficiently**




data system: breaking the blackbox

 *algorithms & operators*

*data & metadata*





# growing environment

ORACLE®

facebook.

**DB**

ACID  
large systems  
complex  
lots of tuning

**noSQL**

BASE  
simple, clean  
“just enough”



IBM

Google



®

>\$200B by 2020, growing at 11.7% every year

[The Forbes, 2016]

# growing environment

**DB**

ACID  
large systems  
complex  
lots of tuning

**noSQL**

BASE  
simple, clean  
“just enough”

**ORACLE®**

**facebook.**



**IBM**

**Google**

>\$200B by 2020, growing at 11.7% every year

[The Forbes, 2016]



**APACHE**



Cockroach Labs



**mongoDB**



**Couchbase**




**SCYLLA**

\$3B by 2020, growing at 20% every year

[Forrester, 2016]

# growing environment



ORACLE®

facebook.



IBM

Google

>\$200B by 2020, growing at 11.7% every year

[The Forbes, 2016]




\$3B by 2020, growing at 20% every year

[Forrester, 2016]

# growing environment

ORACLE®

facebook.



>\$200B by 2020, growing at 11.7% every year

[The Forbes, 2016]



newSQL



\$3B by 2020, growing at 20% every year

[Forrester, 2016]


# *growing need for tailored systems*




more data



new hardware





new applications




new performance  
goals

data systems & the hardware

# memory hierarchy




# memory hierarchy (by Jim Gray)





Jim Gray, IBM, Tandem, Microsoft, DEC

**ACM Turing Award 1998**

**ACM SIGMOD Edgar F. Codd Innovations award 1993**



# memory hierarchy (by Jim Gray)




45TB @ \$150


tape?  
sequential-only magnetic storage  
still a multi-billion industry



# Jim Gray (a great scientist and engineer)





Jim Gray, IBM, Tandem, Microsoft, DEC  
**ACM Turing Award 1998**  
**ACM SIGMOD Edgar F. Codd Innovations award 1993**




*the first collection of  
technical visionary research on  
a data-intensive scientific discovery*


# memory wall



# memory wall



*be careful when you go below the green line*



# cache/memory misses

*computations  
happen here*



CPU

on-chip cache

on-board cache

main memory

flash storage

magnetic disks

*be careful when you go below the green line*

**cache miss:** looking  
for something that  
is not in the cache

**what happens if I miss?**




**memory miss:** looking  
for something that  
is not in memory

**what happens if I miss again?**




*be very careful when you go below the green line*

# data movement



data goes through  
all necessary levels

also read  
*unnecessary* data




*Photo by Gary Dineen/NBAE via Getty Images*

need to read only X  
read the whole page



# data movement



data goes through  
all necessary levels

also read  
*unnecessary* data




Photo by Gary Dineen/NBAE via Getty Images

need to read only X  
read the whole page



## remember!

disk is millions (mem, hundreds) of times slower than CPU

# page-based access & random access

**query x<7**



size=120 bytes

**memory (memory level N)**

**disk (memory level N+1)**

1, 5, 12, 24, 23

2, 7, 13, 9, 8

10, 11, 6, 14, 15

page size =  $5 * 8 = 40$  bytes

\$ 40 bytes

# page-based access & random access



**memory (memory level N)**

**disk (memory level N+1)**

1, 5, 12, 24, 23


2, 7, 13, 9, 8

10, 11, 6, 14, 15

page size =  $5 * 8 = 40$  bytes

\$ 40 bytes

# page-based access & random access



size=120 bytes

**memory (memory level N)**


**disk (memory level N+1)**



page size =  $5 * 8 = 40$  bytes

\$ 40 bytes

# page-based access & random access



size=120 bytes

**memory (memory level N)**


**disk (memory level N+1)**



page size =  $5 * 8 = 40$  bytes

\$ 80 bytes

# page-based access & random access



size=120 bytes

**memory (memory level N)**


**disk (memory level N+1)**



page size =  $5 * 8 = 40$  bytes

\$ 80 bytes

# page-based access & random access



**memory (memory level N)**


**disk (memory level N+1)**



page size =  $5 * 8 = 40$  bytes

\$ 80 bytes

# page-based access & random access



size=120 bytes


**memory (memory level N)**

**disk (memory level N+1)**



\$ 120 bytes

# page-based access & random access



size=120 bytes

**memory (memory level N)**

**disk (memory level N+1)**



page size =  $5 * 8 = 40$  bytes

what if we had an oracle (perfect index)?



# page-based access & random access

**query x<7**



size=120 bytes

**memory (memory level N)**

**disk (memory level N+1)**

1, 5, 12, 24, 23


2, 7, 13, 9, 8

10, 11, 6, 14, 15

page size =  $5 * 8 = 40$  bytes


\$ 40 bytes

# page-based access & random access



\$ 40 bytes

# page-based access & random access



size=120 bytes

**memory (memory level N)**


**disk (memory level N+1)**



page size =  $5 * 8 = 40$  bytes

\$ 40 bytes

# page-based access & random access



size=120 bytes

**memory (memory level N)**


**disk (memory level N+1)**



page size =  $5 * 8 = 40$  bytes

\$ 80 bytes

# page-based access & random access



size=120 bytes

**memory (memory level N)**


**disk (memory level N+1)**



page size =  $5 * 8 = 40$  bytes

\$ 80 bytes

# page-based access & random access



size=120 bytes

**memory (memory level N)**

**disk (memory level N+1)**

1, 5, 12, 24, 23


2, 7, 13, 9, 8

10, 11, 6, 14, 15

page size =  $5 * 8 = 40$  bytes

\$ 80 bytes

# page-based access & random access



size=120 bytes

**memory (memory level N)**

**disk (memory level N+1)**



page size =  $5 * 8 = 40$  bytes

# page-based access & random access

\$ 120 bytes



query x<7

oracle

10, 11, 6, 14, 15

2, 7, 13, 9, 8

1, 5, 2, 6

*was the oracle helpful?*

output (32 bytes)

size=120 bytes

**memory (memory level N)**

**disk (memory level N+1)**

1, 5, 12, 24, 23

2, 7, 13, 9, 8

10, 11, 6, 14, 15

page size = 5 \* 8 = 40 bytes

# when is the oracle helpful?



for which query would an oracle help us?



how to decide whether to use the oracle or not?

1, 5, 12, 24, 23

2, 7, 13, 9, 8

10, 11, 6, 14, 15

**every byte counts**


overheads and tradeoffs

**how we store data**

layouts, indexes

**know the query**

access path selection



# rules of thumb

## **sequential access**

read one block; consume it completely; discard it; read next

*hardware can predict and start prefetching*

*prefetching can exploit full memory/disk bandwidth*

## **random access**

read one block; consume it partially; discard it; (may re-use)



are random accesses always bad?

the one that helps us **avoid a large number of accesses** (random or sequential)

# a “simple” database operator

## *select operator (scan)*





how to implement it?

```
result = new array[data.size];  
j=0;  
for (i=0; i<data.size; i++)  
    if (data[i]<x)  
        result[j++]=i;
```

qualifying positions



query: value < x



over an array of N slots

data

what if only 0.1% qualifies?

memory

data


result



how to implement it?

```
result = new array[data.size];  
j=0;  
for (i=0; i<data.size; i++)  
    if (data[i]<x)  
        result[j++]=i;
```

qualifying positions



query: value < x






data

what if only 0.1% qualifies?

memory

data





how to implement it?

```
result = new array[data.size];  
j=0;  
for (i=0; i<data.size; i++)  
    if (data[i]<x)  
        result[j++]=i;
```

```
result = new array[data.size];  
j=0;  
for (i=0; i<data.size; i++)  
    result[j+=(data[i]<x)]=i;
```

qualifying positions



query: value < x

over an array of N slots

data

what if 99% qualifies?




how can we know?

branches (if statements)  
are bad for the processors,  
can we avoid them?


how to bring the values?  
(remember we have the positions)

```
result = new array[data.size];  
j=0;  
for (i=0; i<data.size; i++)  
    if (data[i]<x)  
        result[j++]=i;
```

needs coordination!  
what about result writing?



what about multi-core?  
NUMA? SIMD? GPU?






what about having multiple queries?

query1: value<x1  
query2: value<x2 ...

```
result = new array[data.size];  
j=0;  
for (i=0; i<data.size; i++)  
    if (data[i]<x)  
        result[j++]=i;
```





# zonemaps


file = collection of pages

page 0

page 1

page 2

page 3



# zonemaps

file = collection of pages

page 0

3, 16, 34, 31, 21

page 1

1, 5, 12, 24, 23

page 2

2, 7, 13, 9, 8

page 3

10, 11, 6, 14, 15

# zonemaps

file = collection of pages

page 0

3, 16, 34, 31, 21

3,34

page 1

1, 5, 12, 24, 23

1,24

page 2

2, 7, 13, 9, 8

2,13

page 3

10, 11, 6, 14, 15

6,15

*light-weight  
typically retained  
in memory*

But what if the data is sorted?



# zonemaps

file = collection of pages

page 0

1, 2, 3, 5, 6

page 1

7, 8, 9, 10, 11

page 2

12, 13, 14, 15, 16

page 3

21, 23, 24, 31, 34

1,6

7,11

12,16

21,34

*light-weight  
typically retained  
in memory*

But what if the data is sorted?



# the language of efficient systems: C/C++

***why?***

fewer assumptions

low-level control over hardware

make decisions about physical data placement and consumptions

the language of efficient systems: C/C++

***why?***

fewer assumptions

we want you in the project to make low-level decisions

# CS 561: Data Systems Architectures

class 2

## Data Systems 101

**next :**

modern main-memory data systems

&

semester project