

CS 561: Data Systems Architectures

class 3

Relational Recap & Column-Stores Basics

Prof. Manos Athanassoulis

<https://bu-disc.github.io/CS561/>

what to do now?

- A) read the syllabus and the website
- B) register to Piazza + Gradescope
- C) **finish project 0 (due 2/1)**
- D) **start working on project 1 (due 2/15)**
- E) register for the presentation (week 3)
- F) start reading papers & prepare for tech. questions (week 3)
- G) go over the class project (end of next week will be available)
- H) start working on the proposal (week 3)

How can I prepare?

1) Read background research material

- **Architecture of a Database System.**

By J. Hellerstein, M. Stonebraker and J. Hamilton.

Foundations and Trends in Databases, 2007

- **The Design and Implementation of Modern Column-store Database Systems.**

By D. Abadi, P. Boncz, S. Harizopoulos, S. Idreos, S. Madden.

Foundations and Trends in Databases, 2013

- **Data Structures for Data-Intensive Applications: Tradeoffs and Design Guidelines.**

By M. Athanassoulis, S. Idreos, D. Shasha.

Foundations and Trends in Databases, 2024

2) Start going over the papers


Database Design Abstraction Levels

Logical Design


Physical Design

System Design

Data can be messy!




Data can be messy!




Data can be messy!

what kind of indexes
size of memory buffer
how many threads to use

...



Data can be messy!



Database Design Abstraction Levels

Logical Design

Physical Design

System Design



Logical design

What is our data? How to model them?

Hierarchical? Network? Object-oriented? Flat?

Logical design

What is our data? How to model them?



relational data model
key-value data model

Logical design

What is our data? How to model them?

Hierarchical? Network? Object-oriented? Flat?

Relational & Key-value

A collection of **tables**, each being a collection of **rows and columns**
[**schema**: describes the columns of each table]

Logical Schema of “University” Database

Students

sid: string, name: string, login: string, year_birth: integer, gpa: real

Courses

cid: string, cname: string, credits: integer


Enrolled

sid: string, cid: string, grade: string

attributes for *Enrolled* ?



Relational Model and SQL



Relational Model and SQL

Students

sid: string, name: string, login: string, year_birth: integer, gpa: real

how to create the table students?

create table students (sid:char(10), name:char(40), login:char(8), age:integer, ...)

Courses

cid: string, cname: string, credits: integer

how to add a new student?

insert into students (U1398217312, John Doe, john19, 19, ...)

Enrolled

sid: string, cid: string, grade: string

bring me the names of all students

select name from students where GPA > 3.5

Relational Model and SQL

student

(sid1, name1, login1, year1, gpa1)
(sid2, name2, login2, year2, gpa2)
(sid3, name3, login3, year3, gpa3)
(sid4, name4, login4, year4, gpa4)
(sid5, name5, login5, year5, gpa5)
(sid6, name6, login6, year6, gpa6)
(sid7, name7, login7, year7, gpa7)
(sid8, name8, login8, year8, gpa8)
(sid9, name9, login9, year9, gpa9)

} cardinality: 9

Relational Model and SQL

student

(sid1, name1, login1, year1, gpa1)
(sid2, name2, login2, year2, gpa2)
(sid3, name3, login3, year3, gpa3)
(sid4, name4, login4, year4, gpa4)
(sid5, name5, login5, year5, gpa5)
(sid6, name6, login6, year6, gpa6)
(sid7, name7, login7, year7, gpa7)
(sid8, name8, login8, year8, gpa8)
(sid9, name9, **login9**, year9, gpa9)

cardinality: 9



what if a student does not have a login ID yet?

Relational Model and SQL

student

```
(sid1, name1, login1, year1, gpa1)  
(sid2, name2, login2, year2, gpa2)  
(sid3, name3, login3, year3, gpa3)  
(sid4, name4, login4, year4, gpa4)  
(sid5, name5, login5, year5, gpa5)  
(sid6, name6, login6, year6, gpa6)  
(sid7, name7, login7, year7, gpa7)  
(sid8, name8, login8, year8, gpa8)  
(sid9, name9, NULL, year9, gpa9)
```

cardinality: 9




what if a student does not have a login ID yet?

Relational Model and SQL



how to show all enrollments in CS561?



Relational Model and SQL



how to show all enrollments in DSA?

Students

sid: string, name: string, login: string, year_birth: integer, gpa: real


Courses

cid: string, cname: string, credits: integer

Enrolled

sid: string, cid: string, grade: string

foreign keys



using foreign keys we can join information of all three tables

```
select student.name  
from students, courses, enrolled  
where course cname="DSA"  
and course.cid=enrolled.cid  
and student.sid=enrolled.sid
```

Database Design Abstraction Levels


Logical Design

Physical Design

System Design

Physical Design

File Organization



Physical Design

File Organization

heap files

sorted files

clustered files

more ...

Indexes

should I build an index?

on which attributes/tables?

what index structure?

B-Tree Trie

Hash Bitmap

Zonemap


Physical Design

File Organization



Indexes

should I build an index?
on which attributes/tables?
what index structure?
B-Tree Trie
Hash Bitmap
Zonemap



k-ary prefix tree

Physical Design

File Organization

heap files

sorted files

clustered files

more ...

Indexes

should I build an index?
on which attributes/tables?
what index structure?
B-Tree Trie
Hash Bitmap
Zonemap



bitmap?

rid	Column	rid	10	20	30
1	30	1	0	0	1
2	20	2	0	1	0
3	30	3	0	0	1
4	10	4	1	0	0
5	20	5	0	1	0
6	10	6	1	0	0
7	30	7	0	0	1
8	20	8	0	1	0

data bitmap


works great for columns with
few distinct values

Data systems are declarative!

ask ***what*** you want

data system

system decides ***how***
to store & access



design decisions, physical design
indexing, tuning knobs

research to automate!

adaptivity

autotuning


Database Design Abstraction Levels

Logical Design


Physical Design

System Design


*algorithms
and
operators*




select max(B) from R where A>5 and C<10



modules



select max(B) from R where A>5 and C<10



registers/CPU

on chip cache

on board cache

memory

disk



tape

	registers/CPU	my head ~0
2x	on chip cache	this room 1min
10x	on board cache	this building 10min
100x	memory	Washington, DC 5 hours
10^6 x	disk	Pluto 2 years
10^9 x	tape	Andromeda 2000 years


memory wall

cache miss: looking for something that is not in the cache

memory miss: looking for something that is not in memory



data movement & page-based access




data go through
all necessary levels

also read
unnecessary data

need to read only X
read the whole page



access granularity



file system and DBMS “pages”

understanding data placement

data storage




how to physically place data?


Student (sid: string, name: string, login: string, year_birth: integer, gpa: real)

student


(sid1, name1, login1, year1, gpa1)
(sid2, name2, login2, year2, gpa2)
(sid3, name3, login3, year3, gpa3)
(sid4, name4, login4, year4, gpa4)
(sid5, name5, login5, year5, gpa5)
(sid6, name6, login6, year6, gpa6)
(sid7, name7, login7, year7, gpa7)
(sid8, name8, login8, year8, gpa8)
(sid9, name9, login9, year9, gpa9)



slotted page



slotted page




querying over slotted pages




schema: R (A,B,C,D)

file

querying over slotted pages



querying over slotted pages



select A,B,C,D from R


**select A from R
any drawbacks?**

column store

querying over slotted pages



schema: R (A,B,C,D)




select A,B,C,D from R

select A from R

select (A+B) from R

each page contains **columns!**

querying over slotted pages



select A,B,C,D from R

select A from R


select (A+B) from R where A>10

each page contains **columns!**

querying over slotted pages



schema: R (A,B,C,D)




select A,B,C,D from R

select A from R

select (A+B) from R where A>10

each page contains **columns** or *groups of columns*!

querying over slotted pages



select A,B,C,D from R

select A from R

select (A+B) from R where A>10

each page contains **columns** or *groups of columns!*

what if I had all three queries?

what if only inserts/updates?

can there be something in between?

A B C D

A B C D

A B C D

A B C D

A

B

C

D

A, B

C

D

A B C D

A B C D


A B C D

A

B

C

D



A B C D

A B C D

A B C D


A

B

C

D

the way we physical store data dictates
what are the possible efficient access methods



A, B

C

D

A B C D


A

B

C

D

column-stores history line



query evaluation

A B C D

A B C D

A B C D

A B C D

A B C D

A B C D

select max(B) from R where A>5 and C<10

tuple reconstruction/early materialization



A B C D

one row at a time

A

B

C

D

select max(B) from R where A>5 and C<10

tuple reconstruction/early materialization

A B C D

one row at a time




A

late materialization


column at a time

A B C D




select max(B) from R where A>5 and C<10

```
int* input=A;  
int* output; /*needs allocation*/  
for (i=0; i<num_tuples; i++,input++)  
    if (*input>5)  
    {  
        *output=i;  
        output++;  
    }
```



select max(B) from R where A>5 and C<10



what is the benefit?

read only useful data

easy to code: working over fixed width and dense columns

scan

```
for (i=0,j=0; i<size; i++)  
    if (column[i] qualifies)  
        res[j++]=i;
```

no complex checks
no function calls
no aux metadata
easy to prefetch
as few ifs as possible


fetch

```
for (i=0,j=0; i<fetch_size; i++)  
    intermediate_result[j++]=column[ids[i]];
```

select max(B) from R where A>5 and C<10



alternative query plans



start from C (why?)

scan A & C in parallel and merge

select max(B) from R where A>5 and C<10




whole column?

row at a time

column at a time

block/vector at a time



select max(B) from R where A>5 and C<10




whole column?

row at a time

column at a time

block/vector at a time



why column-stores are here now?

late materialization – no need to reconstruct tuples

read only useful data

minimize data movement across the memory hierarchy


but it required a complete re-write

why not before?

legacy technology to catch up

more important: **analytical workloads** (as opposed to only OLTP)

new hardware: **larger memories & memory wall**



Project details are now online (more to come)



detailed discussion in next class

Readings for the project

The Log-Structured Merge-Tree (LSM-Tree) by Patrick E. O'Neil, Edward Cheng, Dieter Gawlick, Elizabeth J. O'Neil. Acta Inf. 33(4): 351-385, 1996

Monkey: Optimal Navigable Key-Value Store by Niv Dayan, Manos Athanassoulis, Stratos Idreos. SIGMOD Conference 2017

More readings (for some research projects)

Measures of Presortedness and Optimal Sorting Algorithms by Heikki Mannila. IEEE Trans. Computers 34(4): 318-325 (1985)

Small Materialized Aggregates: A Light Weight Index Structure for Data Warehousing by Guido Moerkotte. VLDB 1998

The adaptive radix tree: ARTful indexing for main-memory databases by Viktor Leis, Alfons Kemper, Thomas Neumann. ICDE 2013: 38-49

programming language: C/C++

it gives you **control over exactly** what is happening
it helps you **learn the impact** of design decisions

 avoid using libraries unless asked to do,
 so you can control storage and access patterns

CS 561: Data Systems Architectures

class 3

Relational Recap & Column-Stores Basics

Prof. Manos Athanassoulis

<https://bu-disc.github.io/CS561/>