

CS 561: Data Systems Architectures

class 3

Relational Recap & Column-Stores Basics

Prof. Manos Athanassoulis

https://bu-disc.github.io/CS561/

what to do now?

- A) read the syllabus and the website
- B) register to Piazza + Gradescope
- C) finish project 0 (due 1/31)
- D) start working on project 1 (due 2/14)
- E) register for the presentation (week 2-3)
- F) start reading papers & prepare for tech. questions (week 3)
- G) go over the class project (end of next week will be available)
- H) start working on the proposal (week 3)



How can I prepare?

1) Read background research material

- Architecture of a Database System.
 By J. Hellerstein, M. Stonebraker and J. Hamilton.
 Foundations and Trends in Databases, 2007
- The Design and Implementation of Modern Column-store Database Systems. By D. Abadi, P. Boncz, S. Harizopoulos, S. Idreos, S. Madden. Foundations and Trends in Databases, 2013
- Data Structures for Data-Intensive Applications: Tradeoffs and Design Guidelines. By M. Athanassoulis, S. Idreos, D. Shasha. Foundations and Trends in Databases, 2024
- 2) Start going over the papers



Database Design Abstraction Levels

Logical Design

Physical Design

System Design











what kind of indexes size of memory buffer how many threads to use

•••









Database Design Abstraction Levels

Logical Design

Physical Design

System Design



Logical design

What is our data? How to model them?


Hierarchical? Network? Object-oriented? Flat?



Logical design



What is our data? How to model them?



relational data model key-value data model

Logical design

What is our data? How to model them?

Hierarchical? Network? Object-oriented? Flat?

Relational & Key-value

A collection of tables, each being a collection of rows and columns

[schema: describes the columns of each table]



Logical Schema of "University" Database

Students

sid: string, name: string, login: string, year_birth: integer, gpa: real

Courses


cid: string, cname: string, credits: integer

Enrolled

sid: string, cid: string, grade: string









how to create the table students?

create table students (sid:char(10), name:char(40), login:char(8), age:integer, ...)

Students

sid: string, **name**: string, **login**: string, **year_birth**: integer, **gpa**: real

how to add a new student?

insert into students (U1398217312, John Doe, john19, 19, ...)

Courses

cid: string, cname: string, credits: integer

Enrolled

sid: string, cid: string, grade: string

select name from students where GPA > 3.5



student

```
(sid1, name1, login1, year1, gpa1)
(sid2, name2, login2, year2, gpa2)
(sid3, name3, login3, year3, gpa3)
(sid4, name4, login4, year4, gpa4)
(sid5, name5, login5, year5, gpa5)
(sid6, name6, login6, year6, gpa6)
(sid7, name7, login7, year7, gpa7)
(sid8, name8, login8, year8, gpa8)
(sid9, name9, login9, year9, gpa9)
```

cardinality: 9



student

```
(sid1, name1, login1, year1, gpa1)
(sid2, name2, login2, year2, gpa2)
(sid3, name3, login3, year3, gpa3)
(sid4, name4, login4, year4, gpa4)
(sid5, name5, login5, year5, gpa5)
(sid6, name6, login6, year6, gpa6)
(sid7, name7, login7, year7, gpa7)
(sid8, name8, login8, year8, gpa8)
(sid9, name9, login9, year9, gpa9)
```

cardinality: 9





student

```
(sid1, name1, login1, year1, gpa1)
(sid2, name2, login2, year2, gpa2)
(sid3, name3, login3, year3, gpa3)
(sid4, name4, login4, year4, gpa4)
(sid5, name5, login5, year5, gpa5)
(sid6, name6, login6, year6, gpa6)
(sid7, name7, login7, year7, gpa7)
(sid8, name8, login8, year8, gpa8)
(sid9, name9, NULL, year9, gpa9)
```

cardinality: 9







how to show all enrollments in CS561?

keys Students **sid**: string, **name**: string, **login**: string, **year_birth**: integer, **gpa**: real Courses cid: string, chame: string, credits: integer Enrolled **sid**: string, **cid**: string, **grade**: string





how to show all enrollments in DSA?

Students

sid: string, name: string, login: string, year_birth: integer, gpa: real

Courses

cid: string, cname: string, credits: integer

Enrolled

sid: string, cid: string, grade: string

foreign keys

using foreign keys we can join information of all three tables

select student.name
from students, courses, enrolled
where course.cname="DSA"
and course.cid=enrolled.cid
and student.sid=enrolled.sid



Database Design Abstraction Levels

Logical Design

Physical Design

System Design



Physical Design


File Organization

heap files

sorted files

clustered files

more ...





Physical Design

File Organization

heap files

sorted files

clustered files

more ...

Indexes

should I build an index?

on which attributes/tables?

what index structure?

B-Tree Trie

Hash Bitmap

Zonemap





Indexes

trie?

heap files

sorted files

clustered files

more ...


should I build an index?

on which attributes/tables?

what index structure?

B-Tree Trie Hash Bitmap

Zonemap



k-ary prefix tree







Indexes

heap files

sorted files

clustered files

more ...

should I build an index?

on which attributes/tables?

what index structure?

B-Tree Trie
Hash Bitmap
Zonemap


rid	Column			rid	10		20		30
1		30		1	0		0		1
2		20		2	0		1		0
3		30		3	0		0		1
4		10		4	1		0		0
5		20		5	0		1		0
6		10		6	1		0		0
7		30		7	0		0		1
8		20		8	0		1		0
data			-	bitmap					

works great for columns with few distinct values



Data systems are declarative!


ask what you want



system decides *how* to store & access



design decisions, physical design indexing, tuning knobs






Database Design Abstraction Levels

Logical Design

Physical Design

System Design







modules

Parser

Optimizer

Evaluation

Storage

registers/CPU

on chip cache


on board cache

memory

disk

tape







memory wall

cache miss: looking for something that is not in the cache

for something that

is not in memory

cheaper/larger

faster

CPU

on-chip cache

on-board cache


memory miss: looking

main memory

flash storage

disks

flash





data movement & page-based access

CPU

on-chip cache

on-board cache

main memory

flash storage

disks

flash

data go through all necessary levels


also read unnecessary data

need to read only X read the whole page





access granularity





file system and DBMS "pages"

understanding data placement

data storage


Student (**sid**: string, **name**: string, **login**: string, **year_birth**: integer, **gpa**: real)

student

(sid1, name1, login1, year1, gpa1) (sid2, name2, login2, year2, gpa2) (sid3, name3, login3, year3, gpa3) (sid4, name4, login4, year4, gpa4) (sid5, name5, login5, year5, gpa5) (sid6, name6, login6, year6, gpa6) (sid7, name7, login7, year7, gpa7) (sid8, name8, login8, year8, gpa8) (sid9, name9, login9, year9, gpa9)




how to physically place data?






slotted page






slotted page













schema: R (A,B,C,D)

select A,B,C,D from R

select A from R

ABCD

file

<u>ABCD</u>

ABCD

ABCD

ABCD

<u>ABCD</u>

ABCD

ABCD

rows are **contiguous** (with possible free space at the end)

each page contains **entire** rows (all their columns)



pages



schema: R (A,B,C,D)




select A,B,C,D from R

select A from R

any drawbacks?

each page contains columns!




column store





schema: R (A,B,C,D)



select A,B,C,D from R

select A from R

select (A+B) from R

each page contains columns!






schema: R (A,B,C,D)

select A,B,C,D from R

select A from R

select (A+B) from R where A>10


each page contains columns!







schema: R (A,B,C,D)



select A,B,C,D from R

select A from R

select (A+B) from R where A>10

each page contains columns or groups of columns!





schema: R (A,B,C,D)

select A,B,C,D from R

select A from R


select (A+B) from R where A>10

each page contains columns or groups of columns!


what if I had all three queries?


what if only inserts/updates?

can there be something in between?











the way we physical store data dictates what are the possible efficient access methods




column-stores history line





query evaluation




tuple reconstruction/early materialization



ABCD

one row at a time





tuple reconstruction/early materialization

ABCD


one row at a time




late materialization

column at a time











easy to code: working over fixed width and dense columns

scan


for (i=0,j=0; i<size; i++)
if (column[i] qualifies)
res[j++]=i;</pre>

no complex checks
no function calls
no aux metadata
easy to prefetch
as few ifs as possible

fetch

```
for (i=0,j=0; i<fetch_size; i++)
  intermediate_result[j++]=column[ids[i]];</pre>
```








alternative query plans

start from C (why?)





whole column?


row at a time

column at a time

block/vector at a time







whole column?

row at a time

column at a time

block/vector at a time



why column-stores are here now?

late materialization – no need to reconstruct tuples read only useful data minimize data movement across the memory hierarchy but it required a complete re-write

why not before?

legacy technology to catch up

more important: analytical workloads (as opposed to only OLTP)

new hardware: larger memories & memory wall





Project details are now online (more to come)



detailed discussion in next class



Readings for the project

The Log-Structured Merge-Tree (LSM-Tree) by Patrick E. O'Neil, Edward Cheng, Dieter Gawlick, Elizabeth J. O'Neil. Acta Inf. 33(4): 351-385, 1996

Monkey: Optimal Navigable Key-Value Store by Niv Dayan, Manos Athanassoulis, Stratos Idreos. SIGMOD Conference 2017

More readings (for some research projects)

Measures of Presortedness and Optimal Sorting Algorithms by Heikki Mannila. IEEE Trans. Computers 34(4): 318-325 (1985)

Small Materialized Aggregates: A Light Weight Index Structure for Data Warehousing by Guido Moerkotte. VLDB 1998

The adaptive radix tree: ARTful indexing for main-memory databases by Viktor Leis, Alfons Kemper, Thomas Neumann. ICDE 2013: 38-49



programming language: C/C++

it gives you control over exactly what is happening it helps you learn the impact of design decisions

avoid using libraries unless asked to do, so you can control storage and access patterns





CS 561: Data Systems Architectures

class 3

Relational Recap & Column-Stores Basics

Prof. Manos Athanassoulis

https://bu-disc.github.io/CS561/