BOSTON UNIVERSITY

# CS 561: **Data Systems Architectures**

## class 4

## Systems & Research Project

Prof. Manos Athanassoulis

https://bu-disc.github.io/CS561/

# Let’s revisit Zonemaps

- Light-weight auxiliary data structure (*“scan accelerator”*)

# Let's revisit Zonemaps

zonemaps

file = collection of pages

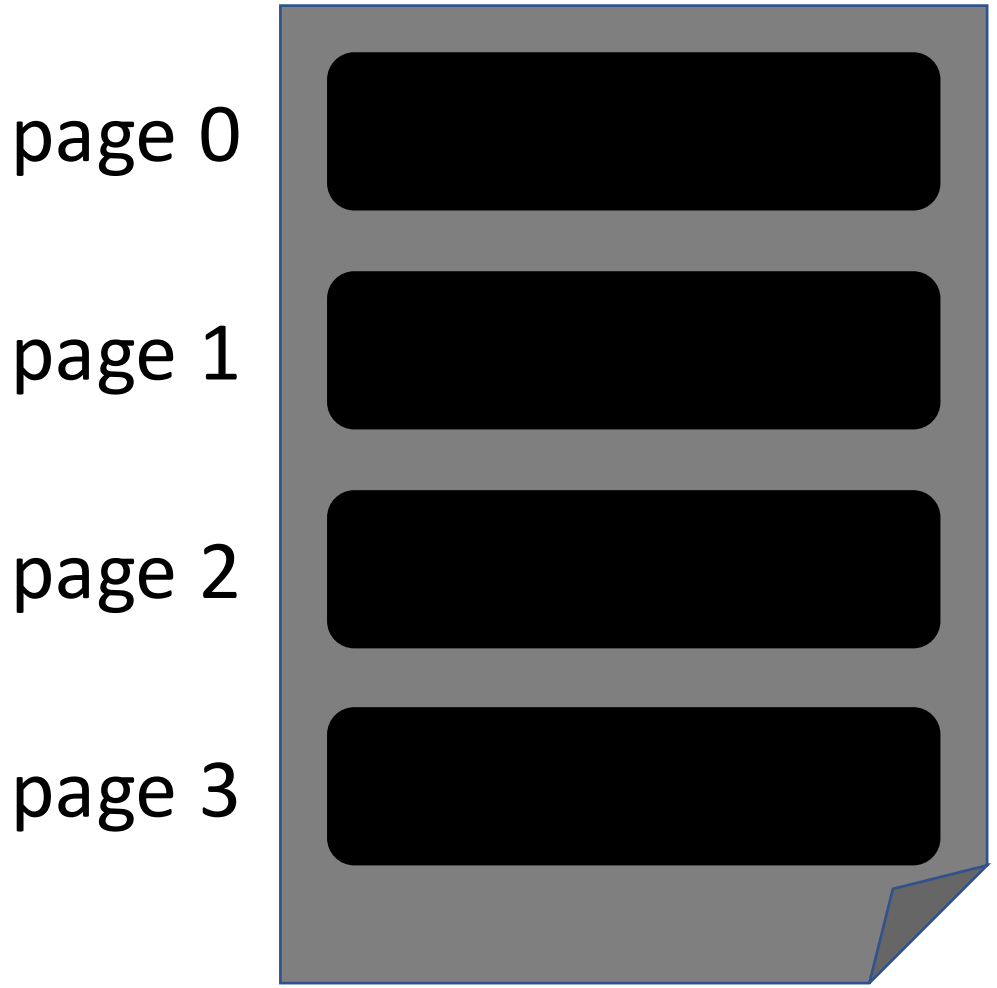

# zonemaps

file = collection of pages

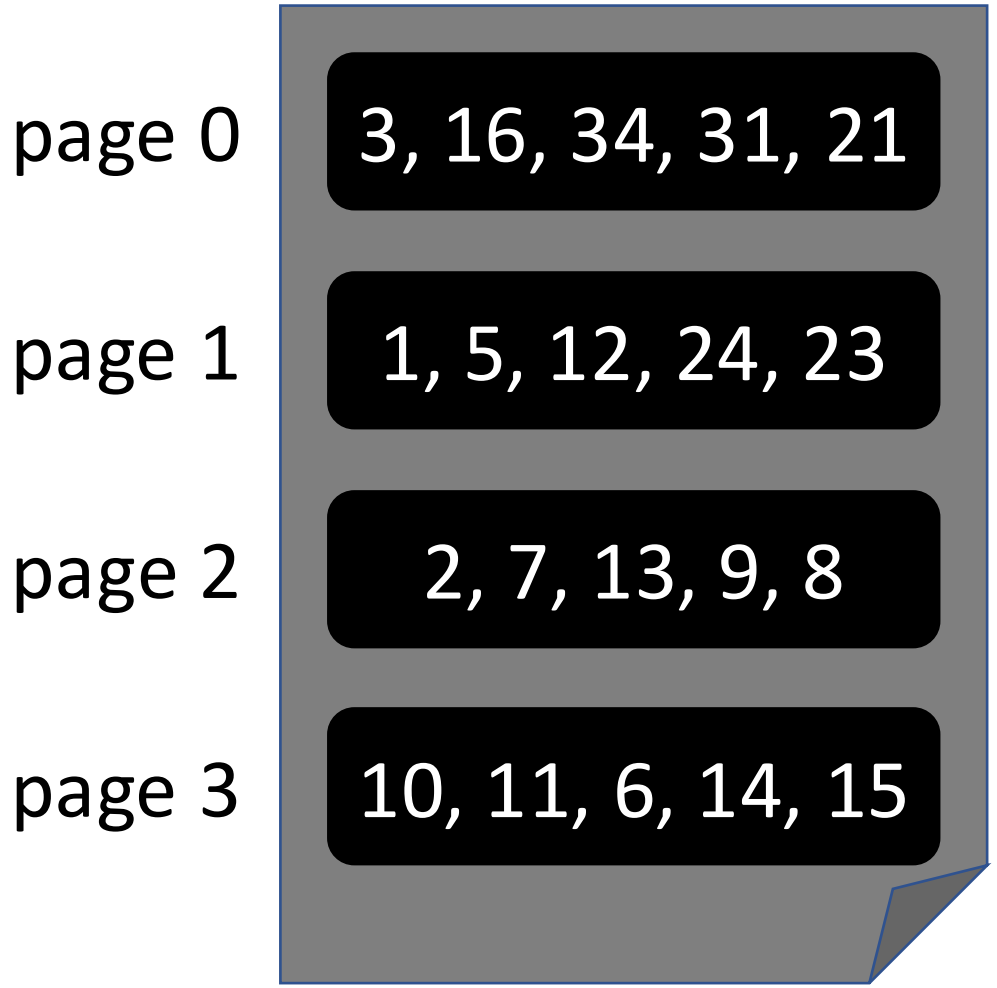

# zonemaps

file = collection of pages

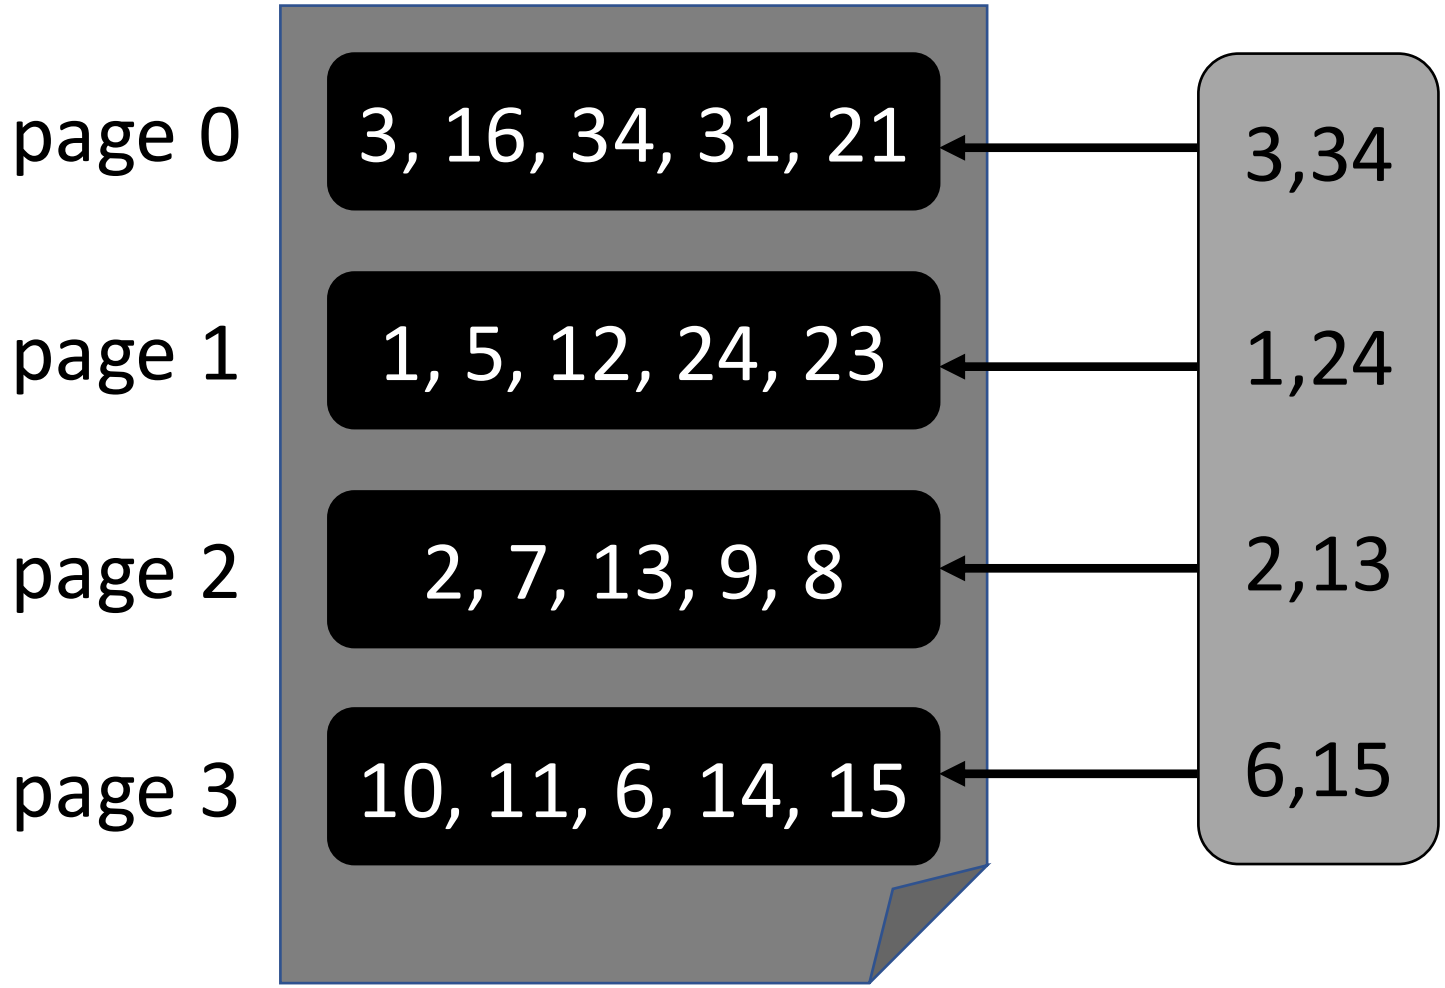


*light-weight*

*typically retained in memory*

But what if the data is sorted?

# zonemaps

file = collection of pages

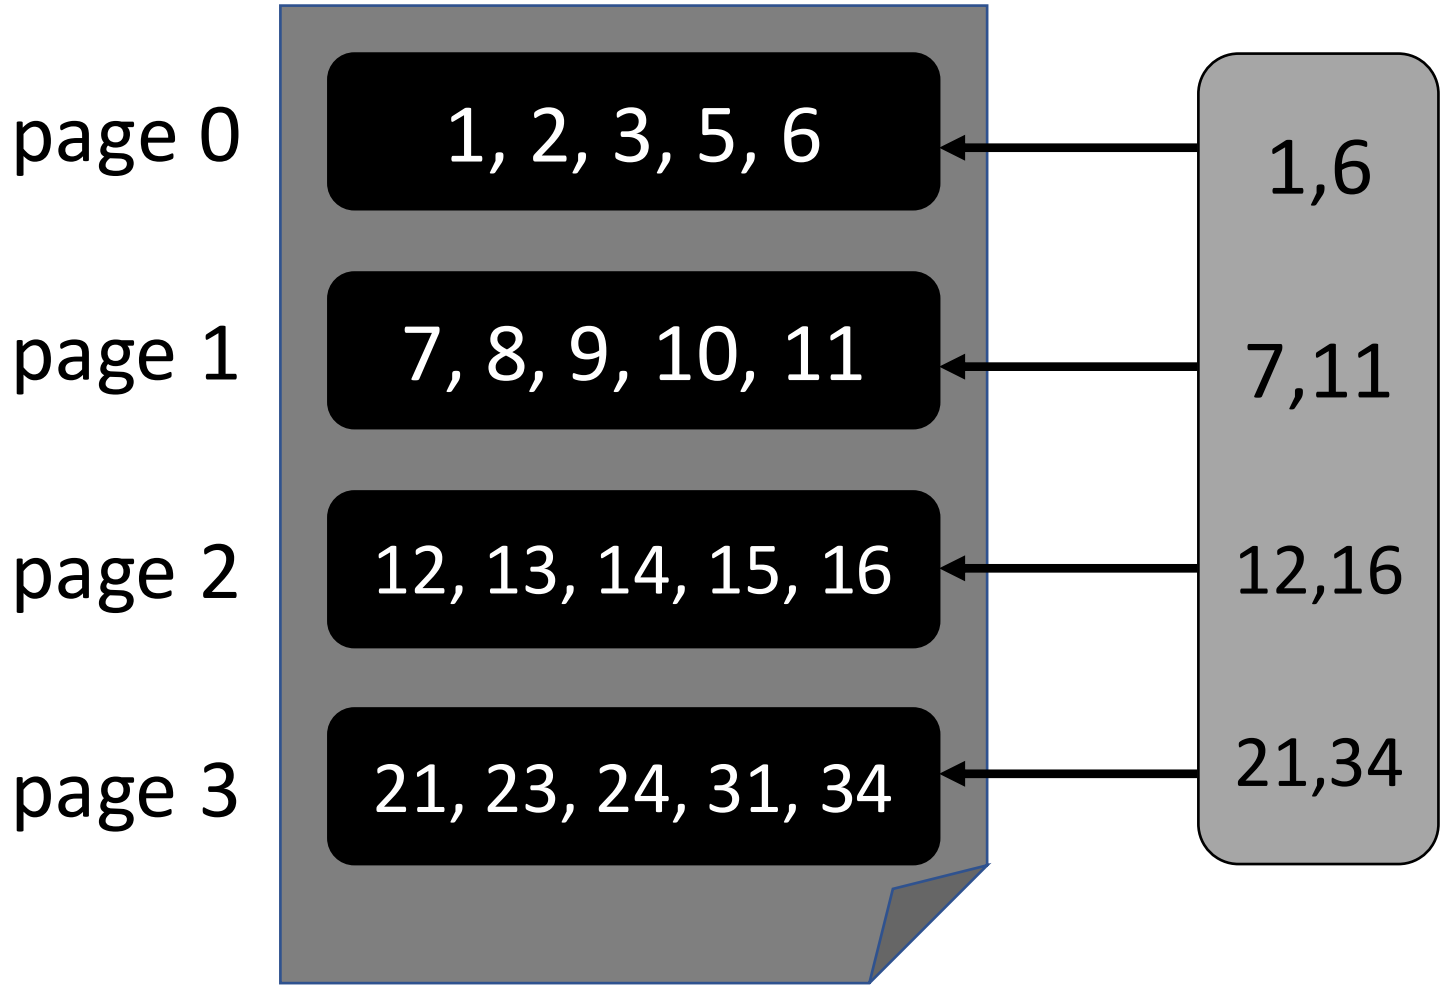


*light-weight*

*typically retained in memory*

But what if the data is sorted?

# zonemaps

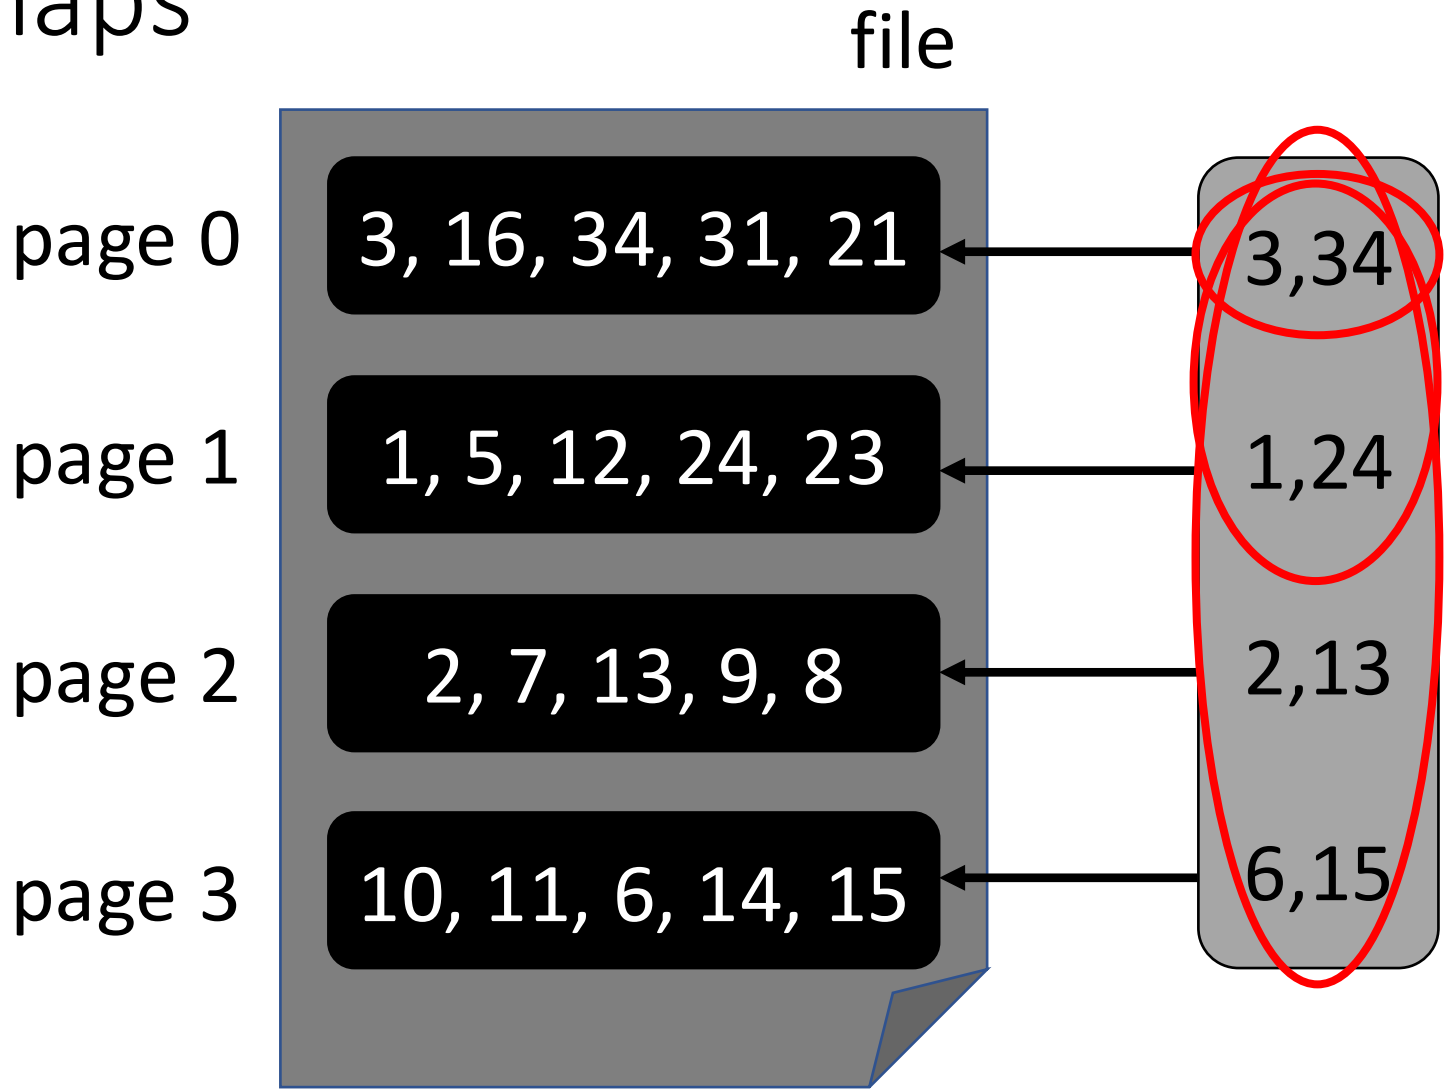


Range Queries

[25, 100] 1 page

[20, 25] 2 pages

[10, 13] 4 pages

# zonemaps

But what if the data is sorted?

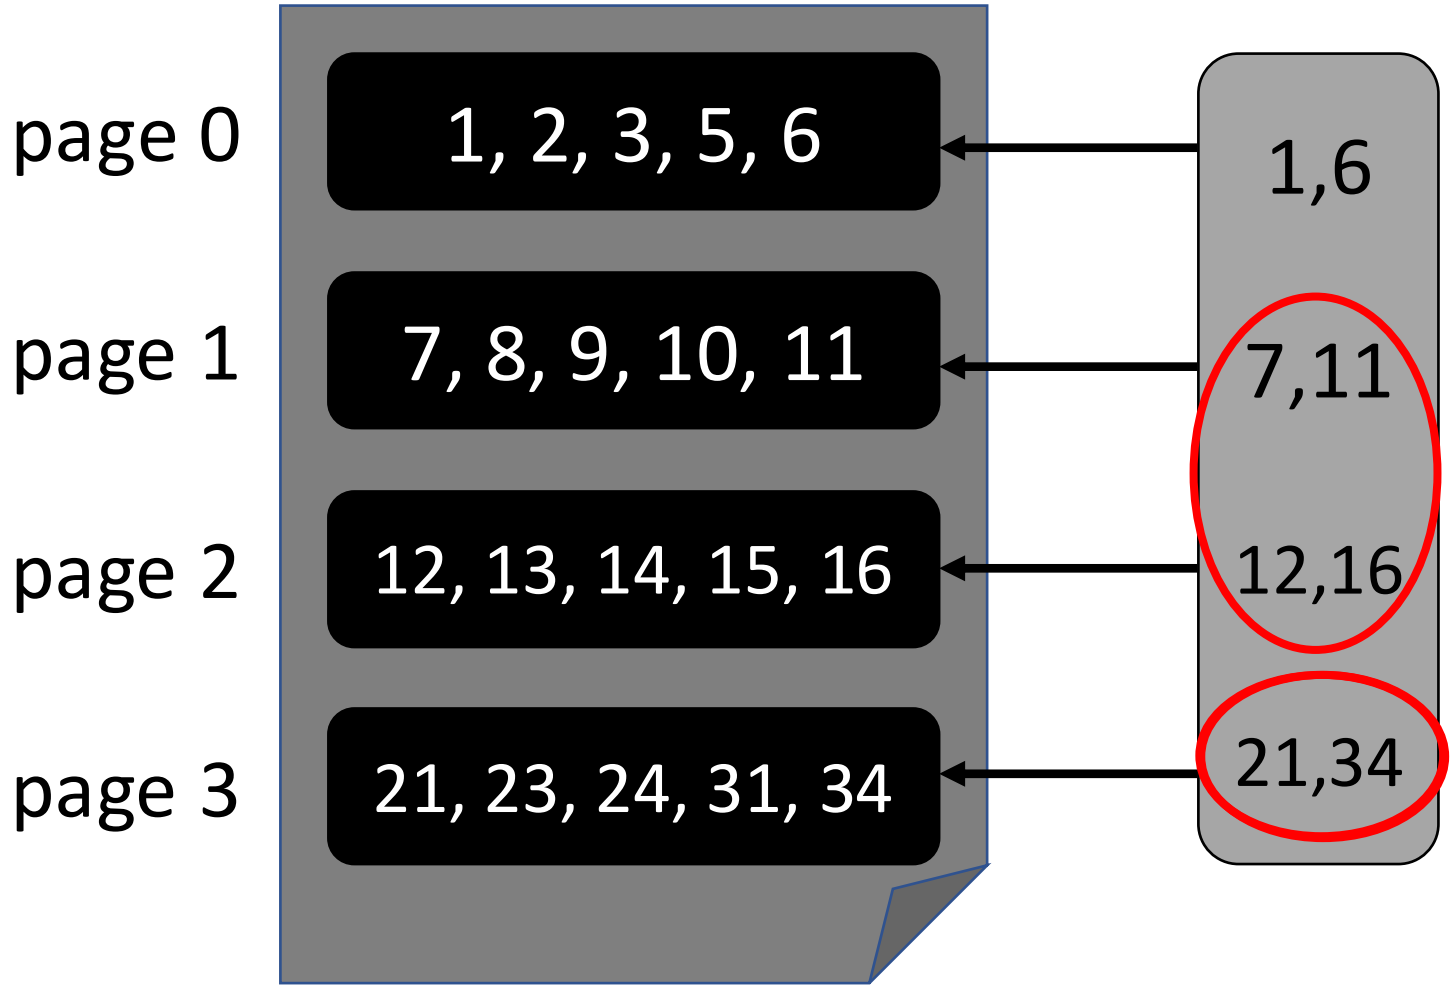


Range Queries

[25, 100] 1 page

[20, 25] 1 page

[10, 13] 2 pages

Zonemaps efficiency depends on data & queries!

# data systems










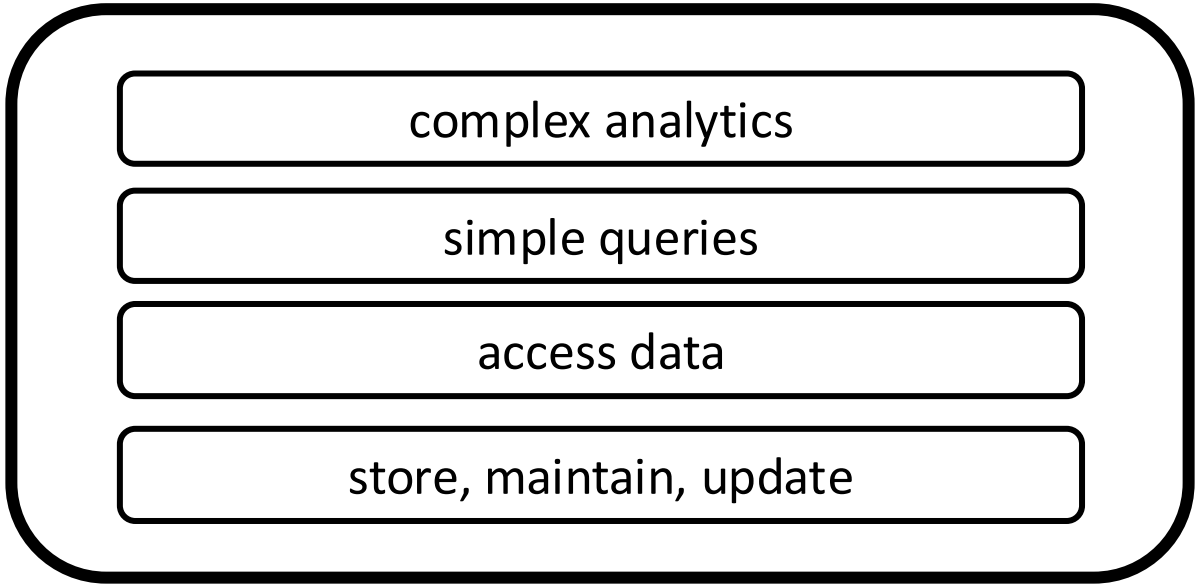

# data systems

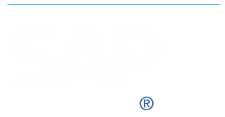





>$200B by 2020, growing at 11.7% every year

[The Forbes, 2016]

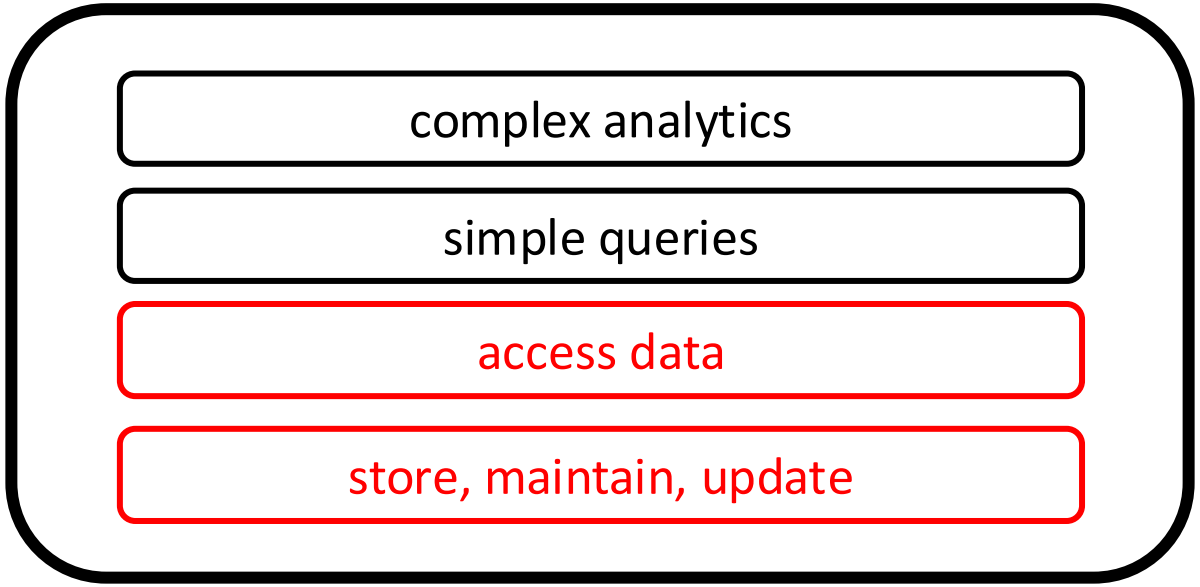


***access methods****

***algorithms** and **data structures**
for organizing and accessing data

# data systems core: storage engines

## main decisions

how to ***store*** data?

how to ***access*** data?

how to ***update*** data?

# let's simplify: **key-value** storage engines

collection of keys-value pairs

query on the key, return both key and value

***remember***








***state-of-the-art*** design

# how general is a key value store?

can we store relational data?

yes! {<primary_key>,<rest_of_the_row>}

example: { **student_id**, { **name**, **login**, **yob**, **gpa** } }

what is the caveat?

**how to index these attributes?**

other problems?

index: { **name**, { **student_id** } }

index: { **yob**, { **student_id**$_1$, **student_id**$_2$, … } }

# how general is a key value store?

can we store relational data?

yes! {<primary_key>,<rest_of_the_row>}

how to efficiently code if we do not know the structure of the "***value***"

index: { **yob**, { **student_id$_1$, student_id$_2$, …** } }

# how to use a key-value store?

***basic interface***

put(k,v)

{v} = get(k)  {$v_1$, $v_2$, …} = get(k)

{$v_1$, $v_2$, …} = get_range($k_{min}$, $k_{max}$)  {$v_1$, $v_2$, …} = full_scan()

c = count($k_{min}$, $k_{max}$)

*deletes: delete(k)*

*updates: update(k,v)* is it different than put?

get set: {$v_1$, $v_2$, …} = get_set($k_1$, $k_2$, …)

# how to build a key-value store?

if we have only ***put*** operations

***append***

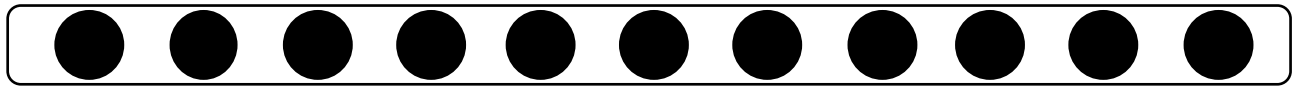

if we mostly have ***get*** operations

***sort***

what about full scan?

range queries?

and then

# can we separate keys and values?

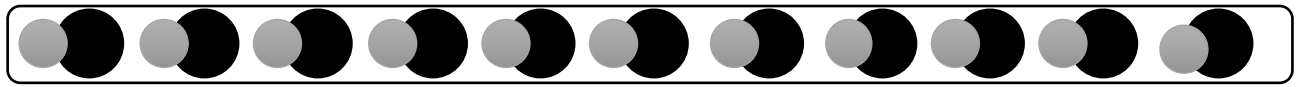

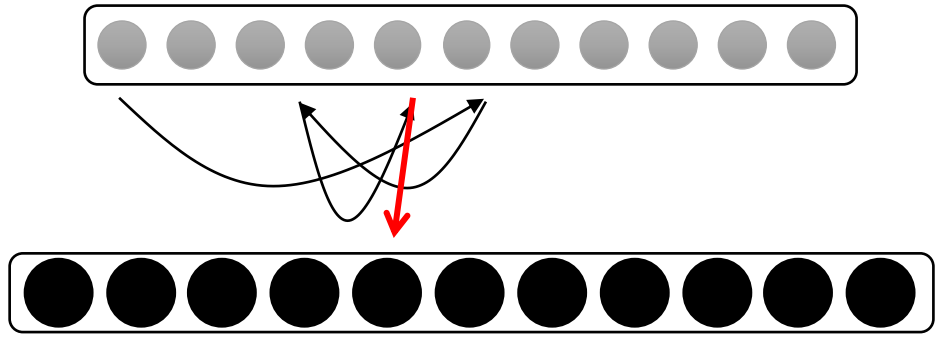

at what price?

locality? code?

read queries
(point or range)

inserts
(or updates)

*sort data*

*simply append*

*amortize sorting cost*

*avoid resorting after every update*

how to bridge?

# LSM-tree
# Key-Value Stores

What are they really?

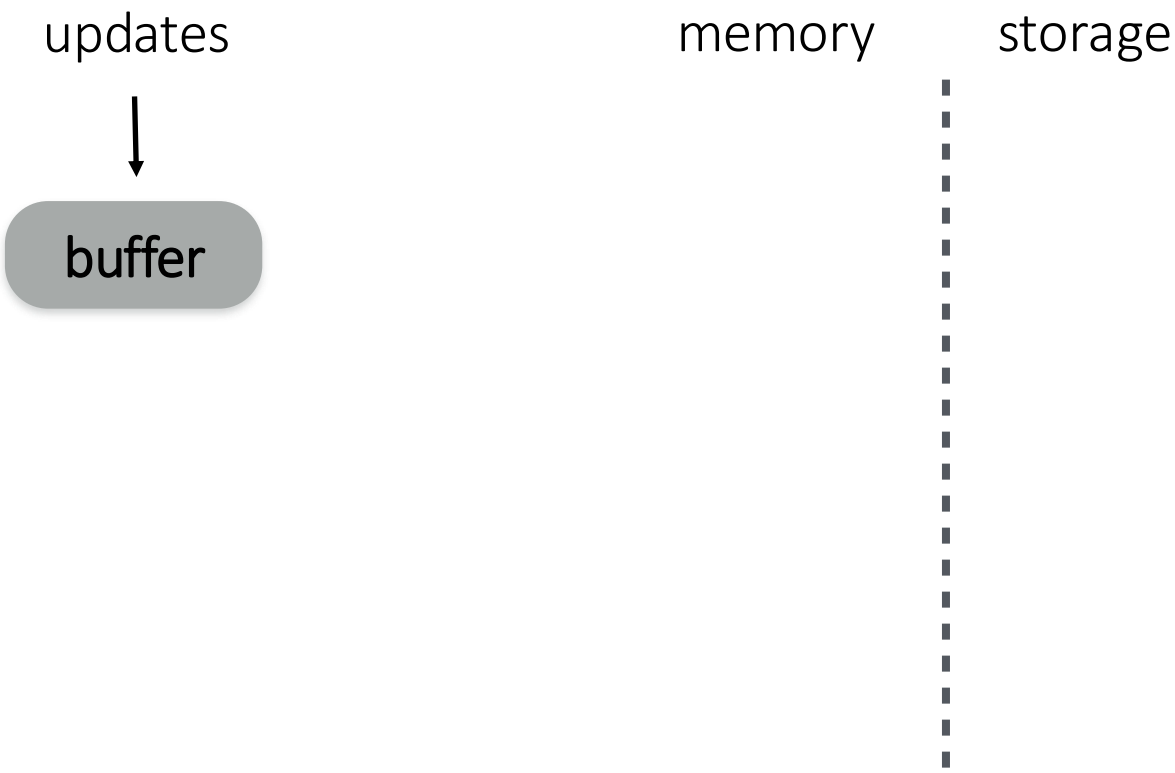

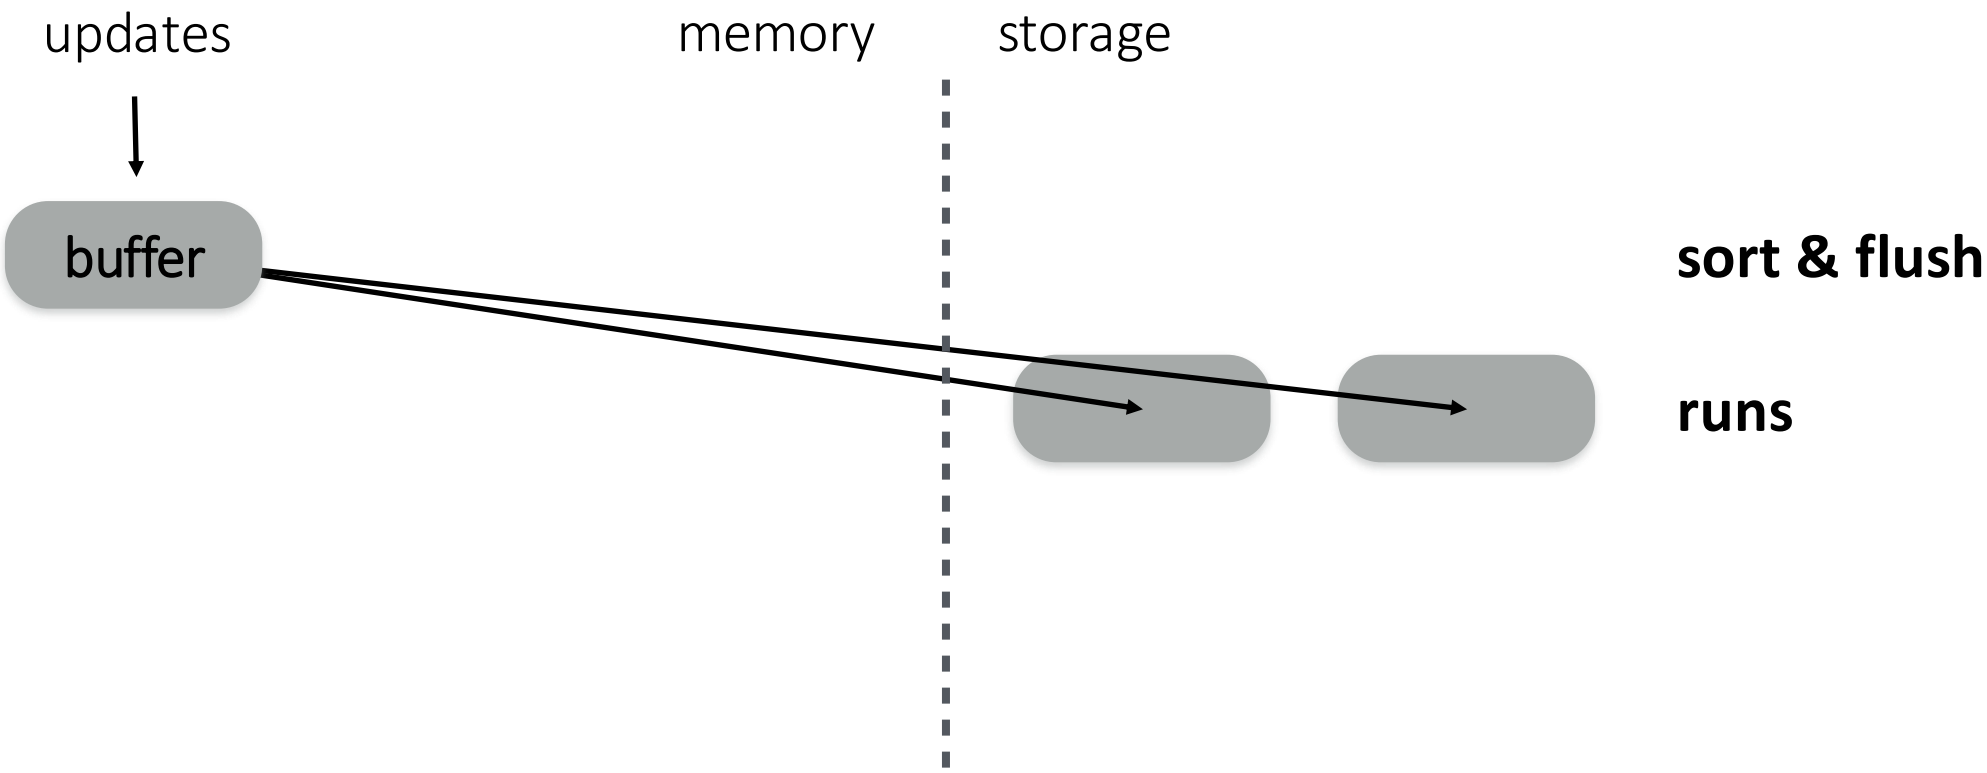

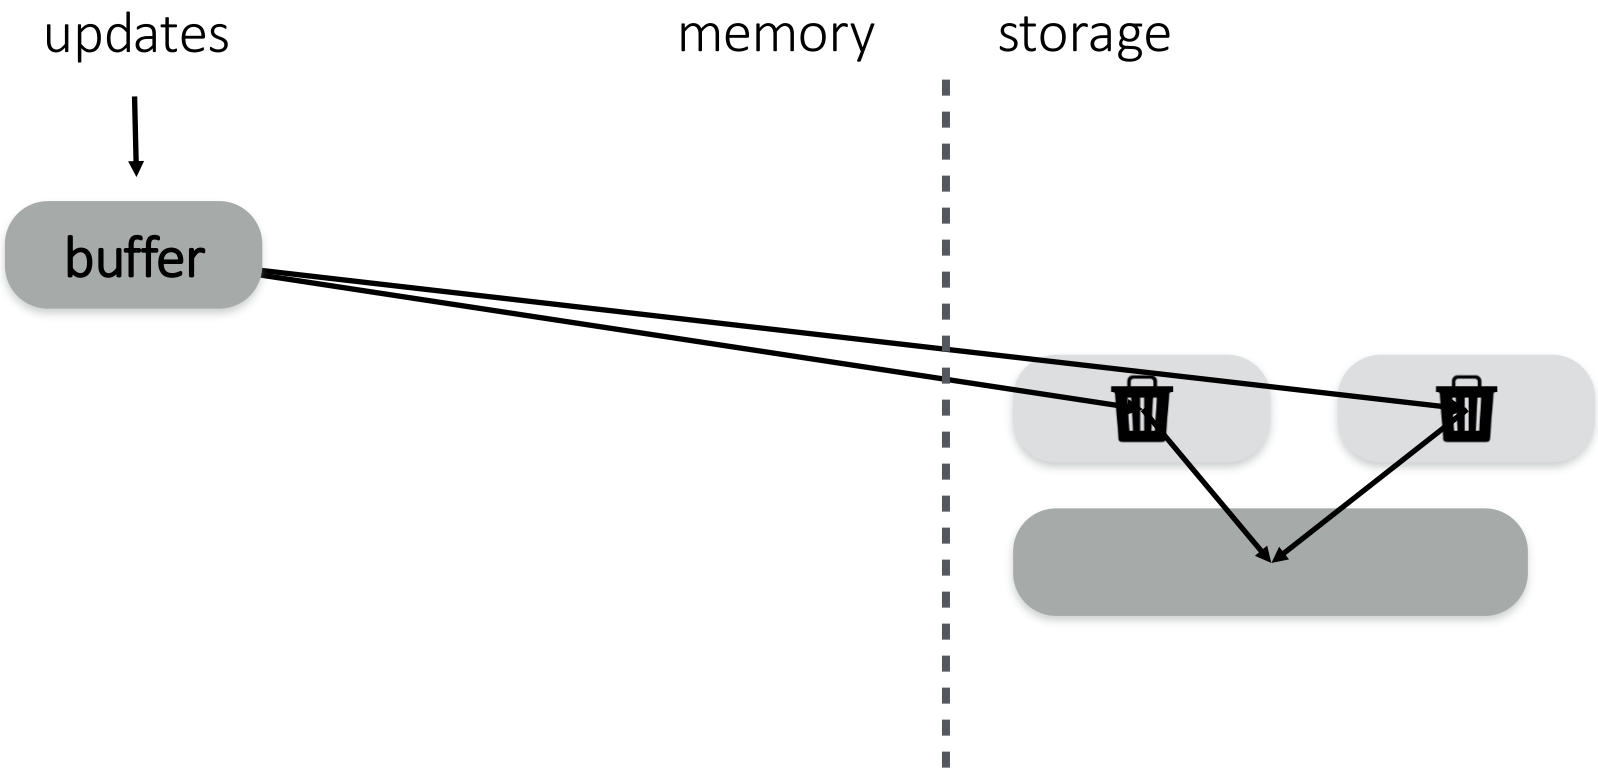


sort & flush

runs

**sort-merge**

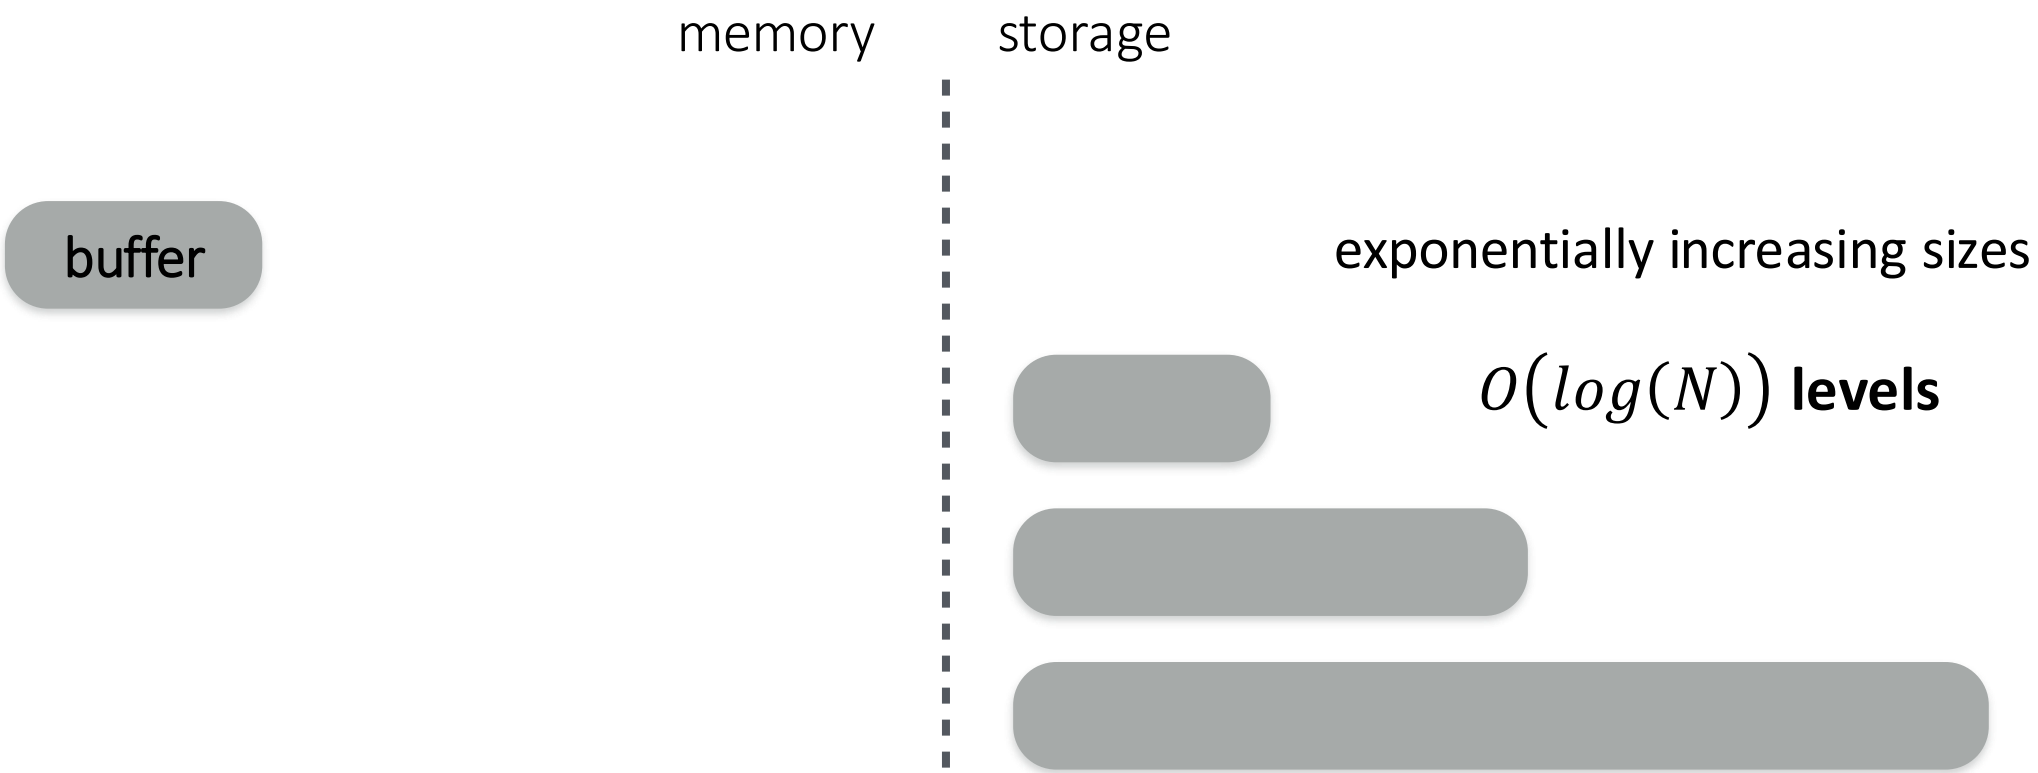

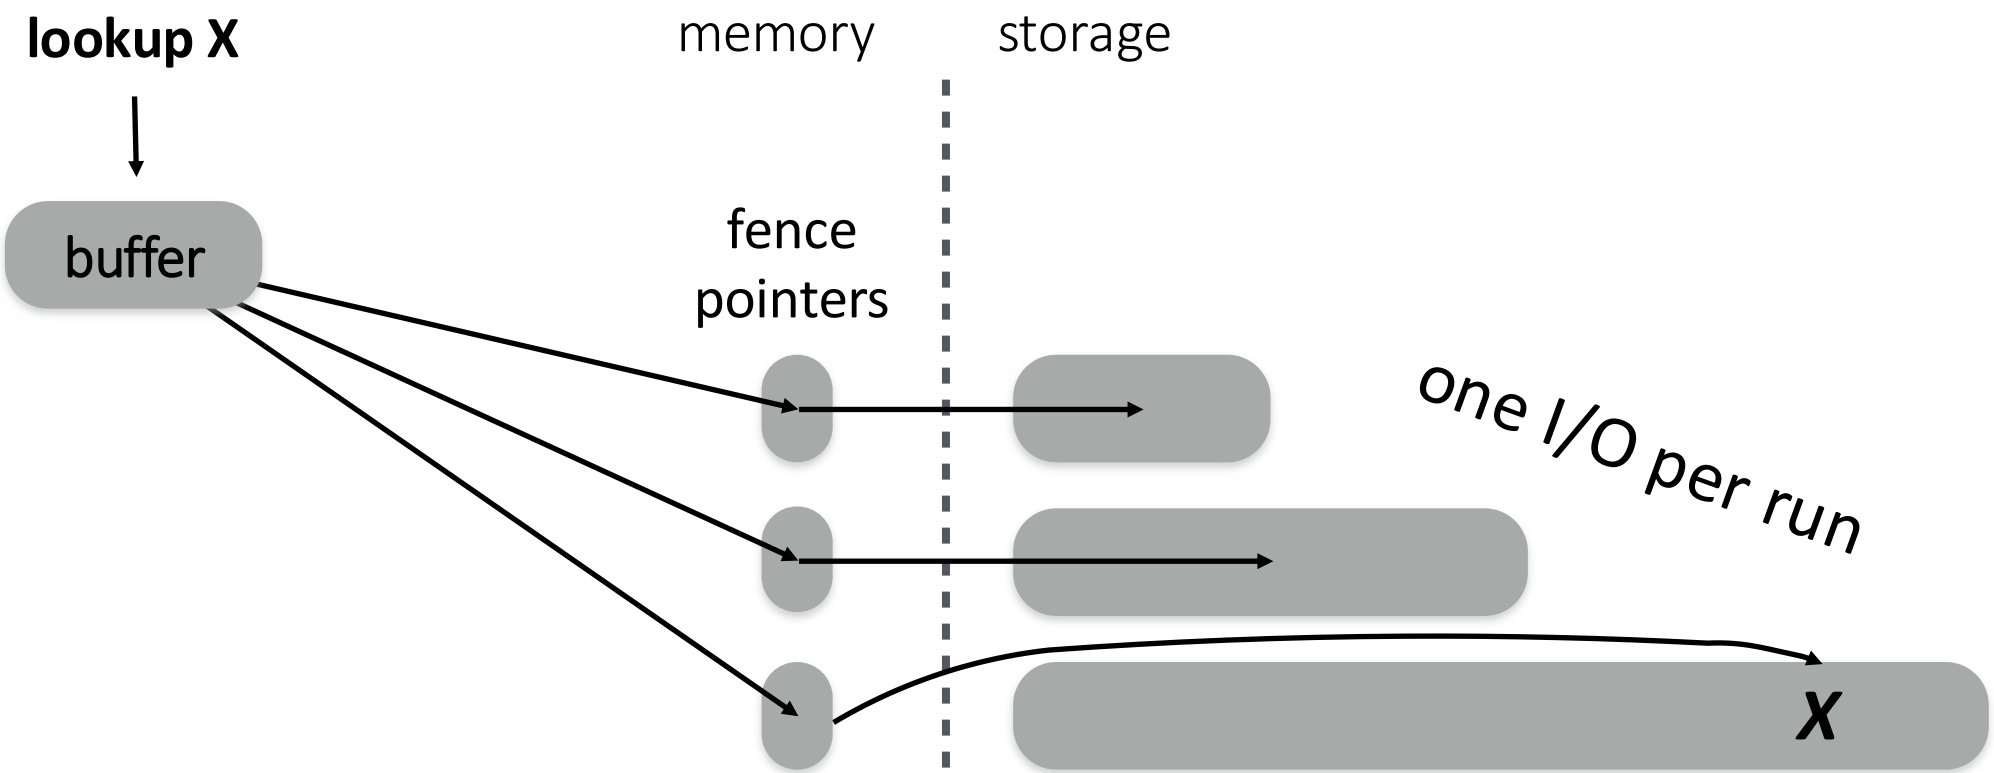

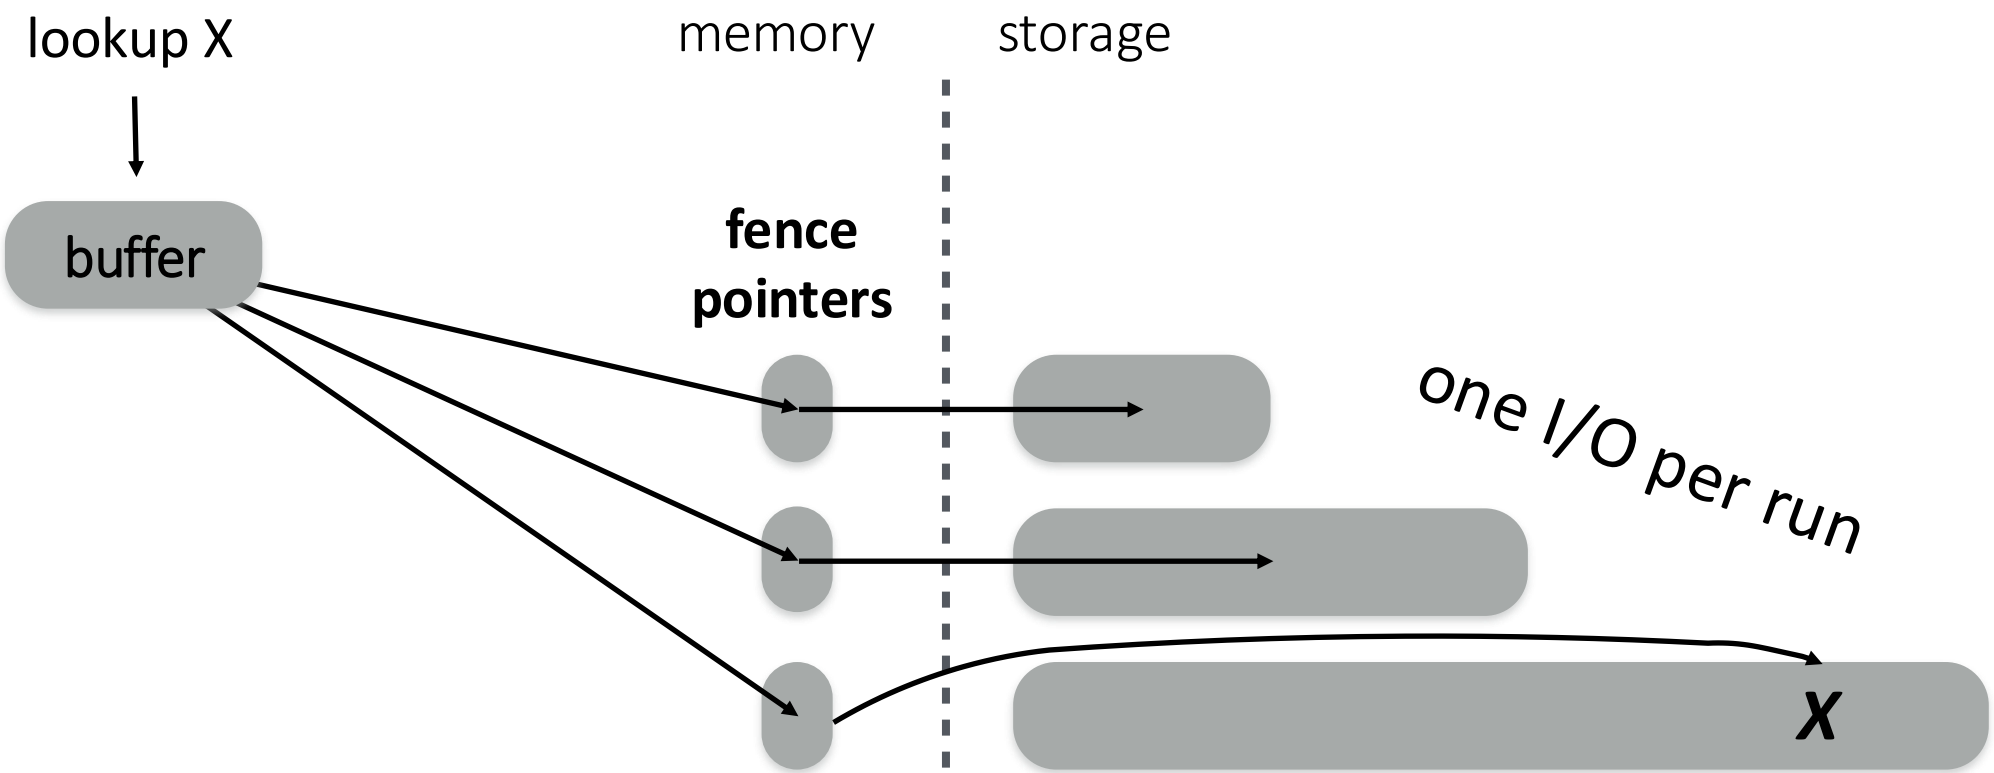

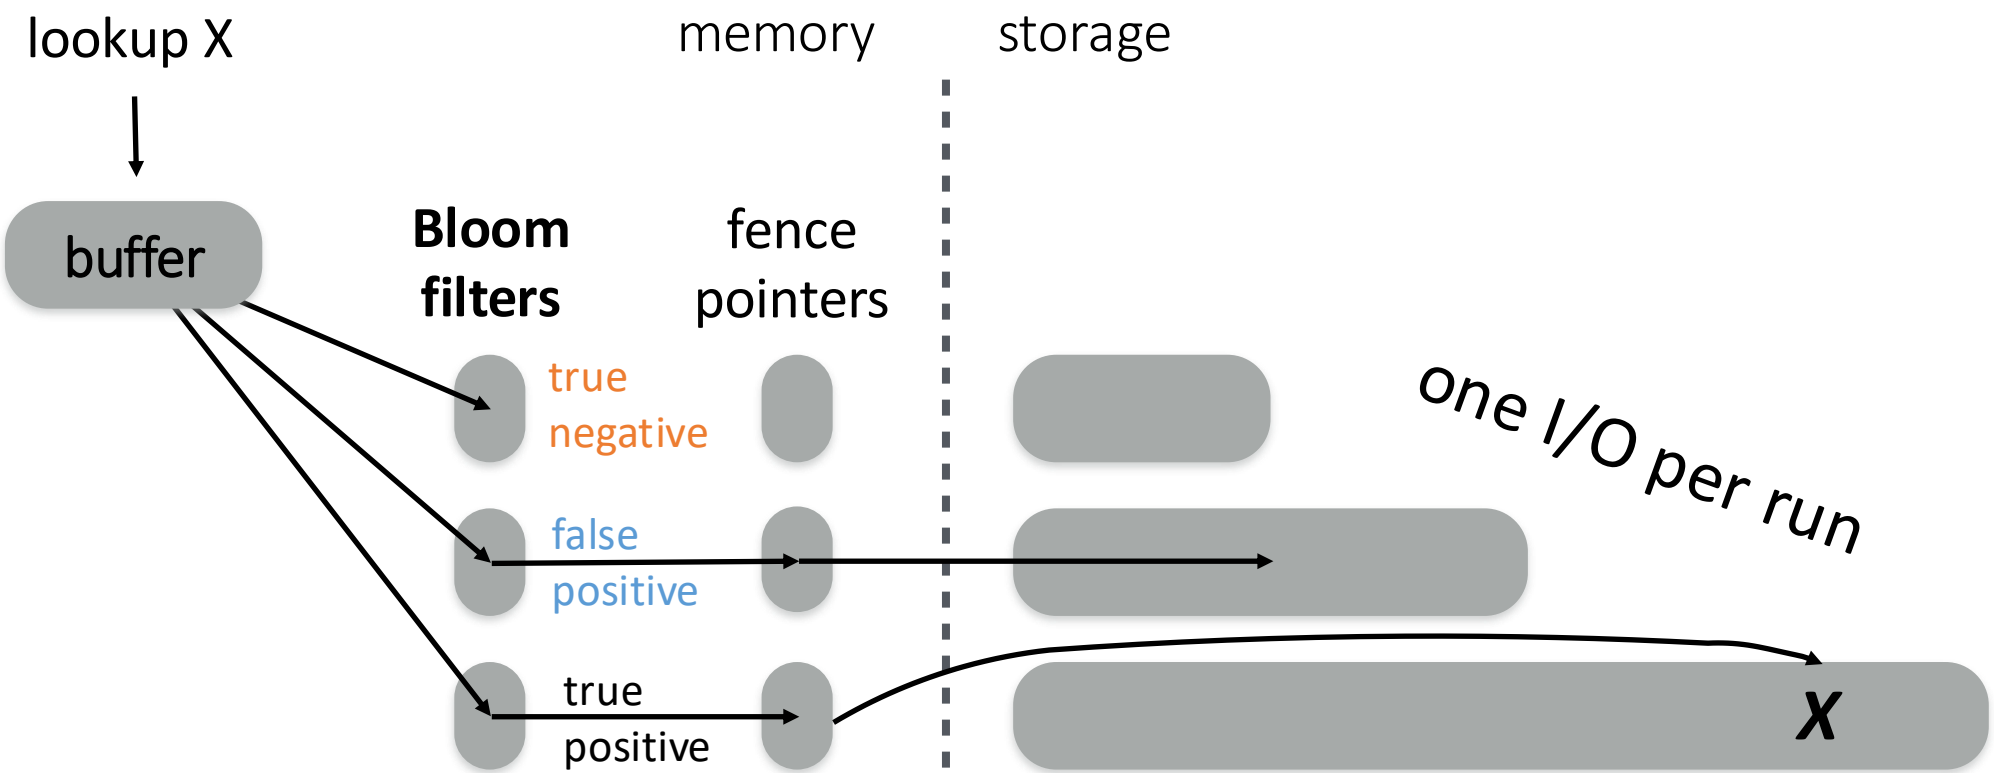

# performance & cost trade-offs

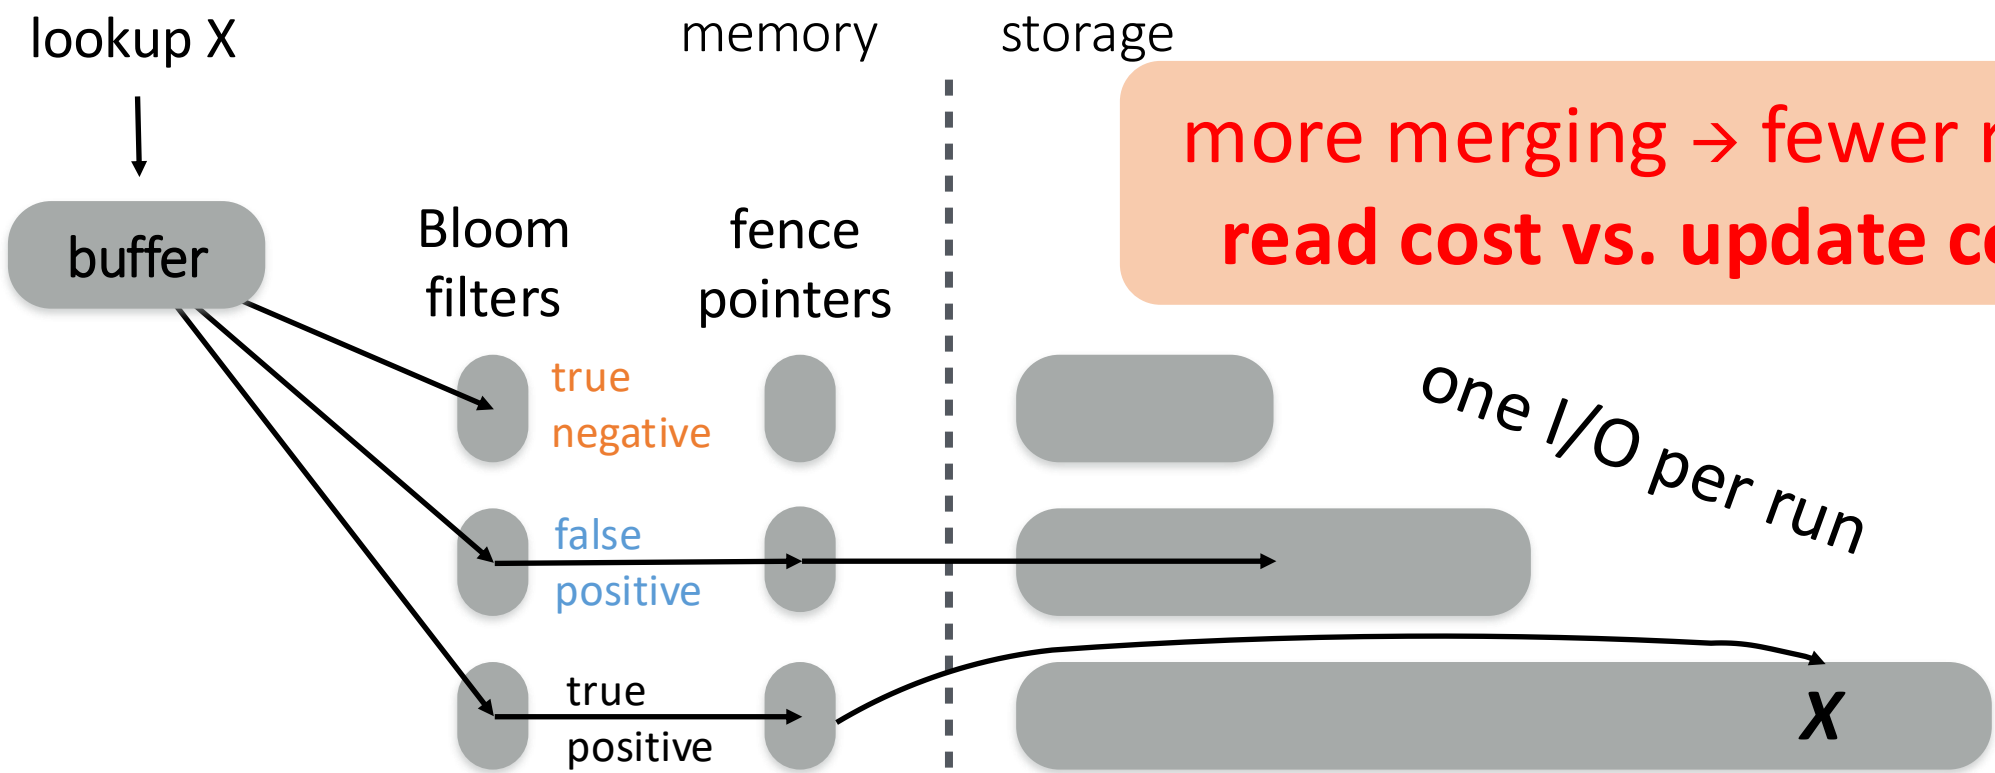


more merging → fewer runs

**read cost vs. update cost**

bigger filters → fewer false positives

**memory space vs. read cost**

# other operations

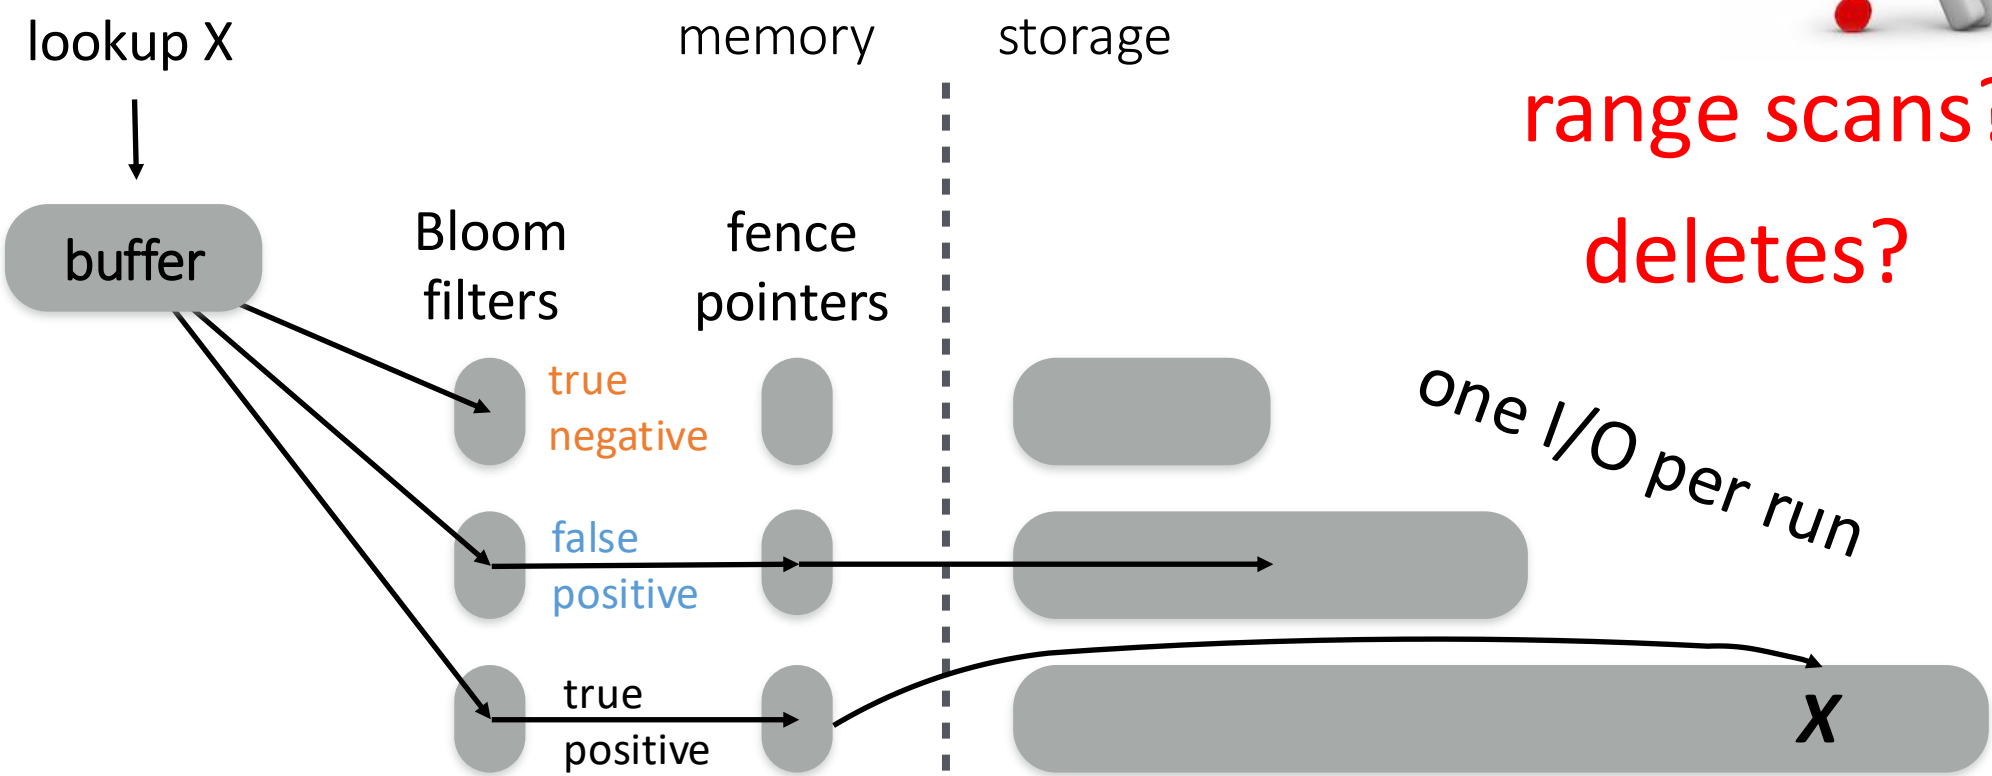


range scans?

deletes?

# remember merging?

updates

memory

storage

buffer

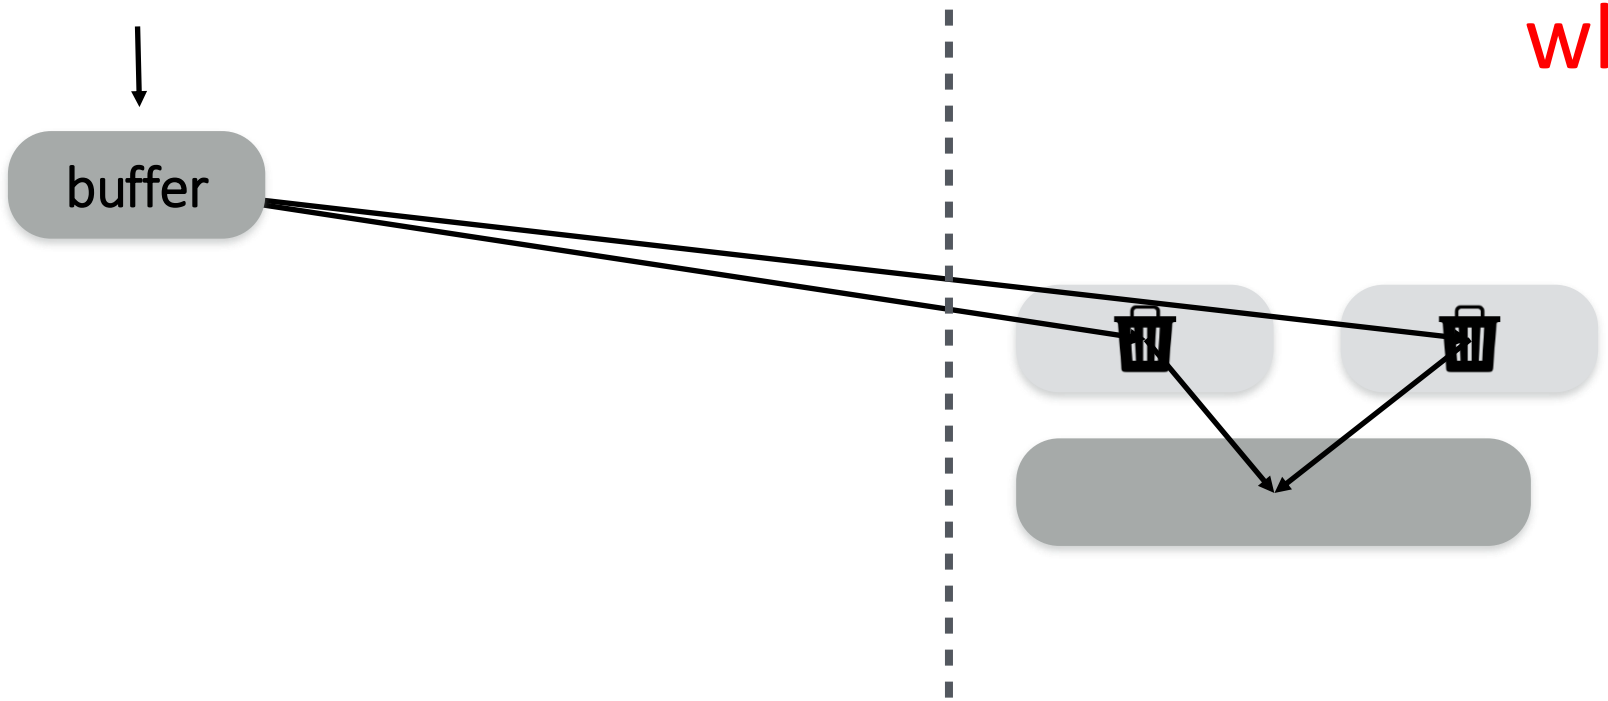

## what strategies?

sort & flush

runs

**sort-merge**

# Merge Policies

**Tiering**
write-optimized

**Leveling**
read-optimized

## Tiering
## write-optimized

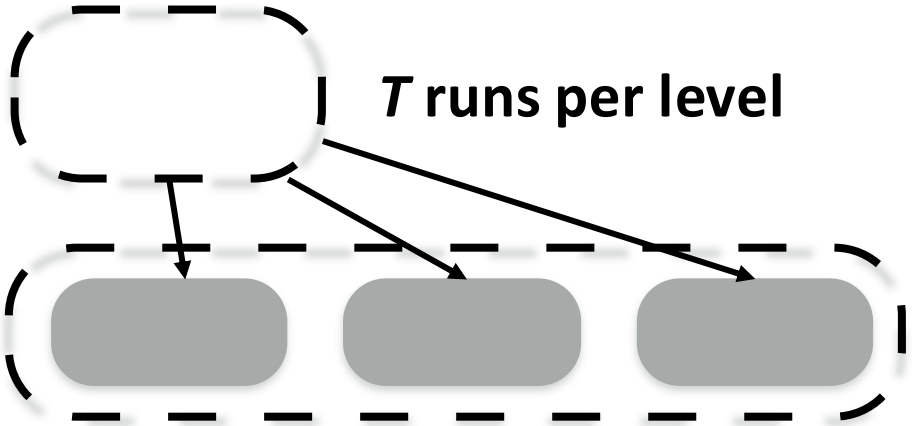


## Leveling
## read-optimized

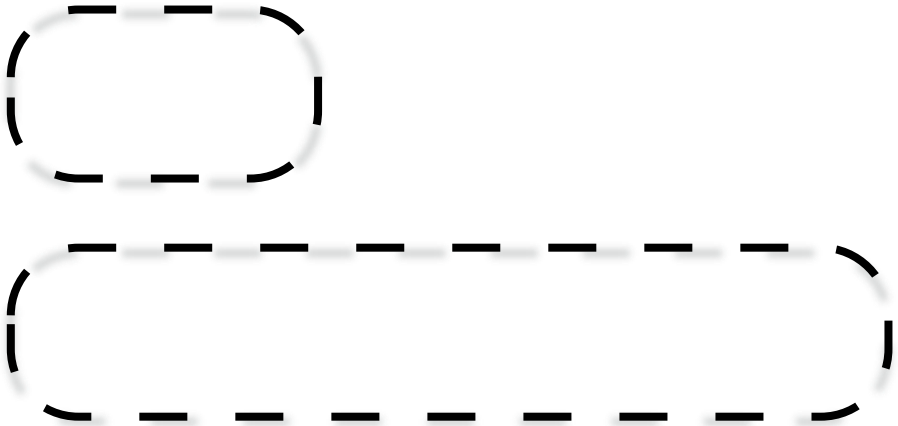

## Tiering
## write-optimized

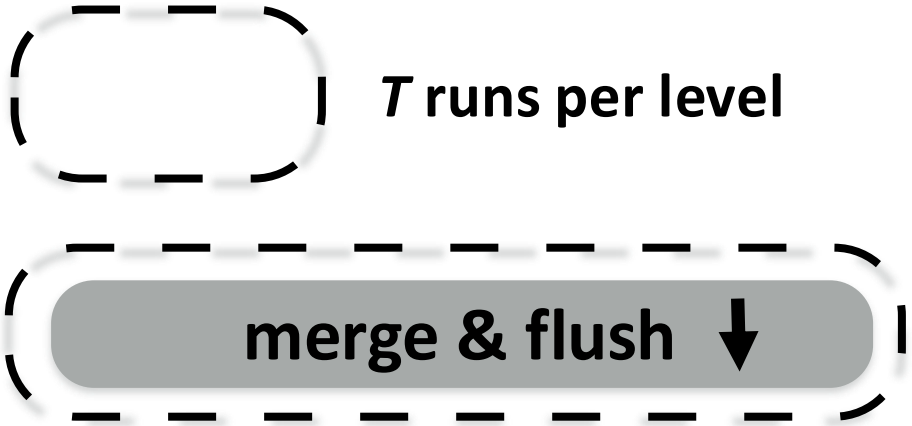


## Leveling
## read-optimized

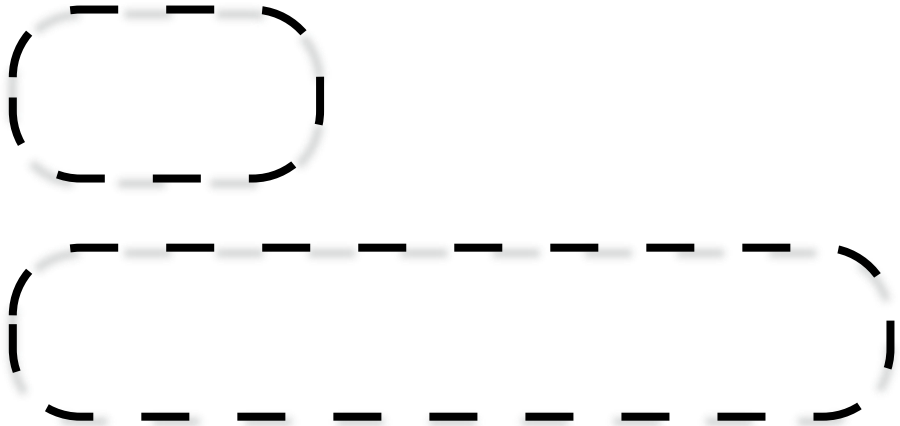

## Tiering
## write-optimized

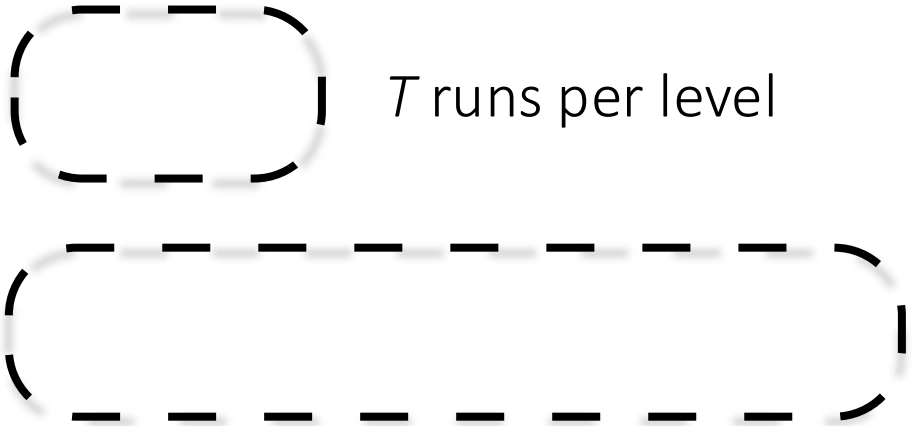


## Leveling
## read-optimized

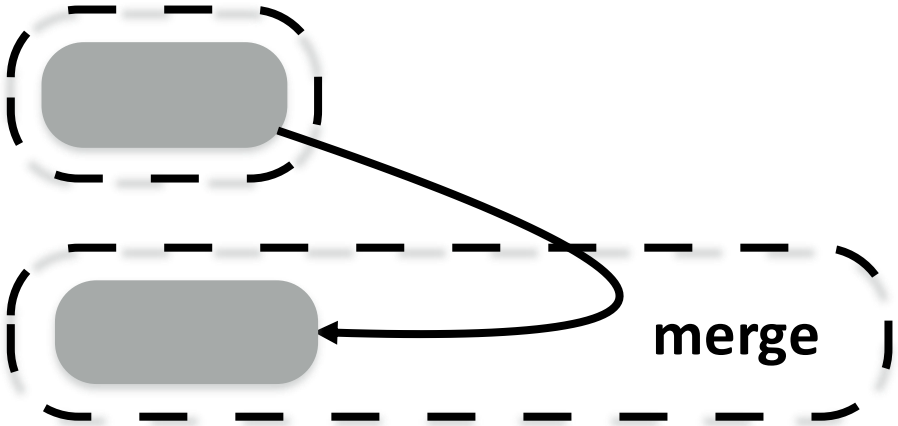

## Tiering
## write-optimized

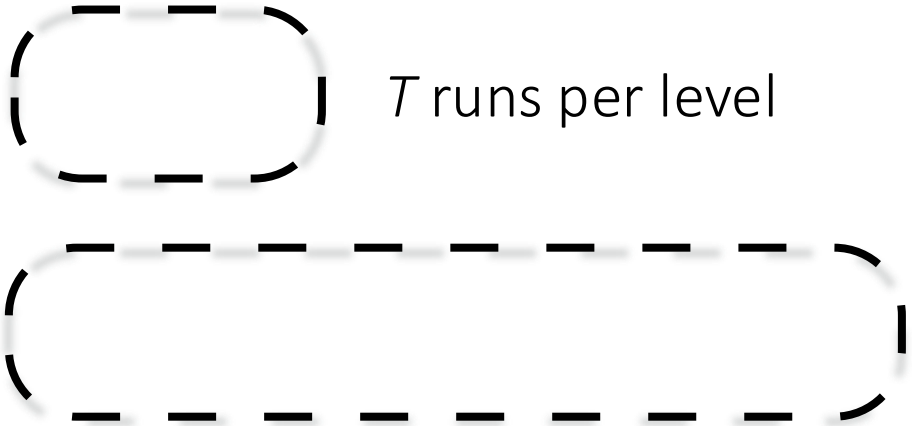


## Leveling
## read-optimized

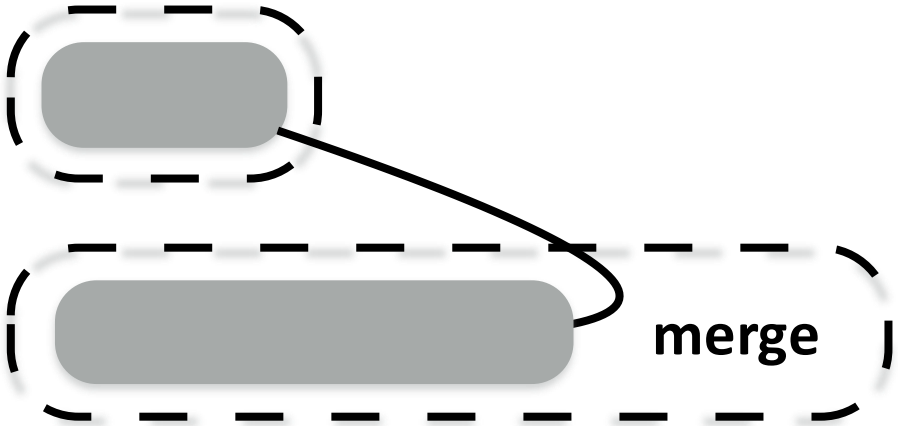

## Tiering
## write-optimized

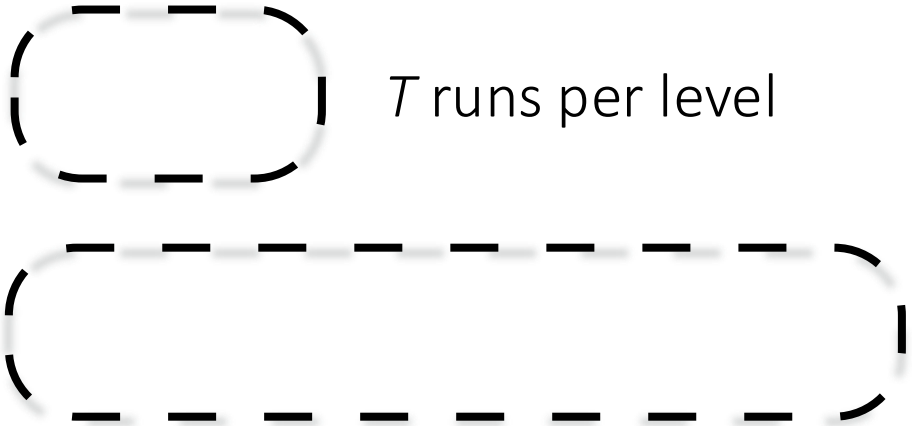


## Leveling
## read-optimized

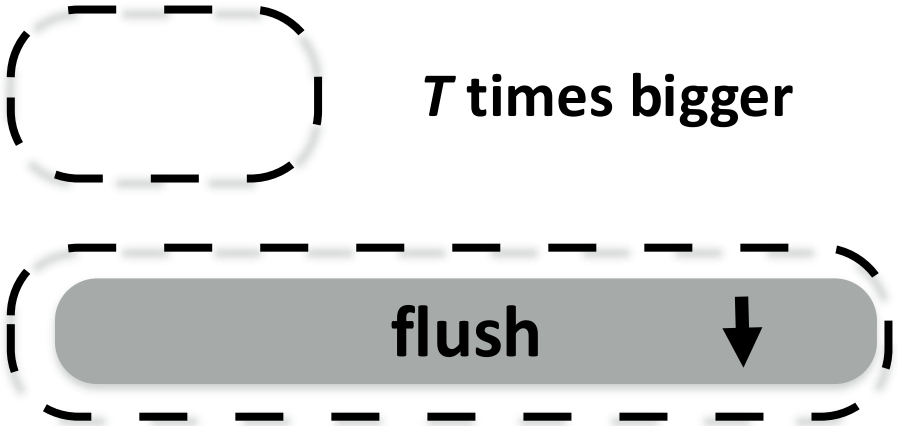

# Systems Project: LSM-Trees

*tuning knobs*

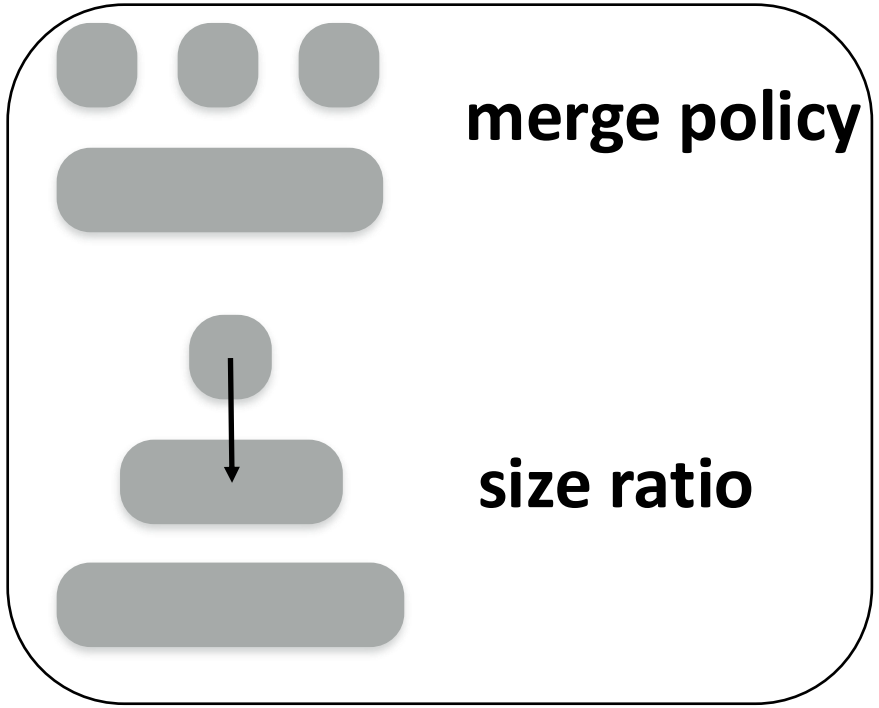


lookup X

memory

storage

buffer

**Bloom filters**

fence pointers

true negative

false positive

true positive

one I/O per run

***X***

# more on LSM-Tree performance

## Tiering
## write-optimized

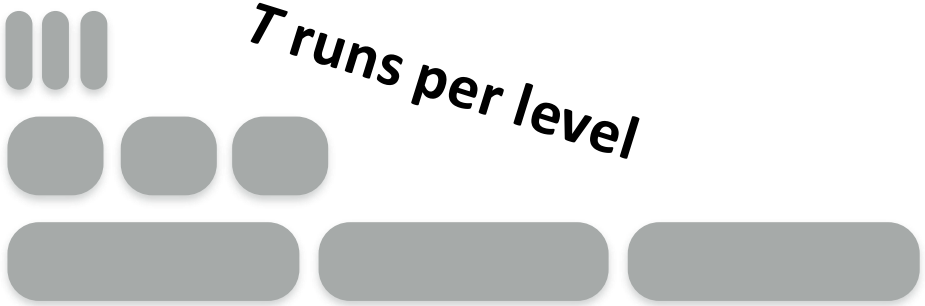


## Leveling
## read-optimized


**lookup cost:** $O\left(T \cdot log_T(N) \cdot e^{-M/N}\right)$

runs per level — levels — false positive rate

$O\left(log_T(N) \cdot e^{-M/N}\right)$

levels — false positive rate

BOSTON UNIVERSITY

## Tiering
## write-optimized

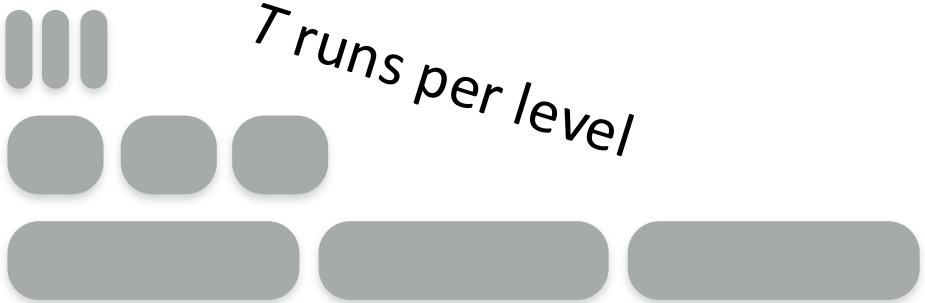


## Leveling
## read-optimized

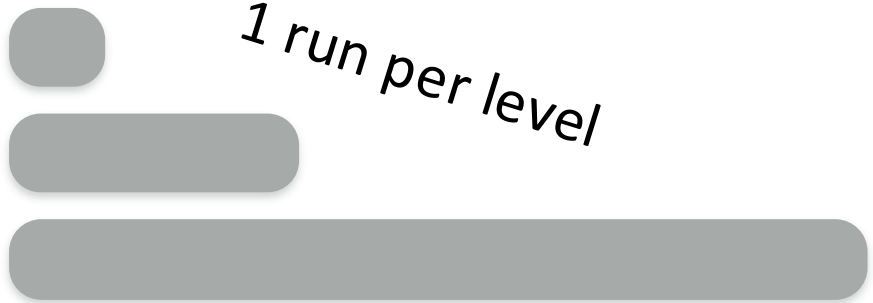


lookup cost: $O\left(T \cdot log_T(N) \cdot e^{-M/N}\right)$

**update cost:** $\boldsymbol{O\left(log_T(N)\right)}$

levels

$O\left(log_T(N) \cdot e^{-M/N}\right)$

$\boldsymbol{O\left(T \cdot log_T(N)\right)}$

merges per level

levels

BOSTON UNIVERSITY

## Tiering
## write-optimized

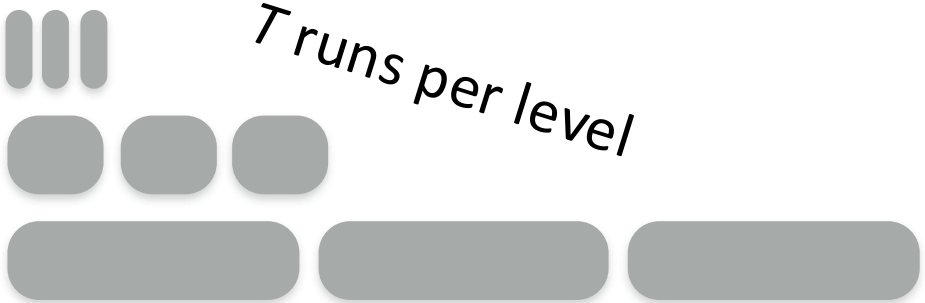


## Leveling
## read-optimized

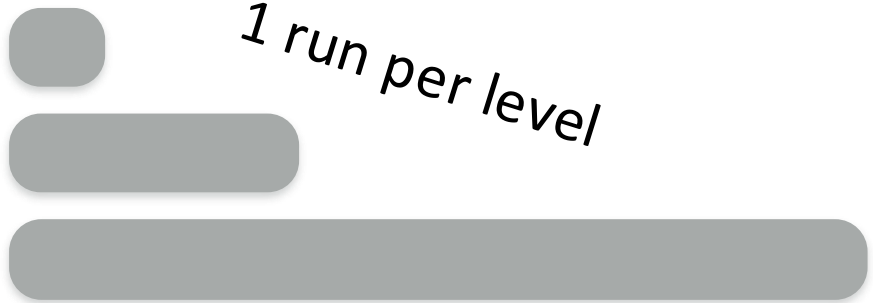


lookup cost: $O\left(T \cdot log_T(N) \cdot e^{-M/N}\right)$ $O\left(log_T(N) \cdot e^{-M/N}\right)$

update cost: $O\left(log_T(N)\right)$ $O\left(T \cdot log_T(N)\right)$

**for size ratio T** ⌄⌄

Tiering
write-optimized

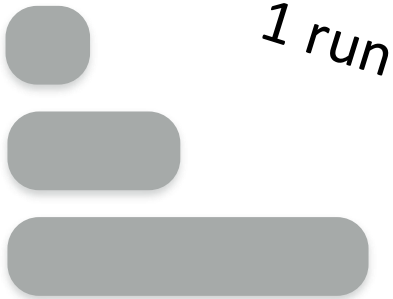


Leveling
read-optimized


lookup cost: $O\left(log_T(N) \cdot e^{-M/N}\right) = O\left(log_T(N) \cdot e^{-M/N}\right)$

update cost: $O\left(log_T(N)\right) = O\left(log_T(N)\right)$

**for size ratio T** ⌄⌄

## Tiering
## write-optimized

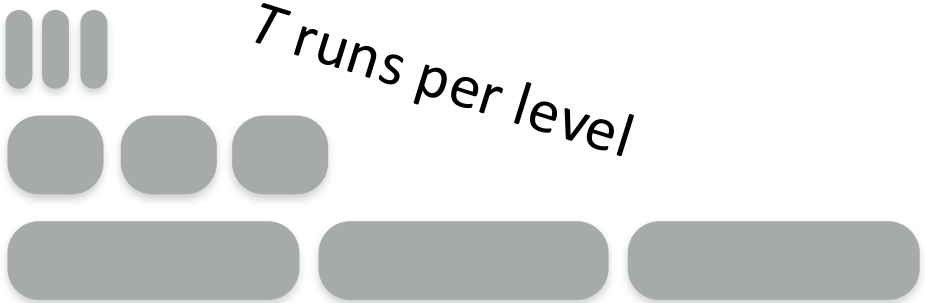


## Leveling
## read-optimized

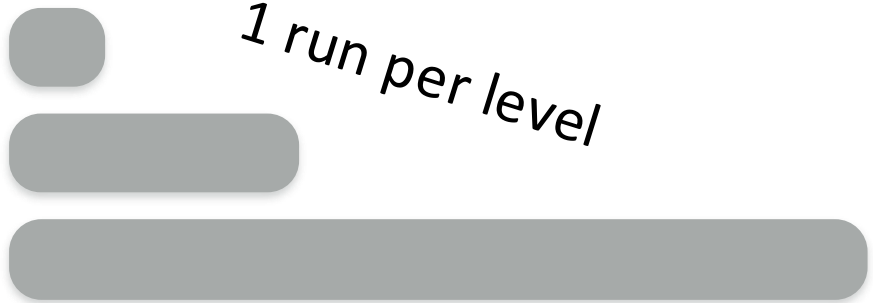


lookup cost: $O\left(T \cdot log_T(N) \cdot e^{-M/N}\right)$ $O\left(log_T(N) \cdot e^{-M/N}\right)$

update cost: $O\left(log_T(N)\right)$ $O\left(T \cdot log_T(N)\right)$

**for size ratio T** ⌃⌃

Tiering
write-optimized

$O(N)$ runs per level

**log**

Leveling
read-optimized

1 run per level

**sorted array**

lookup cost: $O\left(T \cdot log_T(N) \cdot e^{-M/N}\right)$ $O\left(log_T(N) \cdot e^{-M/N}\right)$

update cost: $O\left(log_N(N)\right) = \boldsymbol{O(1)}$ $O\left(N \cdot log_N(N)\right) = \boldsymbol{O(N)}$

**for size ratio T** **N** ≫

BOSTON UNIVERSITY

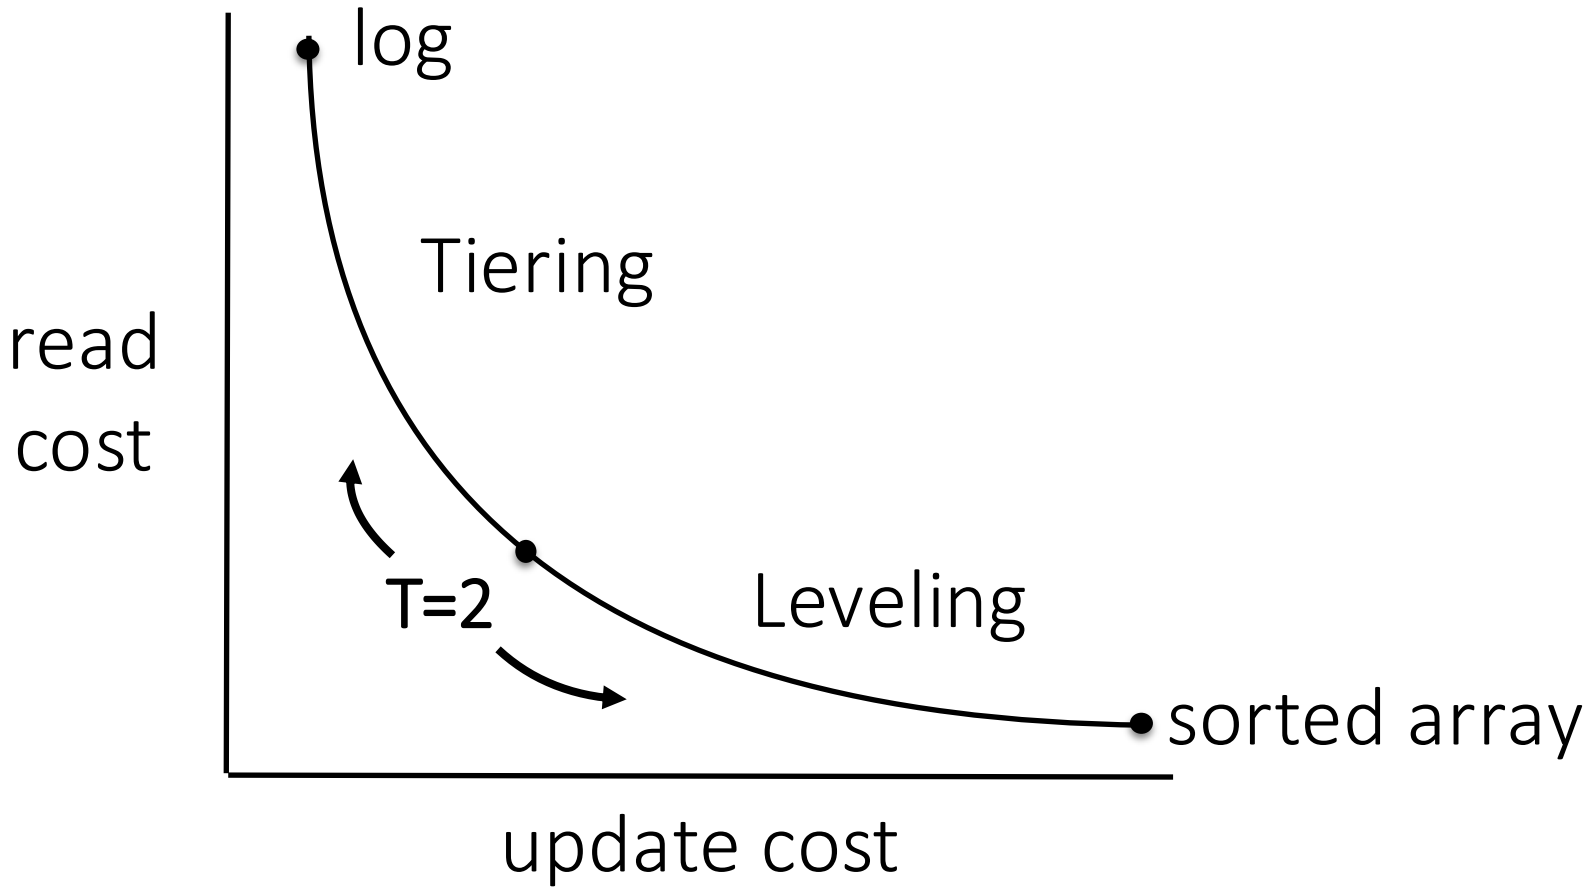


T : size ratio

# Research Question on LSM-Trees

how to reduce temporary space amplification at compaction time?

how to support variable deletion persistent thresholds?

how to address high number of contiguous tombstones?

how to efficiently support range deletes?

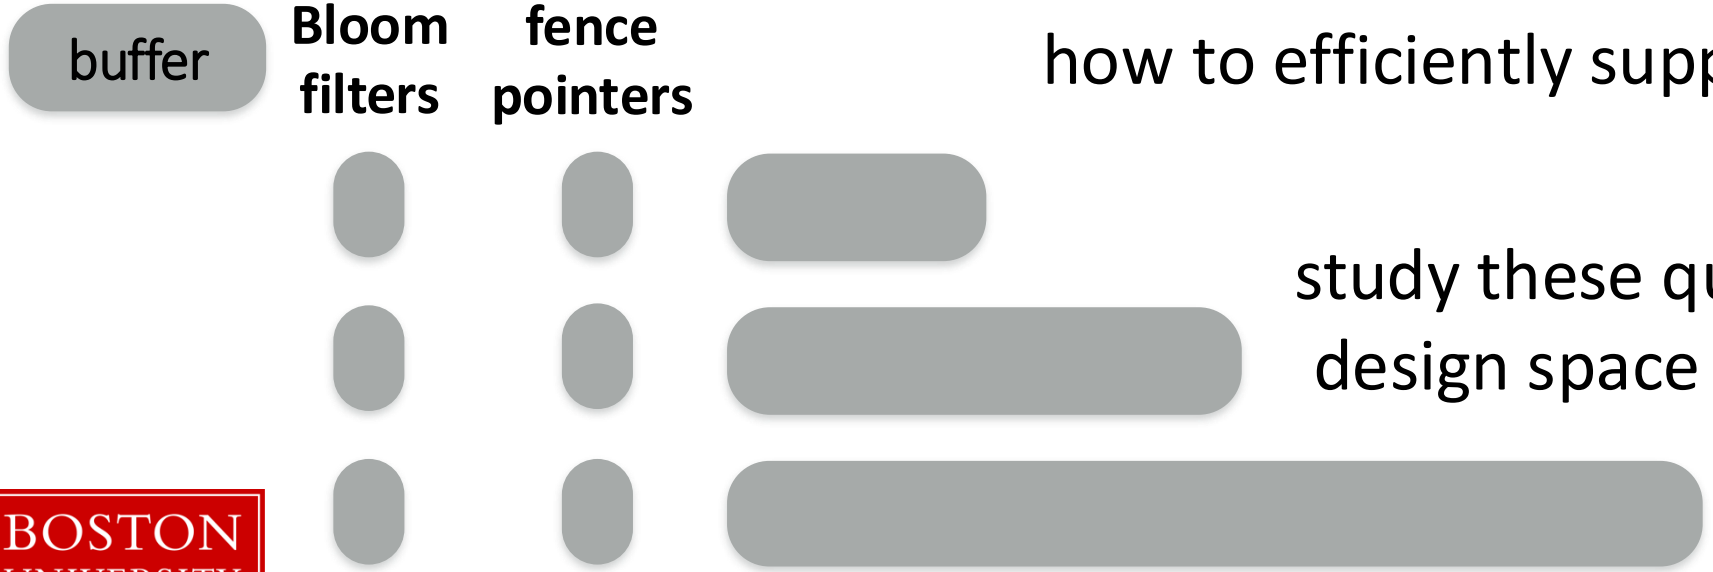


study these questions and navigate LSM design space using Facebook's RocksDB

BOSTON UNIVERSITY

RocksDB

# Systems Project: Bufferpool

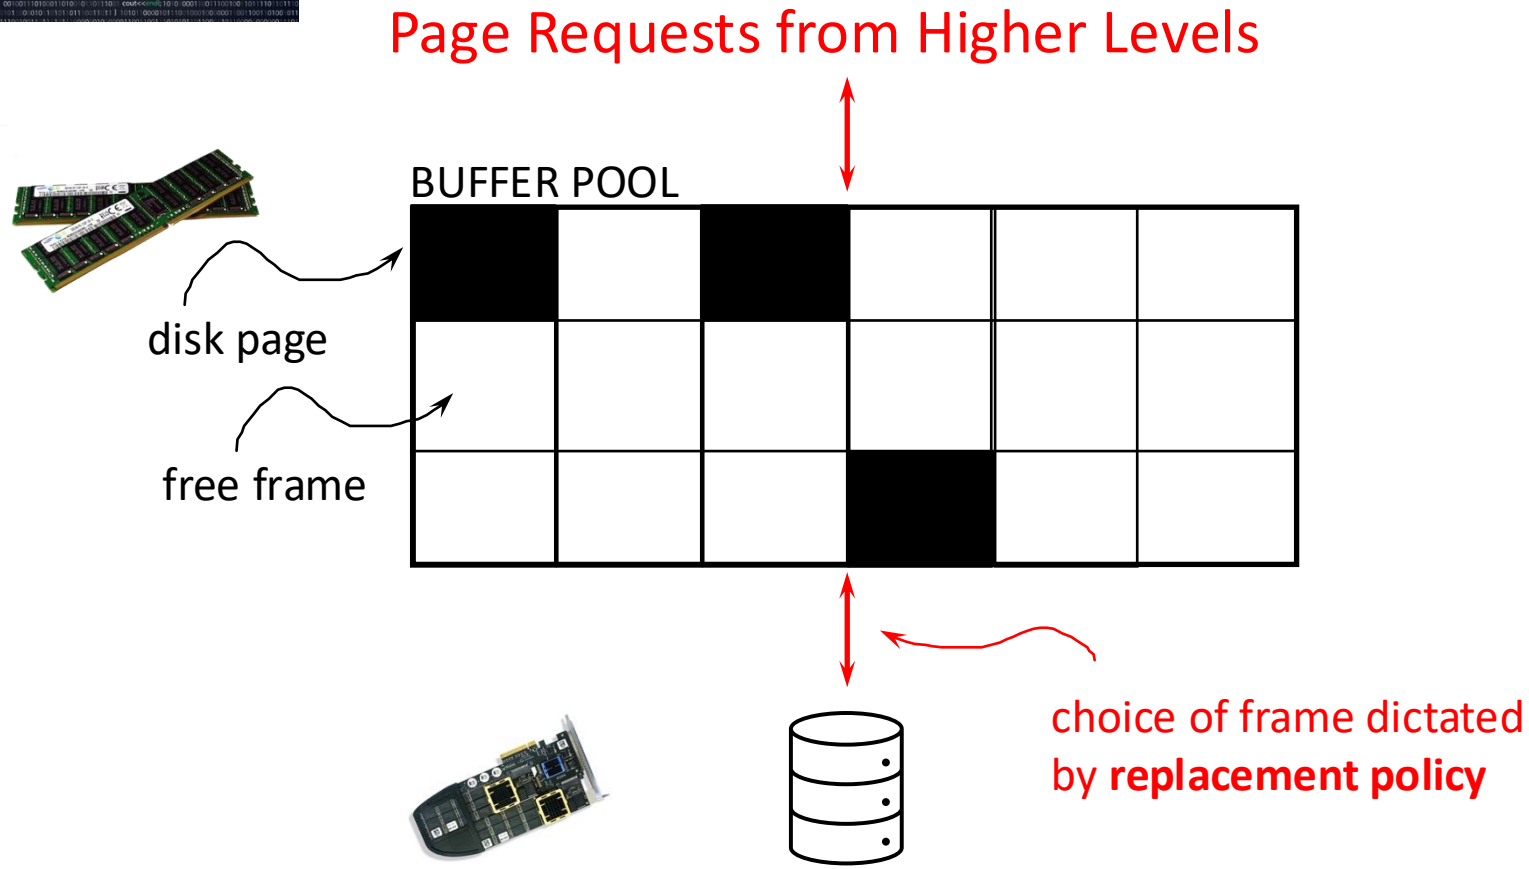


*Implementation of a bufferpool*

- **Application requests** a page
  - If **in the bufferpool** return it
  - If **not in the bufferpool** fetch it from the disk
    - If bufferpool is full select page to **evict**

*Core Idea: Eviction Policy*

- *Least Recently Used*
- *First In First Out*
- *more …*

BOSTON UNIVERSITY

# Research on PostgreSQL


***A state-of-the-art relational database***

How can we implement a sortendess-aware join algorithm?

Integrate light-weight eviction policies in PostgreSQL bufferpool

BOSTON UNIVERSITY

# Research Topics on SSDs

how to design a flexible ZNS SSD interface?

understanding Garbage Collection in applications on ZNS SSDs

understanding the impact of Zone Sizes on ZNS SSDs

# Research Topics on Bitmap Indexing

how to do access path selection in the presence of bitmap indexes?

how to execute index-based joins with bitmap indexes?

# what to do now?

**systems project**
form groups of 3
(speak to me in OH if you want a different arrangement)

**research project**
form groups of 3
pick one of the subjects & read background material
define the behavior you will study and address
sketch approach and success metric
(if LSM-related get familiar with RocksDB)

# what to do now?

**systems project**
form groups of 2
(speak to me in OH if you want to work on your own)

**research project**
form groups of 2

*come to OH/Labs*

submit project 0 this Sunday on 2/1

start working on project 1 (due on 2/15)

submit semester project proposal on 2/22

BOSTON UNIVERSITY

# CS 561: Data Systems Architectures

## class 4

## Systems & Research Project

Prof. Manos Athanassoulis

https://bu-disc.github.io/CS561/