

CS 561: Data Systems Architectures

class 7

Design Tradeoffs in Key-Value Stores

Prof. Manos Athanassoulis

<https://bu-disc.github.io/CS561/>

Do we have a quiz today ... ?

Yes!

what to do now?

- A) read the syllabus and the website
- B) register to Piazza + Gradescope
- C) finish project 0
- D) finish project 1
- E) finish project proposal (due **2/22**)
- F) register for the student-presentations (by **3/6**)

Fast Scans on Key-Value Stores (KVS)

Key-Value Stores are designed for *transactional* workloads (put and get operations)



Analytical workloads require efficient scans and aggregations
(typically offered by column-store systems)



Can we do both in one system?


Why combine KVS and analytical systems?

cheaper and cheaper storage

more data ingestion

need for write-optimized
data structures

what about analytical queries?



Both transactional and analytical systems

Most organizations maintain both

- ***transactional*** systems (often as key-value stores)
- ***analytical*** systems (often as column-stores)

problems?




requires additional expertise and management (e.g., two DBAs)

harder to maintain (more systems, more code)

time consuming data integration/transfer

Log scale!

Scan Time of 50M records (~4GB of data)



Goals of this paper

Bridge the conflicting goals of ***get/put*** and ***scan*** operations

get/put operations need **sparse data structures** → locality is not required, access one object
scans require **locality** (relevant data to be packed together)

we will discuss how to compromise, via the design of ***Tellstore***

how to amend the ***SQL-over-NoSQL*** architecture for mixed workloads

SQL over NoSQL

Elasticity

Snapshot Isolation

e.g., MVCC, no locking, timestamp based

Support for:


Scans

Versioning

Batching

Processing
Layer

Storage
Layer



Scans

selection

projection

(simple) aggregates

shared scans

remember them?

Versioning

multiple versions
through timestamps

garbage collection

discarding old versions
during scans might be costly

Batching

batch several
requests to the
storage layer

amortize the
network time



Challenges

scans vs. get/put

Scans need columnar locality

#1, John, 2/4/88, Boston

get/put need row-wise locality

why?



2/4/88
2/1/87
7/7/93
4/1/92
3/9/91
9/3/96

Challenges

scans vs. get/put

#1, John, 2/4/88, Boston, v1

#1, John, 2/4/88, Cambridge, v2

scans vs. versioning



versioning reduces locality in scans

checking for the latest version in scans needs CPU time

Challenges

scans vs. get/put

scans vs. versioning

scans vs. batching

batching **multiple scans or multiple put/get** requests is ok

but ...

batching scans and puts/gets is a bad idea!

why?



puts/gets need fast predictable performance

scans inherently have high and variable latency

How to design KVS for efficient scans?

Key design decisions

(A) Updates

(B) Layout

(C) Versioning

How to design KVS for efficient scans?

Key design decisions


(A) Updates *in-place*



How to design KVS for efficient scans?

Key design decisions

(A) Updates *in-place* *log-structured*



How to design KVS for efficient scans?


Key design decisions

(A) Updates

in-place

log-structured

delta-main






How to design KVS for efficient scans?

Key design decisions


(A) Updates *in-place* *log-structured* *delta-main*


(B) Layout *column* *PAX (columnar per page)*



Small detour: page layouts

middle ground?





Partition Attributes Across (PAX)

Middle ground?




Decompose a slotted-page internally
in mini-pages per attribute

✓ Cache-friendly

- Brings only relevant attributes to cache



- Compatible with slotted-pages?
- Same update abstraction?
– (insert in a page)




How to design KVS for efficient scans?

Key design decisions

(A) Updates *in-place* *log-structured* *delta-main*

(B) Layout *column (PAX)* *row*




How to design KVS for efficient scans?

Key design decisions

(A) Updates *in-place* *log-structured* *delta-main*

(B) Layout *column (PAX)* *row*

(C) Versioning *clustered*



How to design KVS for efficient scans?


Key design decisions

(A) Updates *in-place* *log-structured* *delta-main*

(B) Layout *column (PAX)* *row*

(C) Versioning *clustered*

chained



How to design KVS for efficient scans?

Key design decisions

(A) Updates *in-place* *log-structured* *delta-main*

(B) Layout *column (PAX)* *row*

(C) Versioning *clustered* *chained*

what comes as a result of versioning?



Garbage Collection (GC)

- (A) Periodic *separate dedicated thread(s)*
- (B) Piggy-backed GC during scans
 - increases scan time* *but frequently read tables benefit*
 - avoids re-reading for GC (since data is already accessed)*

Design Space

| | | | | | | |
|-----------------------|---|---------------------|---|------------------|---|---------------------|
| Updates | X | Layout | X | Versioning | X | GC |
| <i>in-place</i> | | <i>column (PAX)</i> | | <i>clustered</i> | | <i>periodic</i> |
| <i>log-structured</i> | | <i>row</i> | | <i>chained</i> | | <i>piggy-backed</i> |
| <i>delta-main</i> | | | | | | |



hybrid designs are also valid!
should we consider all possible designs?

Design Space

Updates



Layout



Versioning



GC

in-place

column (PAX)

clustered

periodic

log-structured

row

chained

piggy-backed

delta-main

some combinations can be discarded:

log-structured & column **worse than** delta-main & column

log-structured & clustered **worse than** log-structured & chained

note that each combination here represents multiple options

| <i>Dimension</i> | <i>Approach</i> | <i>Advantages</i> | <i>Disadvantages</i> |
|------------------|-------------------------|-------------------------|----------------------------|
| <i>Update</i> | update-in-place | storage | versioning, concurrency |
| | log-structured | storage, concurrency | GC |
| | delta-main | | compromise |
| <i>Layout</i> | column (PAX) | scan | get/put |
| | row | get/put | scan |
| <i>Versions</i> | clustered | get/put | GC |
| | chained | GC | scan |

Table 2: Design Tradeoffs

Design Space

Updates \times Layout \times Versioning \times GC

in-place

log-structured

delta-main

column (PAX)

row

clustered

chained

periodic


piggy-backed

focus on two extremes:

(1) log-structured & row & chained

TellStore-Log

Design Space



focus on two extremes:

- (1) log-structured & row & chained
- (2) delta-main & column & clustered


TellStore-Log

TellStore-Col


TellStore-Log

one log per table (locality for scans)
inserts, updates, and deletes are all logged

lock-free hash table




TellStore-Log Insertion



a record is considered valid after the hash table pointer is updated

the log contains *rows*


TellStore-Log Update



previous pointer


the log contains *rows*

TellStore-Log Delete



the log contains *rows*

TellStore-Log Garbage Collection



invalidated entries hurt scan performance

during scan, health (% of invalid entries) per page is calculated
if $\text{health} < \text{threshold}$, page is re-written in the head of the log &
update hash table & old page is reclaimed

the log contains *rows*

TellStore-Log in a nutshell

log-structure: efficient puts

hash-table: efficient gets (always points to the latest entry)

snapshot Isolation: high throughput, no locks needed

self-contained log: efficient scans (valid from/to needed)


lazy GC: Optimize tables that are scanned

TellStore-Col

four data structures


before inserting we query
new entries go to insert log
updates go to update log

main is read-only
columnar: read-optimized



logs are write-optimized
row-oriented: append-only

TellStore-Col Layout



fixed-size data is stored in columnar format

variable-size data is indexed in columnar format
but stored in row-wise format

why row-wise?




- (1) faster materialization (contiguous copying)
- (2) less metadata (one offset for many columns)

TellStore-Col Versioning

in main storage multiple versions
are **clustered** contiguously


newest pointers
may exist




previous pointers may
exist only in update log

logs work like before
and they are row-wise


TellStore-Col Insertion




TellStore-Col Update



TellStore-Col Update



TellStore-Col Garbage Collection



dedicated thread
(conversion from row to column)

all main pages with invalid entries

all pages from insert log + update
to main

run GC frequently + truncate logs

TellStore-Col in a nutshell

delta-main: compromise between puts and scans

hash-table: efficient gets (always points to the latest entry, may need one more pointer to follow)


PAX layout: minimize disk I/O, maintain locality for scans

separate insert/update logs: efficient GC

eager GC: improve scans

Implementation Details

scans are assigned to dedicated threads




scan coordinator for shared scans

Implementation Details

efficient predicate evaluation via code generation and predicate pushdown

all queries in CNF

reuse work



Yahoo! Cloud Serving Benchmark# (YCSB#)

based on YSCB, a put/get benchmark

main_table (P, A, B, C, D, E, F, G, H, I, J) P: 8-byte key | A-H: 2-bytes, 4-bytes, 8-bytes | I-J: strings 12-16 bytes

- *Query 1:* A simple aggregation on the first floating point column to calculate the maximum value:

```
SELECT max(B) FROM main_table
```


- *Query 2:* The same aggregation as Query 1, but with an additional selection on a second floating point column and selectivity of about 50%:

```
SELECT max(B) FROM main_table  
WHERE H > 0 and H < 0.5
```


- *Query 3:* A selection with approximately 10% selectivity:

```
SELECT * FROM main_table  
WHERE F > 0 and F < 26
```

Experiments: Transactional Workload



(a) Get/Put Workload




(b) Get-only workload

Figure 8: Exp 1, Throughput: YCSB, TellStore Variants and Kudu, Vary Storage Nodes


Kudu is used as it was the most competitive to begin with

All TellStore approaches are not that far!

Experiments: Scans



(a) YCSB# Query 1




(b) YCSB# Query 3


Figure 10: Exp 3, Response Time: YCSB#, Vary Storage Nodes

Q3 does not have projections,
so no benefit from columnar


Experiments: Mixed Workload



(a) TellStore vs Kudu



(b) TellStore, Scale Get Requests



(c) TellStore, Scale Get/Put Requests

Figure 11: Exp 4, Response Time: YCSB# Query 1, 4 Storage Nodes

Contrary to competition,
scan perf. is stable with
more gets/puts

In the absence of updates
TellStore scales perfectly:
scans+gets go to different
cores

With 50% updates
eventually logging wins

Things to remember

KVS vs. Scans: how to compromise, navigate the design space

- ✓ delta-main vs. log-structure
- ✓ chained vs. clustered versions
- ✓ row-major vs. column-major
- ✓ lazy vs. eager GC

F2: Designing a Key-Value Store for Large Skewed Workloads

Konstantinos Kanellis*

University of Wisconsin-Madison

kkanellis@cs.wisc.edu

Badrish Chandramouli

Microsoft Research

badrishc@microsoft.com

Shivaram Venkataraman

University of Wisconsin-Madison

shivaram@cs.wisc.edu

A new set of requirements:

point put/get operations (need for high throughput, ideally full in-memory for hot data)

larger-than-memory working set (so full in-memory not possible)

high skew in access patterns (reads/writes)

The target application has **no scans!**

which design should they follow?



Key-value stores

Log-structured Merge (LSM) Trees

- Handle *larger-than-memory* workloads
- Organized in levels; first is in-memory
- Support both *point* & range queries
- Avoid I/O by employing (Bloom) *filters*
- Judicious use of main memory

Point-optimized Stores

- Focus on use-cases like *web caching*
- Large *in-memory* index structures
- Latch-free* concurrent designs
- Saturate I/O (even for NVMe SSDs)
- Very high* throughput (>1M ops/sec)



RocksDB



SPLINTERDB

KVel: the Design and Implementation of a Fast Persistent Key-Value Store

Baptiste Lepers
University of Sydney

Karan Gupta
Nutanix

Abstract

Modern block-addressable NVMe SSDs provide much higher bandwidth and similar performance for random and sequential access. Persistent key-value stores (KV) designed for earlier storage devices, using either Log-Structured Merge (LSM) or B trees, do not take full advantage of these new devices. Logic to avoid random accesses, expensive operations for keeping data sorted on disk, and synchronization

FASTER: A Concurrent Key-Value Store with In-Place Updates

Badrish Chandramouli[†], Guna Prasaad[‡], Donald Kossmann[†], Justin Levandoski[‡], James Hunter[†], Mike Barnett[†]

[†]Microsoft Research [‡]University of Washington

badrishe@microsoft.com, guna@cs.washington.edu, {donaldk, justin.levandoski, jahunter, mbarrett}@microsoft.com


ABSTRACT

Over the last decade, there has been a tremendous growth in data-intensive applications and services in the cloud. Data is created on a variety of edge sources, e.g., devices, browsers, and servers, and processed by cloud applications to gain insights or take decisions. Applications and services either work on collected data, or monitor and process data in real time. These applications are typically update intensive and involve a large amount of state beyond what can fit in main memory. However, they display significant temporal locality in their access pattern. This paper presents FASTER, a new key-

Real-world, large skewed workloads

- Point queries and **high** throughput *paramount*
- **Working sets** *larger* than main-memory – most data rarely accessed or updated
- Total **indexed** data order of magnitude *larger* than main-memory
- Natural **skew** in key access pattern – both for *reads* and *writes*
- Memory resources **scarce** – disk **wear** is a practical *concern*

Search Engine Workload




Insert / Update **statistics** (e.g., clicks, statistics)
Millions of **active** users at any time – *critical* path
Many more *inactive*, but we still need to keep data!
We fetch clicks for active users (during browsing)
We count viewed ads, from active users

Limitations of Existing Systems (1)

Log-structured Merge (LSM) Trees

- + Enable and tune (Bloom) *filters*
- + Use hash indices
- + Efficient compaction policies
- Filters may no longer fit in memory
- CPU overhead (10s of filters / query)
- Need tuning



Limitations of Existing Systems (2)

KVell

- Adjust page-cache size
- Large in-memory B-tree (19B per key)
- B-tree continuously paged out to disk

FASTER

- Tune record log in-memory size / hash index
- Reducing index size increases I/O ops
- Log compaction “pollutes” in-memory log


Limitations of Existing Systems (2)

KVell

- Adjust page-cache size
- Large in-memory B-tree (19B per key)
- B-tree continuously paged out to disk


FASTER

- Tune record log in-memory size / hash index
- Reducing index size increases I/O ops
- Log compaction “pollutes” in-memory log




Introducing F2

Key Idea: separate management of records across both read/write and hot/cold domains



Design Space



similar to TellStore-Log, but with **periodic compaction**

compaction aggressively optimized

F2 – Read Cache

Contains *read-hot* records of both **hot** log and **cold** log

- Hash index entries can point to *either* read cache entries or (**hot**) record log (single bit in address)


Reads go from **hot**-log index → (optionally) read cache → **hot** log

Reads from **cold**-log are *always* inserted at the tail of the read cache

Upserts and **RMWs** write *directly* to **hot** log tail, eliminating read cache chain



- If record with same key in read cache, it's **invalidated!**

Periodically, read cache is *evicting* in-memory records (HEAD), by altering the **hash chains**



F2 – Performance Comparison

YCSB: 250M keys, 8B keys, 100B values; **3GiB** mem budget (10% of dataset size); 24 threads, NVMe SSD




Zipfian: SplinterDB (**1.78x**), RocksDB (**4.61x**), FASTER (**4.94x**)


Hotspot 10%: SplinterDB (**1.66x**), RocksDB (**1.88x**), FASTER (**1.92x**)

F2 – I/O Amplification & Scalability Comparison

YCSB: 250M keys, 8B keys, 100B values; **3GiB** mem budget (10% of dataset size); 24 threads, NVMe SSD



F2 achieve less WA than Baselines



F2 saturates at 8-12 threads

→ Varying memory-budget, YCSB skewness; F2 detailed evaluation ←

CS 561: Data Systems Architectures

class 7

Design Tradeoffs in Key-Value Stores

Prof. Manos Athanassoulis

<https://bu-disc.github.io/CS561/>