



# Spanner

A horizontally scalable, highly-available SQL Database

**Ben Vandiver**

**3 April 2025**

**Google**


# What is Spanner?

# Spanner is a SQL Database



```
SELECT COUNT(*) AS Count  
FROM Artists AS a  
WHERE LEFT(a.name,1) = 'M';
```

# Spanner is Transactional & Consistent




**External Consistency**  
(stronger than linearizability and serializability)

# Spanner is Geographically Replicated



# Spanner is Horizontally Scalable

Proprietary + Confidential



# Zanzibar Availability on Spanner

Proprietary + Confidential



Availability over 3 years remained above 99.999% (“five 9’s”)

99.999% is ~5 minutes of downtime per year.



# How Big is Spanner?



# Spanner's Scale



**15 Exabytes**




**5 B QPS**

# What Makes Spanner so Big?



# What Makes Spanner so Big?




Google Cloud

Cloud Control Plane



Cloud Spanner



Zanzibar

google3

Google's internal infrastructure



# Wins!

2024




15 EiB


1976



60 MB? HDD



20? QPS



5B QPS

250,000,000,000x  
11 orders of (decimal) magnitude

250,000,000x  
8 orders of magnitude

```
CALL SEQUEL('UNDERPAID(NAME, SAL) ←
    SELECT  NAME, SAL
    FROM    EMP
    WHERE   JOB = 'PROGRAMMER'
    AND     SAL < 10,000');
```

# Programmers Not So Much

1976



&lt; \$10,000

```
CALL SEQUEL('UNDERPAID(NAME, SAL) ←  
  SELECT  NAME, SAL  
  FROM    EMP  
  WHERE   JOB = 'PROGRAMMER'  
  AND     SAL < 10,000');
```

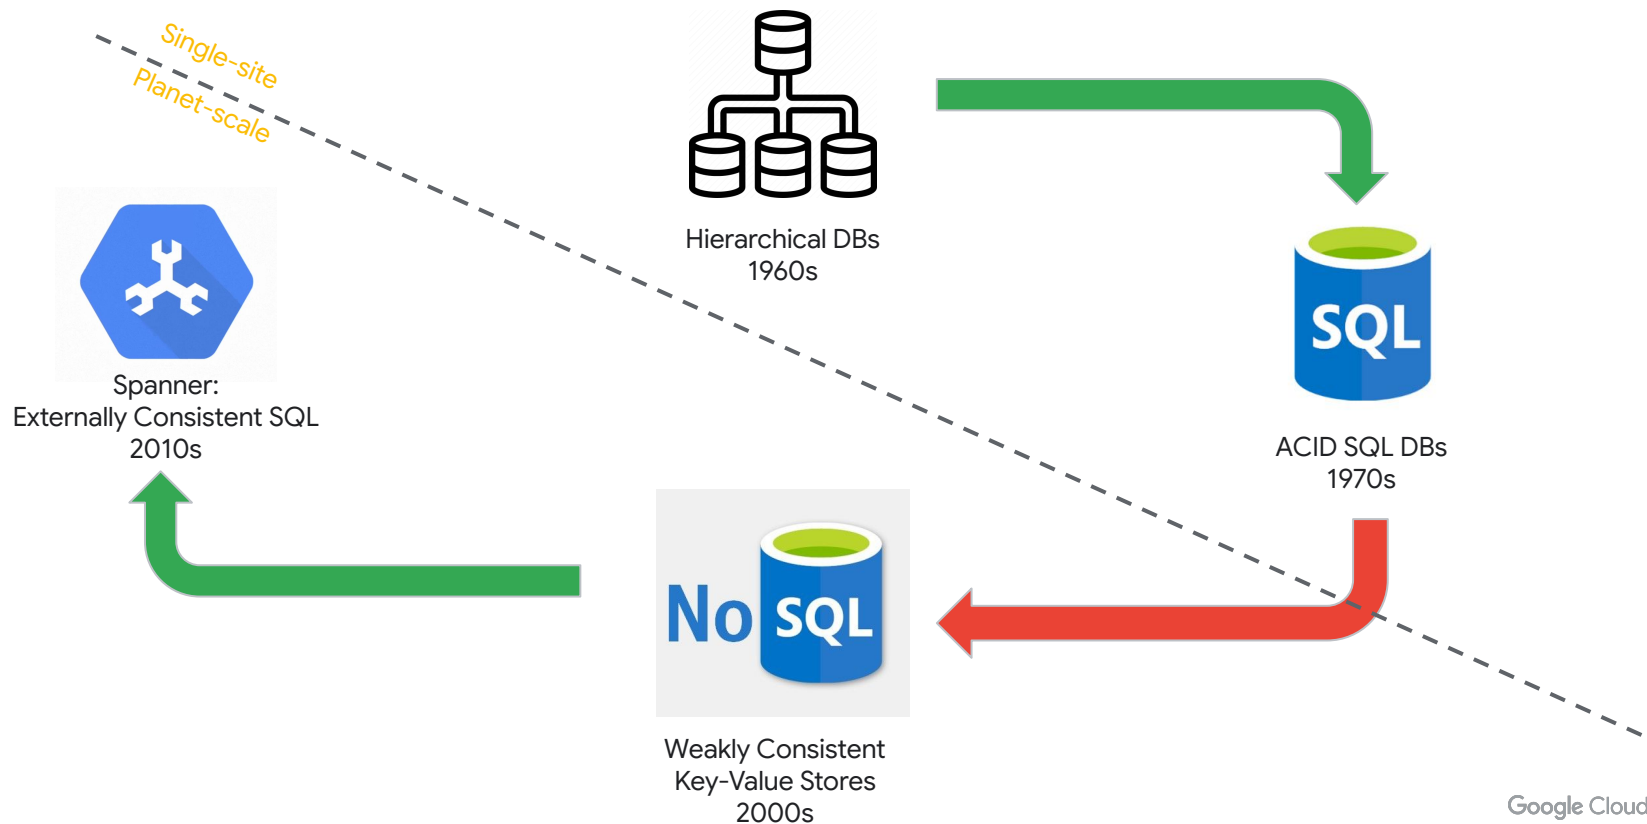
2024



&lt; \$100,000

1 order of magnitude

# Coming Full Cycle



# How Does Spanner Scale?


# SQL Example (not so unrealistic)

```
CREATE TABLE Songs (  
  ArtistId      INT64 NOT NULL,  
  AlbumId       INT64 NOT NULL,  
  TrackId       INT64 NOT NULL,  
  SongName      STRING(MAX),  
  ...  
) PRIMARY KEY (  
  ArtistId, AlbumId, TrackId)
```

ArtistId	AlbumId	TrackId	SongName
1	1	1	A Night in Tunisia
1	2	1	Exuberante
3	1	1	Straight, No Chaser
3	2	1	All Blues
3	2	2	Flamenco Sketches
3	2	3	So What
5	1	1	Gangnam Style



# Sharding




Machine 1

Machine 2

Machine 3

ArtistId	AlbumId	TrackId	SongName
1	1	1	A Night in Tunisia
1	2	1	Exuberante
3	1	1	Straight, No Chaser
3	2	1	All Blues
3	2	2	Flamenco Sketches
3	2	3	So What
5	1	1	Gangnam Style

# Splitting (and Merging)




ArtistId	AlbumId	TrackId	SongName
1	1	1	A Night in Tunisia
1	2	1	Exuberante
3	1	1	Straight, No Chaser
3	2	1	All Blues
3	2	2	Flamenco Sketches
3	2	3	So What
5	1	1	Gangnam Style

# Interleaving

```
CREATE TABLE Albums (  
  ArtistId      INT64 NOT NULL,  
  AlbumId       INT64 NOT NULL,  
  AlbumName     STRING(MAX),  
  ...  
) PRIMARY KEY (  
  ArtistId, AlbumId);
```

```
CREATE TABLE Songs (  
  ArtistId      INT64 NOT NULL,  
  AlbumId       INT64 NOT NULL,  
  TrackId       INT64 NOT NULL,  
  SongName      STRING(MAX),  
  ...  
) PRIMARY KEY (  
  ArtistId, AlbumId, TrackId),  
INTERLEAVE IN PARENT Albums;
```

# Interleaving



Machine 1

Machine 2

Machine 3

ArtistId	AlbumId	AlbumName	TrackId	SongName
1	1	Best of Dizzy Gillespie		
1	1		1	A Night in Tunisia
1	2	Afro-Cuban Moods		
1	2		1	Exuberante
3	1	Milestones		
3	1		1	Straight, No Chaser
3	2	All Blues		
3	2		1	All Blues
3	2		2	Flamenco Sketches
3	2		3	So What
5	1	Psy 6 (Six Rules)		
5	1		1	Gangnam Style


# Transactions & High Availability

# Transactions: Single-Shard

Well-understood:

- Transaction Manager with pessimistic locking
- Strict 2-Phase Locking
- Write-Ahead Logging to persistent storage (disk)

But what about machine failures? Want copies of the data on multiple machines (ideally in different data centers).



Machine 1

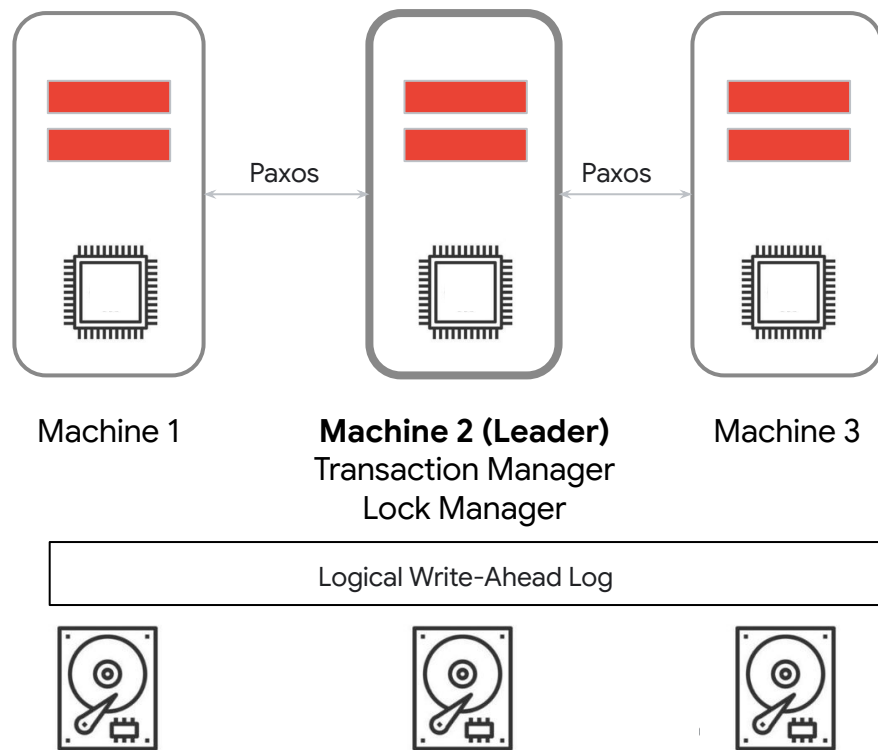


Write-Ahead Log

# Transactions: Paxos

Answer: Paxos (or Raft)

- Synchronous replication protocol
- Gracefully deals with failures
  - Makes progress as long as majority of replicas are alive
- Elects a leader to manage writes
  - Transaction Manager and Lock Manager run here



# Paxos in a Nutshell

Manages a shared log of writes. Every replica applies log entries in order.

Consists of  $N = 2f + 1$  nodes


- Can tolerate  $f$  failures

Examples:

- $N=3$  nodes: tolerate 1 failed node
- $N=5$  nodes: tolerate 2 failed nodes

Quorum: Majority of nodes ( $f + 1$ )

Observation: Intersection between any two quorums is at least one node.





# Distributed Transactions

What about transactions that span multiple shards?


Answer: 2-Phase Commit (2PC)

- Participants are paxos groups (not individual machines)
- One participant picked as Coordinator
- All participants (durably) prepare and notify coordinator
  - Promise not to release locks
- Coordinator commits transaction once all participants have prepared, then notifies participants

Biggest down-side of 2PC:

- Gets stuck if coordinator becomes unavailable

We just made every participants (and coordinator) highly available through Paxos!



# CAP Theorem

## C - Consistent

- All nodes have identical data, strong reads

## A - Available


- Every request gets a response

## P - Partition Tolerant

- Continue to function (correctly) even if a network partition occurs.

Pick 2!

Spanner is a **CP** system.





# SQL interface

# Common SQL dialect

- Standards-compliant
- Type system aligned with programming languages
  - INT64, FLOAT, STRING (UTF8), TIMESTAMP (nanoseconds)
  - Reduces impedance mismatch
- First-class support for nested data
  - ARRAY and STRUCT types
  - Protocol Buffers: schematized binary objects

Significant language design work across teams
- Shared with other Google systems: BigQuery, F1, etc.

# Sample query: name & titles



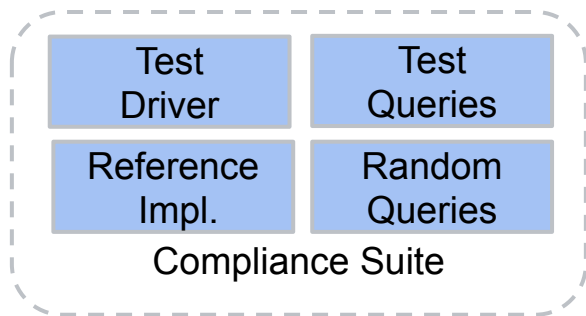

```
SELECT s.SingerName,  
       ARRAY(SELECT a.AlbumTitle  
             FROM Albums a  
             WHERE a.SingerId = s.SingerId) titles  
FROM Singers s  
WHERE s.SingerId BETWEEN 1 AND 5
```

- Easier to use than outer joins or multiple roundtrips

SingerName STRING	titles ARRAY<STRING>
Beatles	[Help!, Abbey Road]
U2	[ ]
Pink Floyd	[The Wall]

# Same query semantics across systems

Input	
Shared	
Engine	





# Distributed Execution

# Distributed query execution

- Tightly coupled architecture
  - Query processor inside the database server
  - Typical design for standalone DBMSes (vs. distributed systems)
- Challenge of scale: data never sits still
  - Continuous resharding (due to load, capacity, config changes, ...)
  - Shard boundaries may change while query is running
  - Shards may become temporarily unavailable during query execution
  - Alternative replicas: near/far, loaded/idle, caught-up/behind
- Mechanisms used in Spanner
  - Query routing: key-range rpcs + range extraction
  - Parallelizing execution: partition work by shards, push it down
  - Dealing with failures: restartable query processing



# Query routing: key-range rpcs

- Routes requests to row ranges
  - E.g., WHERE SingerId BETWEEN @low AND @high
- Hides complexity of locating data
- Finds nearest, sufficiently up-to-date replica for given concurrency mode
- Retries automatically
  - Unavailability, data movement, schema updates, ...
- Clients cache sharding information
- Clients cache "location hints" for queries
  - Send query to right server without extra hops or query analysis
  - E.g., Singers/SingerId[@low]

# Query routing: range extraction

```
SELECT * FROM Albums
WHERE (SingerId = 1 AND AlbumId >= 10) OR
      (SingerId IN (2,3) AND AlbumId != 0)
```

- Also used for restricting scan ranges
- Computed at runtime
  - May access data
- Uses efficient data structure
  - Filter tree (in the paper)


SingerId	AlbumId
[1..1]	[10, +INF)
[2..2]	(-INF, 0)
[2..2]	(0, +INF)
[3..3]	(-INF, 0)
[3..3]	(0, +INF)



# Parallelizing Execution

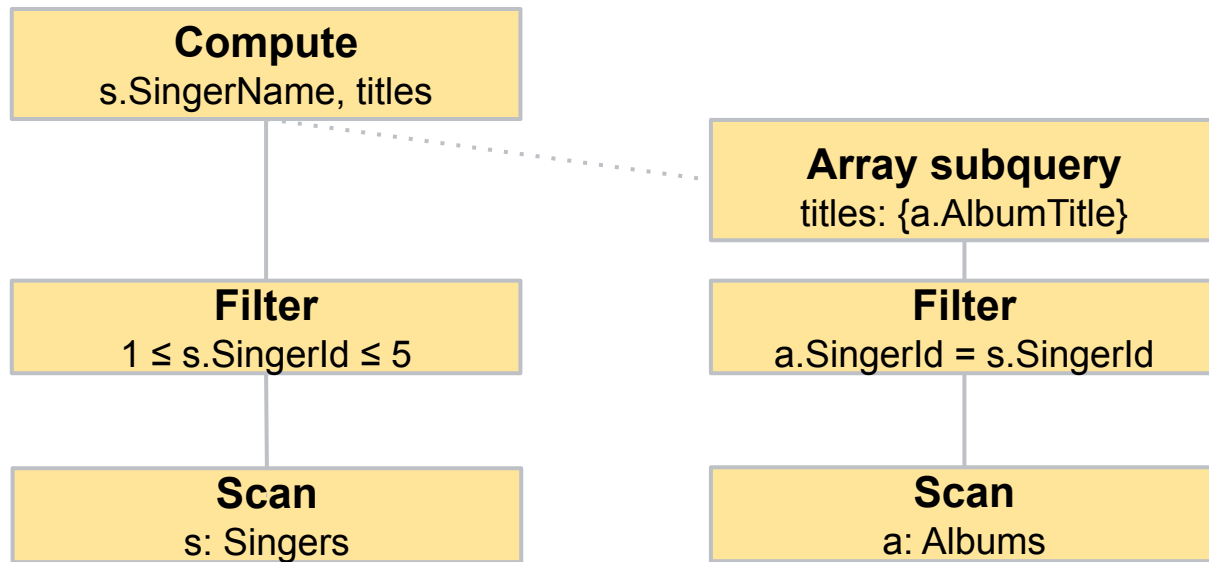
# Parallelizing execution

```
SELECT SingerName, ARRAY(SELECT ...) titles  
FROM Singers WHERE SingerId BETWEEN 1 AND 5
```

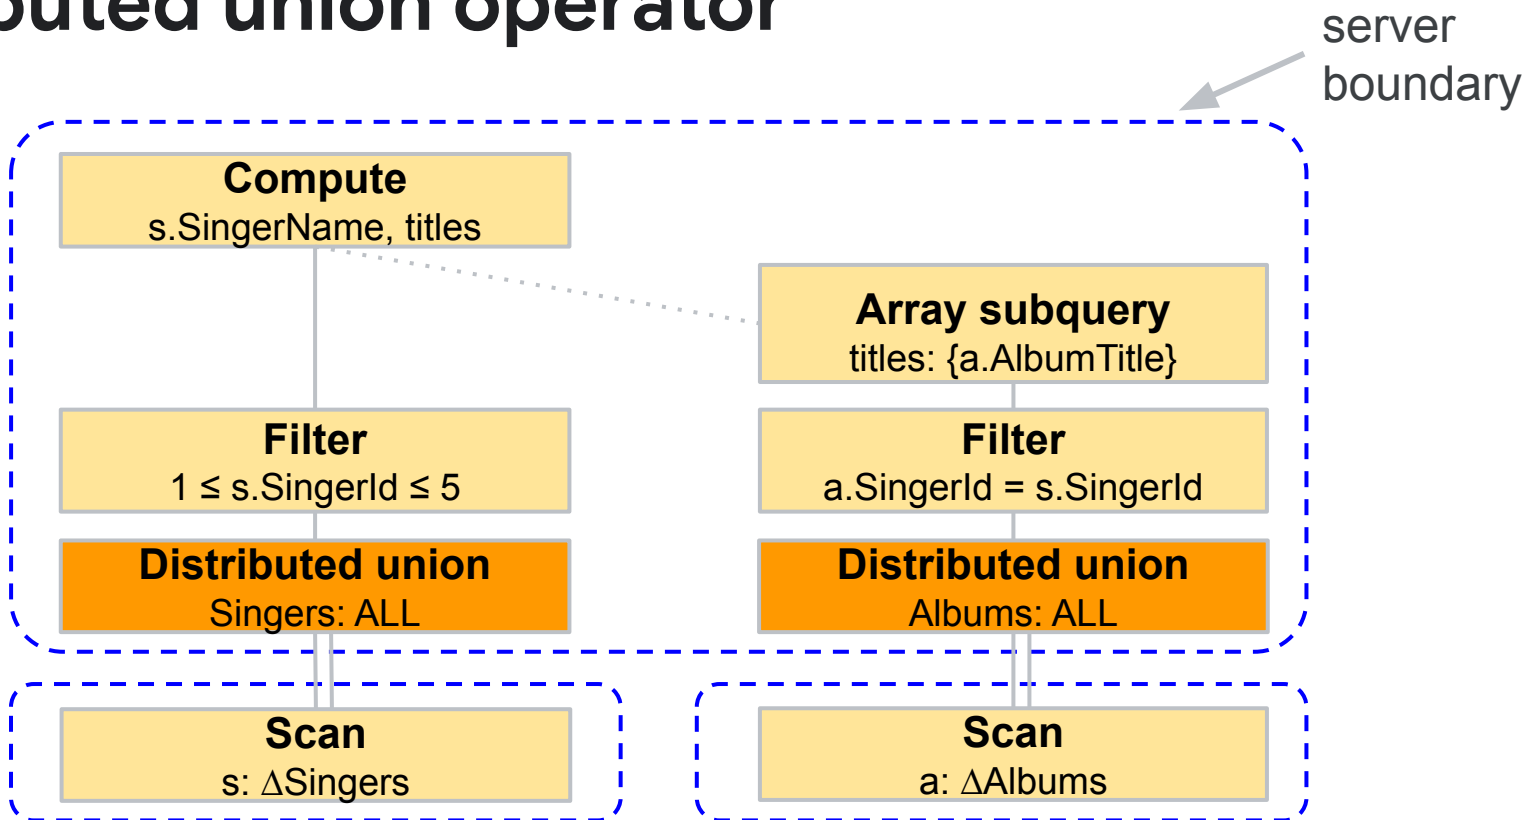


- Assume fixed shard boundaries for now

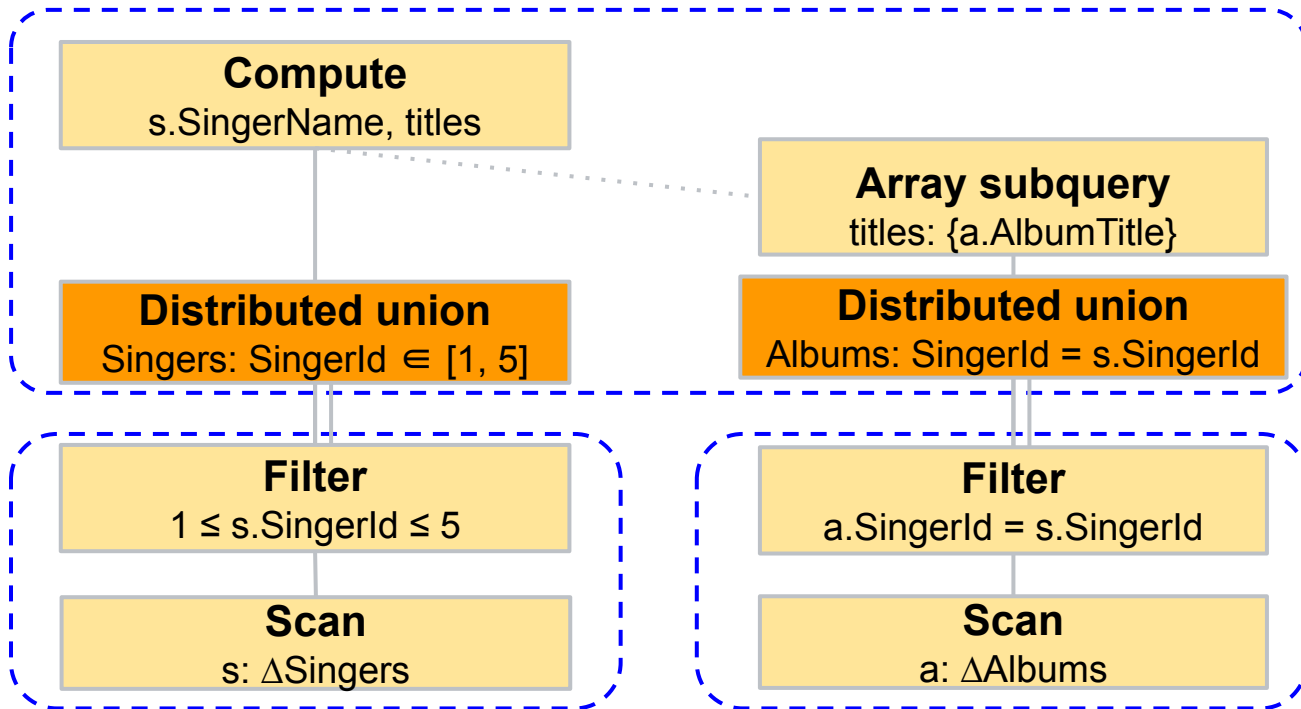
# Initial logical plan



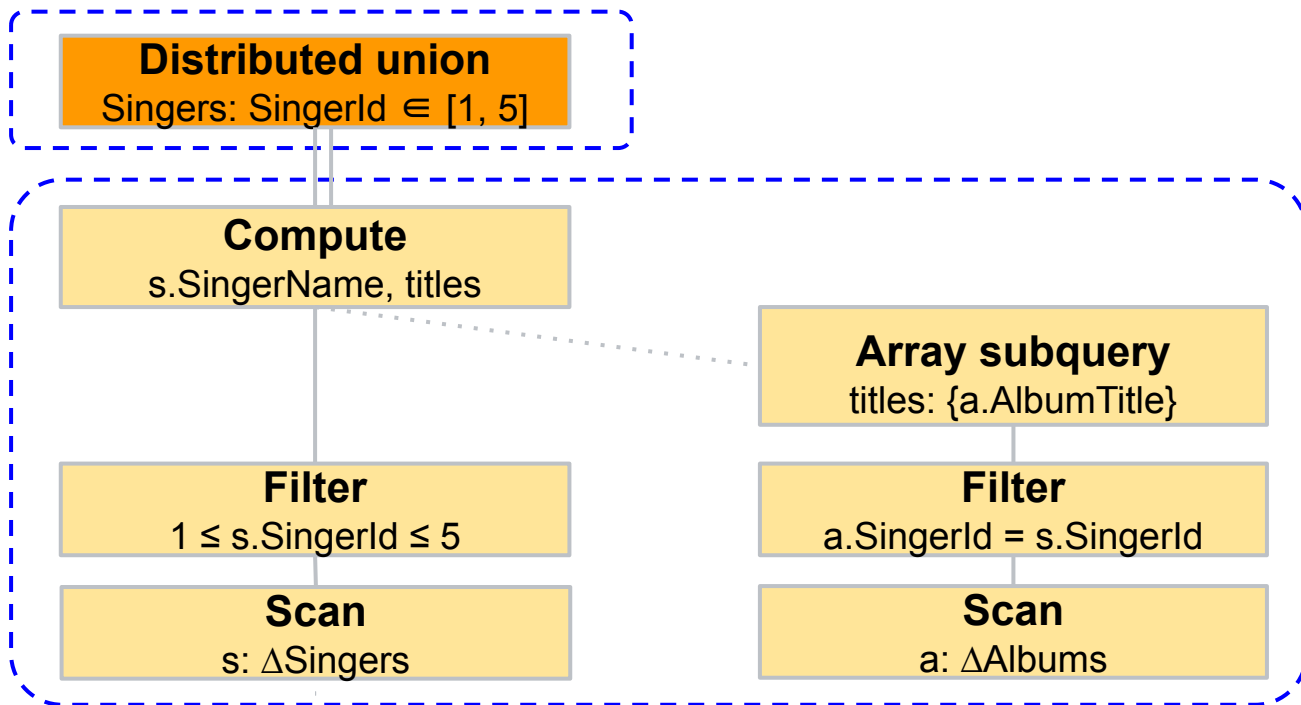
# Distributed union operator



# Push work to shards, extract distribution ranges



# Exploiting co-location






# Parallel-consumer API

```
SELECT SingerName, ARRAY(SELECT ...) titles
FROM Singers WHERE SingerId BETWEEN 1 AND 5
```

- Root-partitionable query:  
 $Q(\text{Union of } \Delta T) = \text{Union of } Q(\Delta T)$
- Same result up to order of rows
- Another main distribution operator: Distributed Cross Apply





# Restartable snapshot queries

## Query restarts: overview

- Automatic compensation for failures
- For snapshot queries only
- Server yields "restart token" with each result batch
- Client resumes query execution after consuming partial results
- Contract: omit previously returned rows
  - No repeatability guarantee for subsequent rows

<b>SingerName</b> STRING	<b>titles</b> ARRAY<STRING>
Beatles	[Help!, Abbey Road]
U2	[ ]
Pink Floyd	[The Wall]

restart



# Query restarts: implementation challenges

- Naive solutions don't work well for "large" queries
  - Buffer final result, persist intermediate results, count rows, etc.
- Instead: efficiently capture distributed state of query execution
- Dynamic resharding
  - May restart on different row range
- Non-determinism
  - Memory size, parallelism, computer architecture, numerics, ...
- Restarts across server versions
  - Query plans, execution algorithms


# Query restarts: hard but worth it

- Hide transient failures
- No retry loops: simpler programming model
- Streaming pagination
- Ensure forward progress for important class of long-running queries
- Improve tail latency of online requests
- Low-impact rolling server upgrades



# Building Database Systems

# Location



# Database System Contract

**Goal: Absorb complexity from application designers**

- **Failures - Replication, Transactions**
- **Performance - Efficient data layout, Query Optimization**
- **Scalability - Sharding, Routing**
- **Security & Compliance - Regulatory, Fine-grained ACLs**



# What are we building?

- Operating System + Compiler + Distributed System
- Interface surface area is immense
- SQL is a standard - OK, not really.

## Hyrum's Law

With a sufficient number of users of an API,  
it does not matter what you promise in the contract:  
all observable behaviors of your system  
will be depended on by somebody



EVERY CHANGE BREAKS SOMEONE'S WORKFLOW.

# Complexity

- Respect Complexity
  - No one person fully understands Spanner. Not even close.
- Spend "complexity budget" where it matters
- Databases are heavily interconnected systems
  - Hard to tightly componentize - "leaky abstractions"
  - Everything breaks Backup/Restore

# Tests, Tests, Tests

- Randomized Testing

- **Random Query, Data, and Schema generators**

- Generate syntactically and semantically valid queries and data  
(eg queries actually return rows)

- **Inject Faults**

- Anything that can fail or change should

- **Validate with simpler emulator**

- Check the answers against a much simpler system with the same semantics

For more: "Randomized Testing of Cloud Spanner" Jay Corbett (Medium post)

# Tests, Tests, Tests

- Integrity Checks

- **Validate structures against each other internally**

Index vs table, generated columns vs expressions, check constraints, foreign keys, etc.

- **Use disaggregated compute to avoid workload impact**

- Log like crazy

- It may take a while to notice an issue

# Plan Stability: Avoiding Regressions

- *"We have millions of query plans in production, how do we update the optimizer safely?"*
- Pin Plans? Pin Optimizer version? Pin feature flags?
- Disaggregation helps:
  - pin plans
  - async validate nothing regresses
  - unpin as you validate

**Questions?**