

CS 561: Data Systems Architectures

class 22

Machine Learning & Data Systems

Prof. Manos Athanassoulis


<https://bu-disc.github.io/CS561/>

Machine learning algorithms improve *automatically* through *experience* and by the use of ***data***.

Machine learning algorithms build a model based on ***training data***, in order to make ***predictions*** or ***decisions*** without being explicitly programmed to do so.




Which database systems components can benefit/be replaced by ML algorithms?






*application/SQL
access patterns
complex queries*



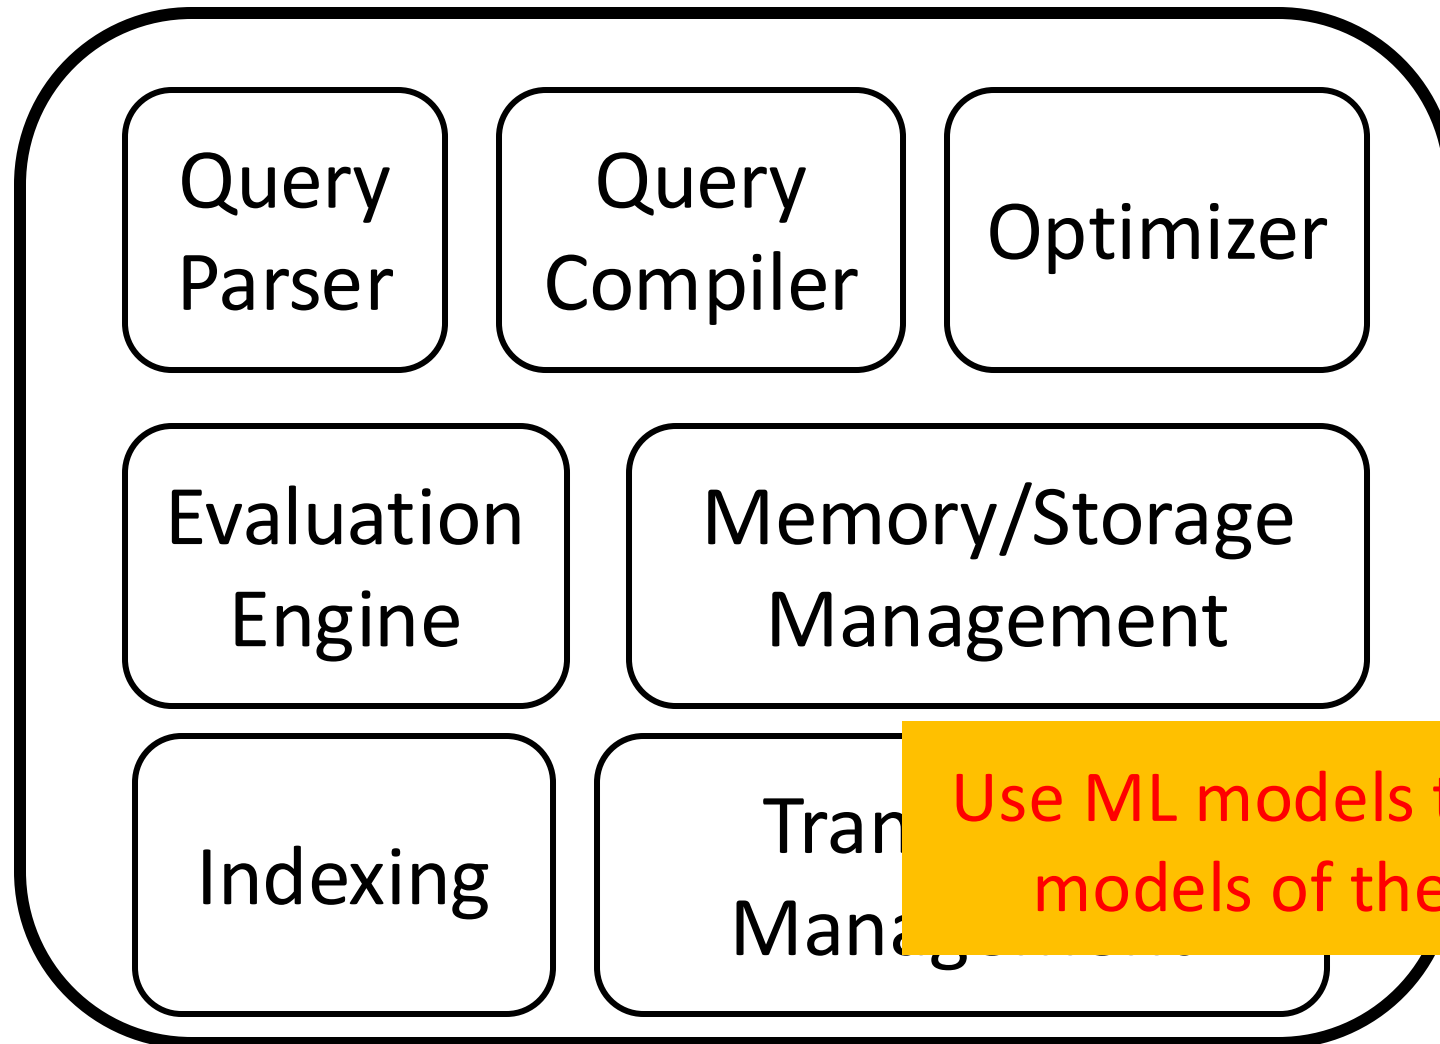
modules






*application/SQL
access patterns
complex queries*

modules



Use ML models to replace the cost-models of the database ***Tuner***



*application/SQL
access patterns
complex queries*

modules

Query
Parser

Query
Compiler

Optimizer


Memory/Storage
Management

Transaction
Management

Indexing

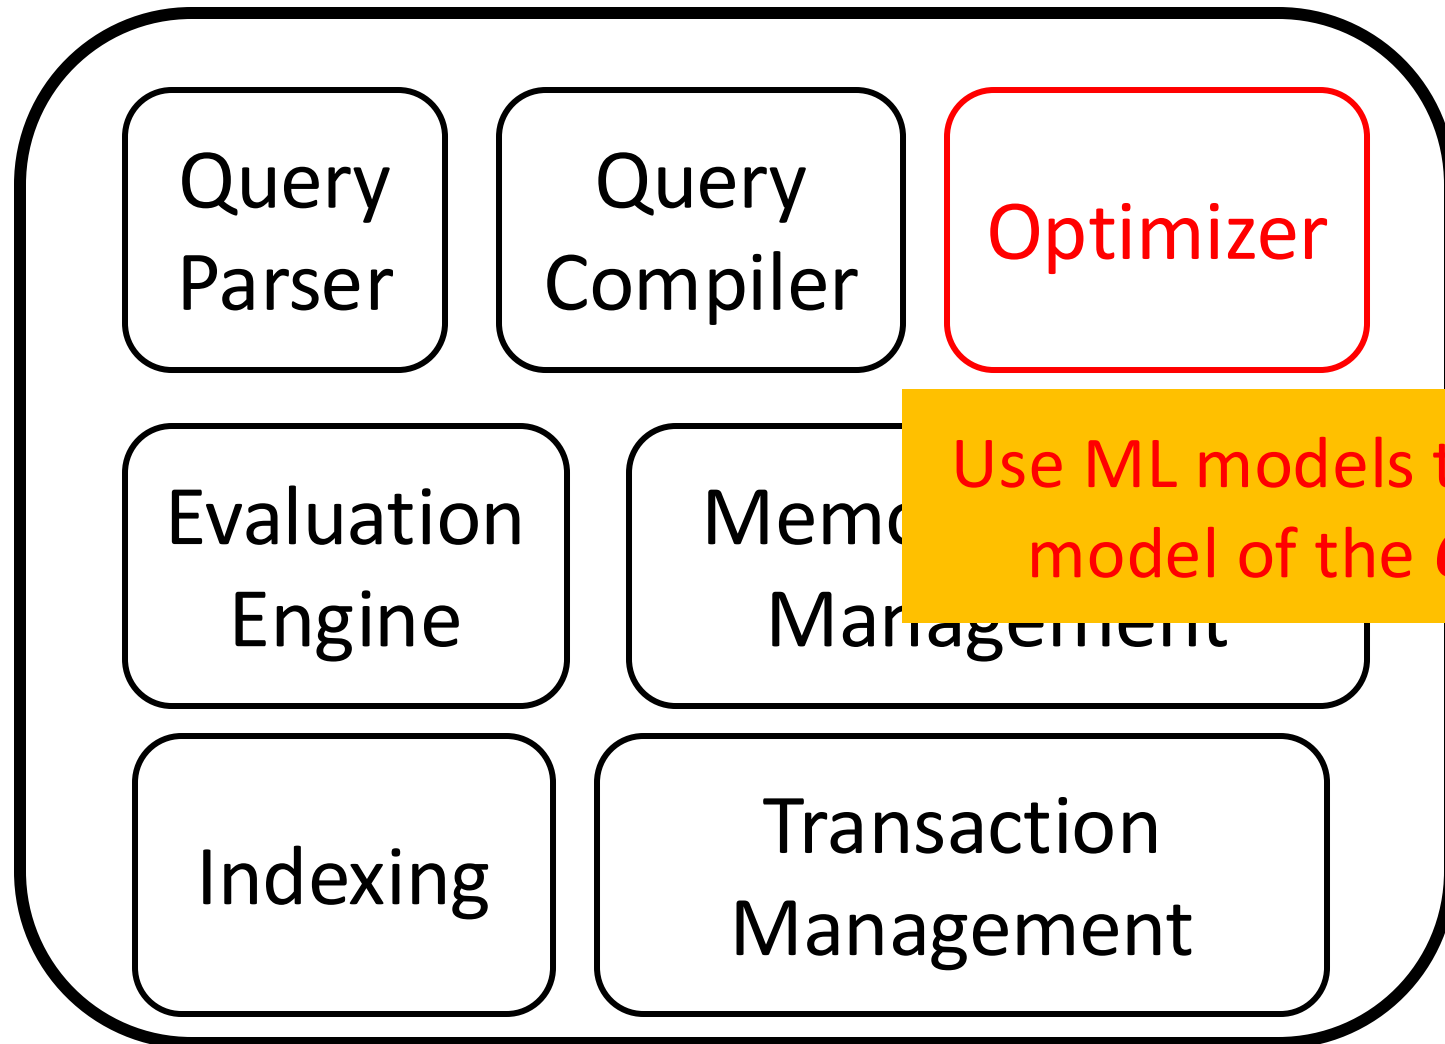
Tuner 

Use ML models to replace the
*navigational part of an **Index***




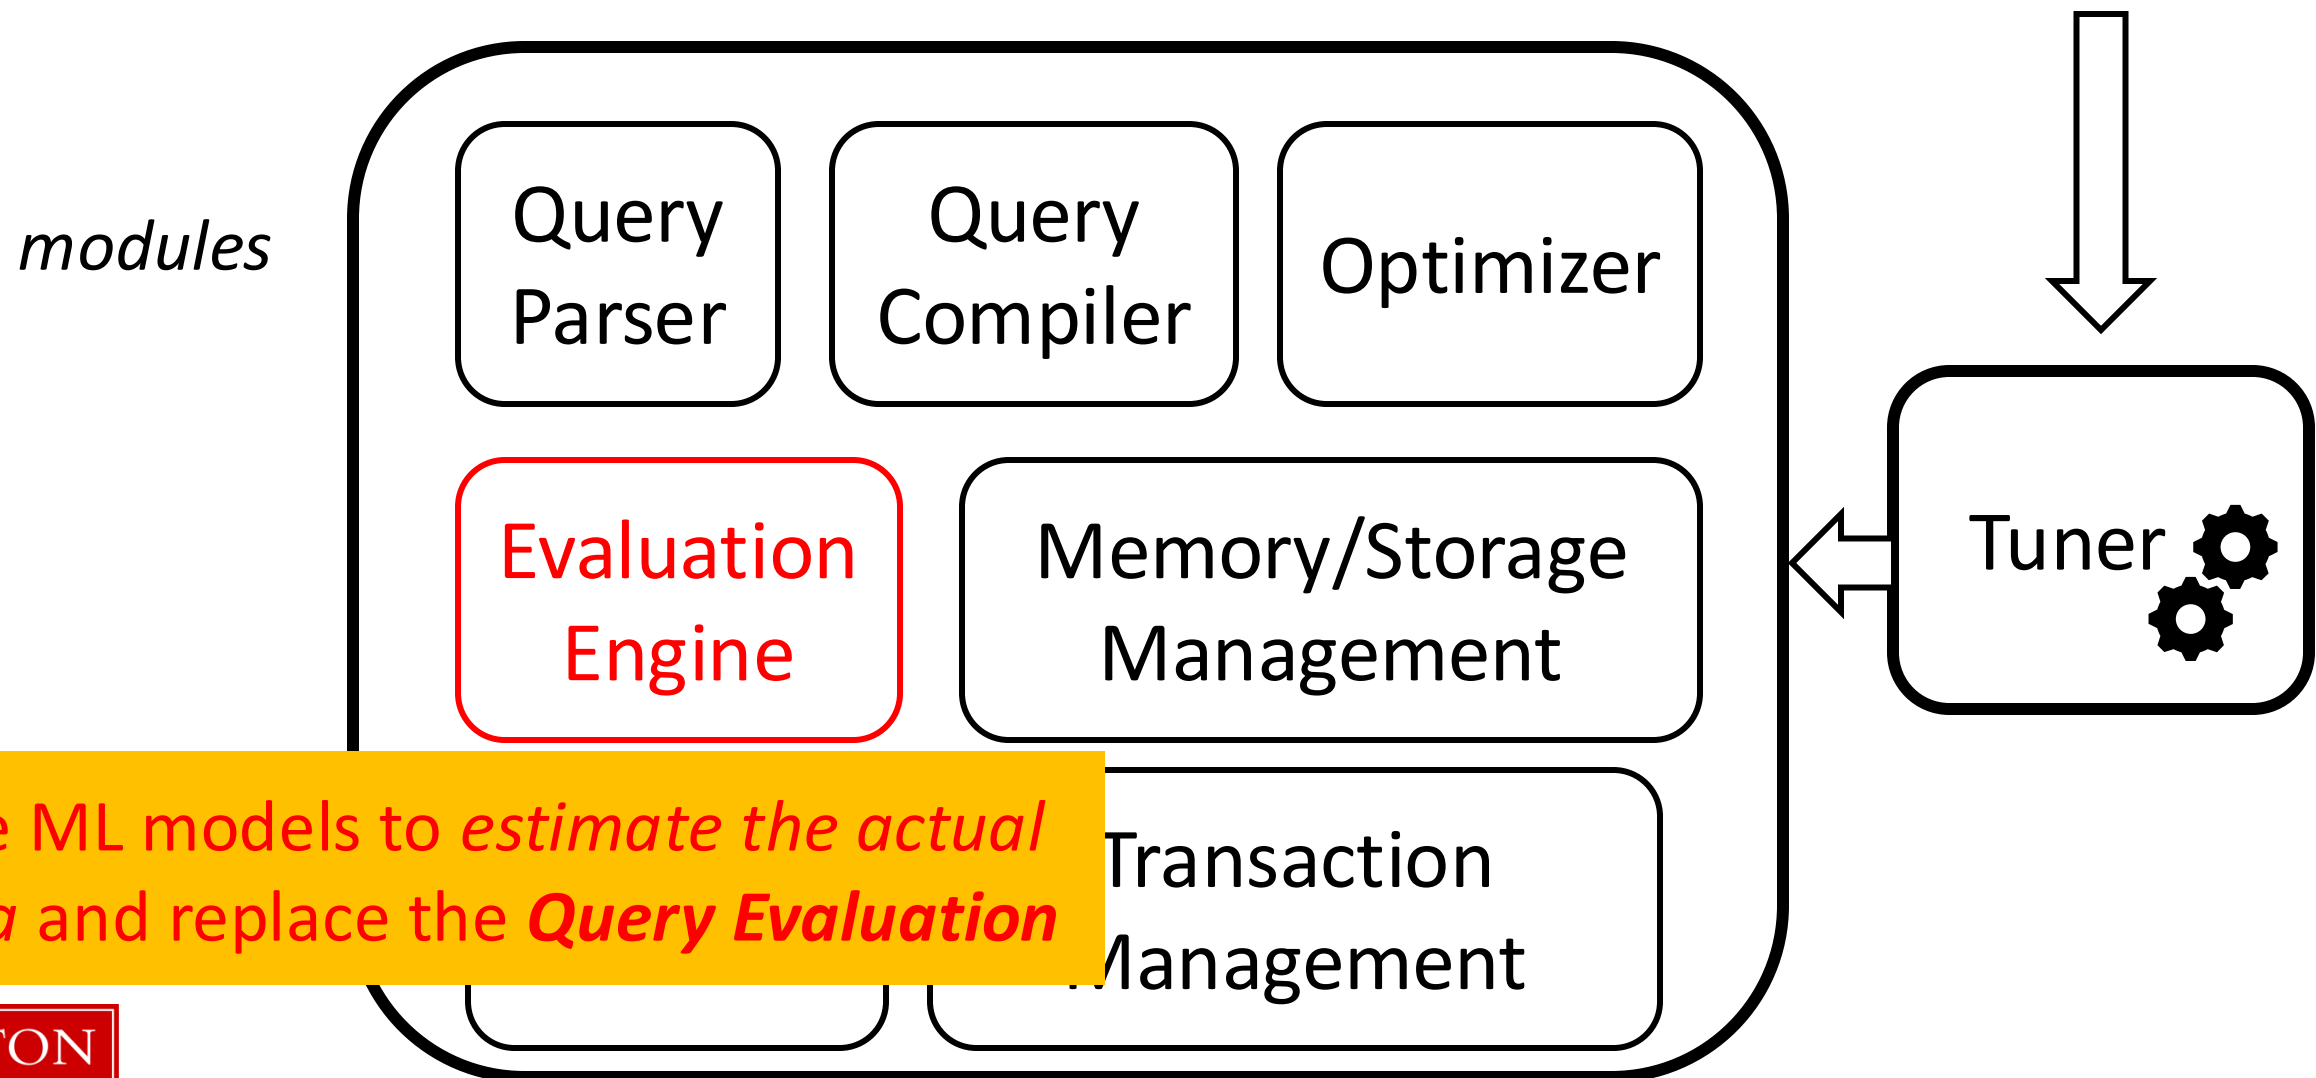
*application/SQL
access patterns
complex queries*

modules




Use ML models to replace the cost-model of the **Query Optimizer**






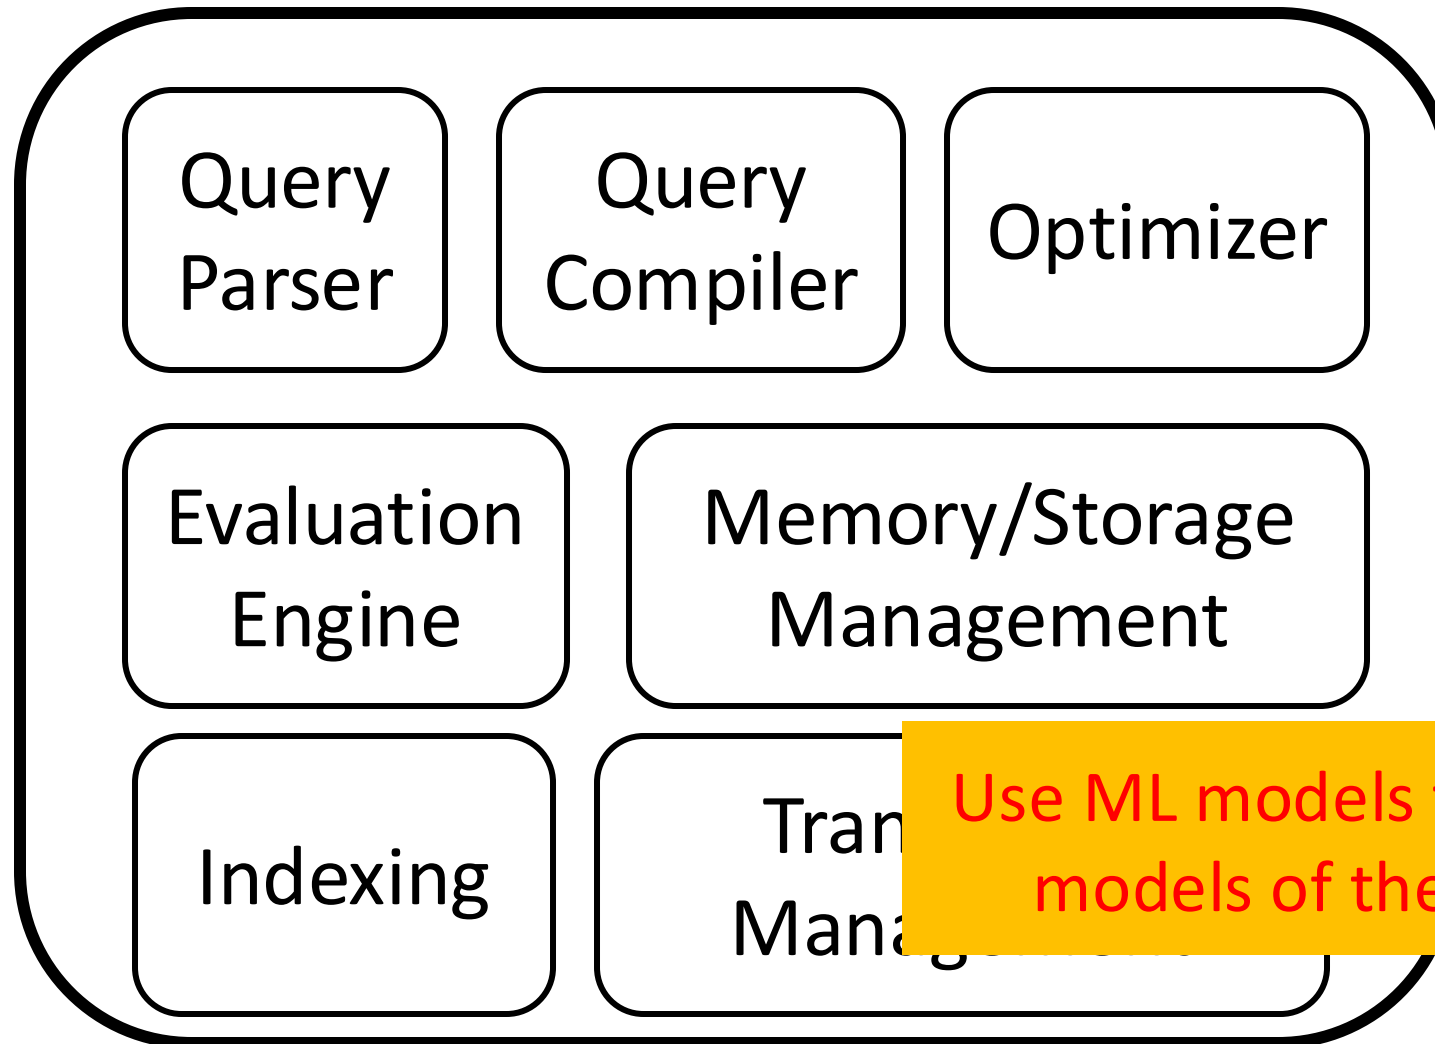


Use ML models to *estimate the actual data* and replace the **Query Evaluation**



*application/SQL
access patterns
complex queries*

modules



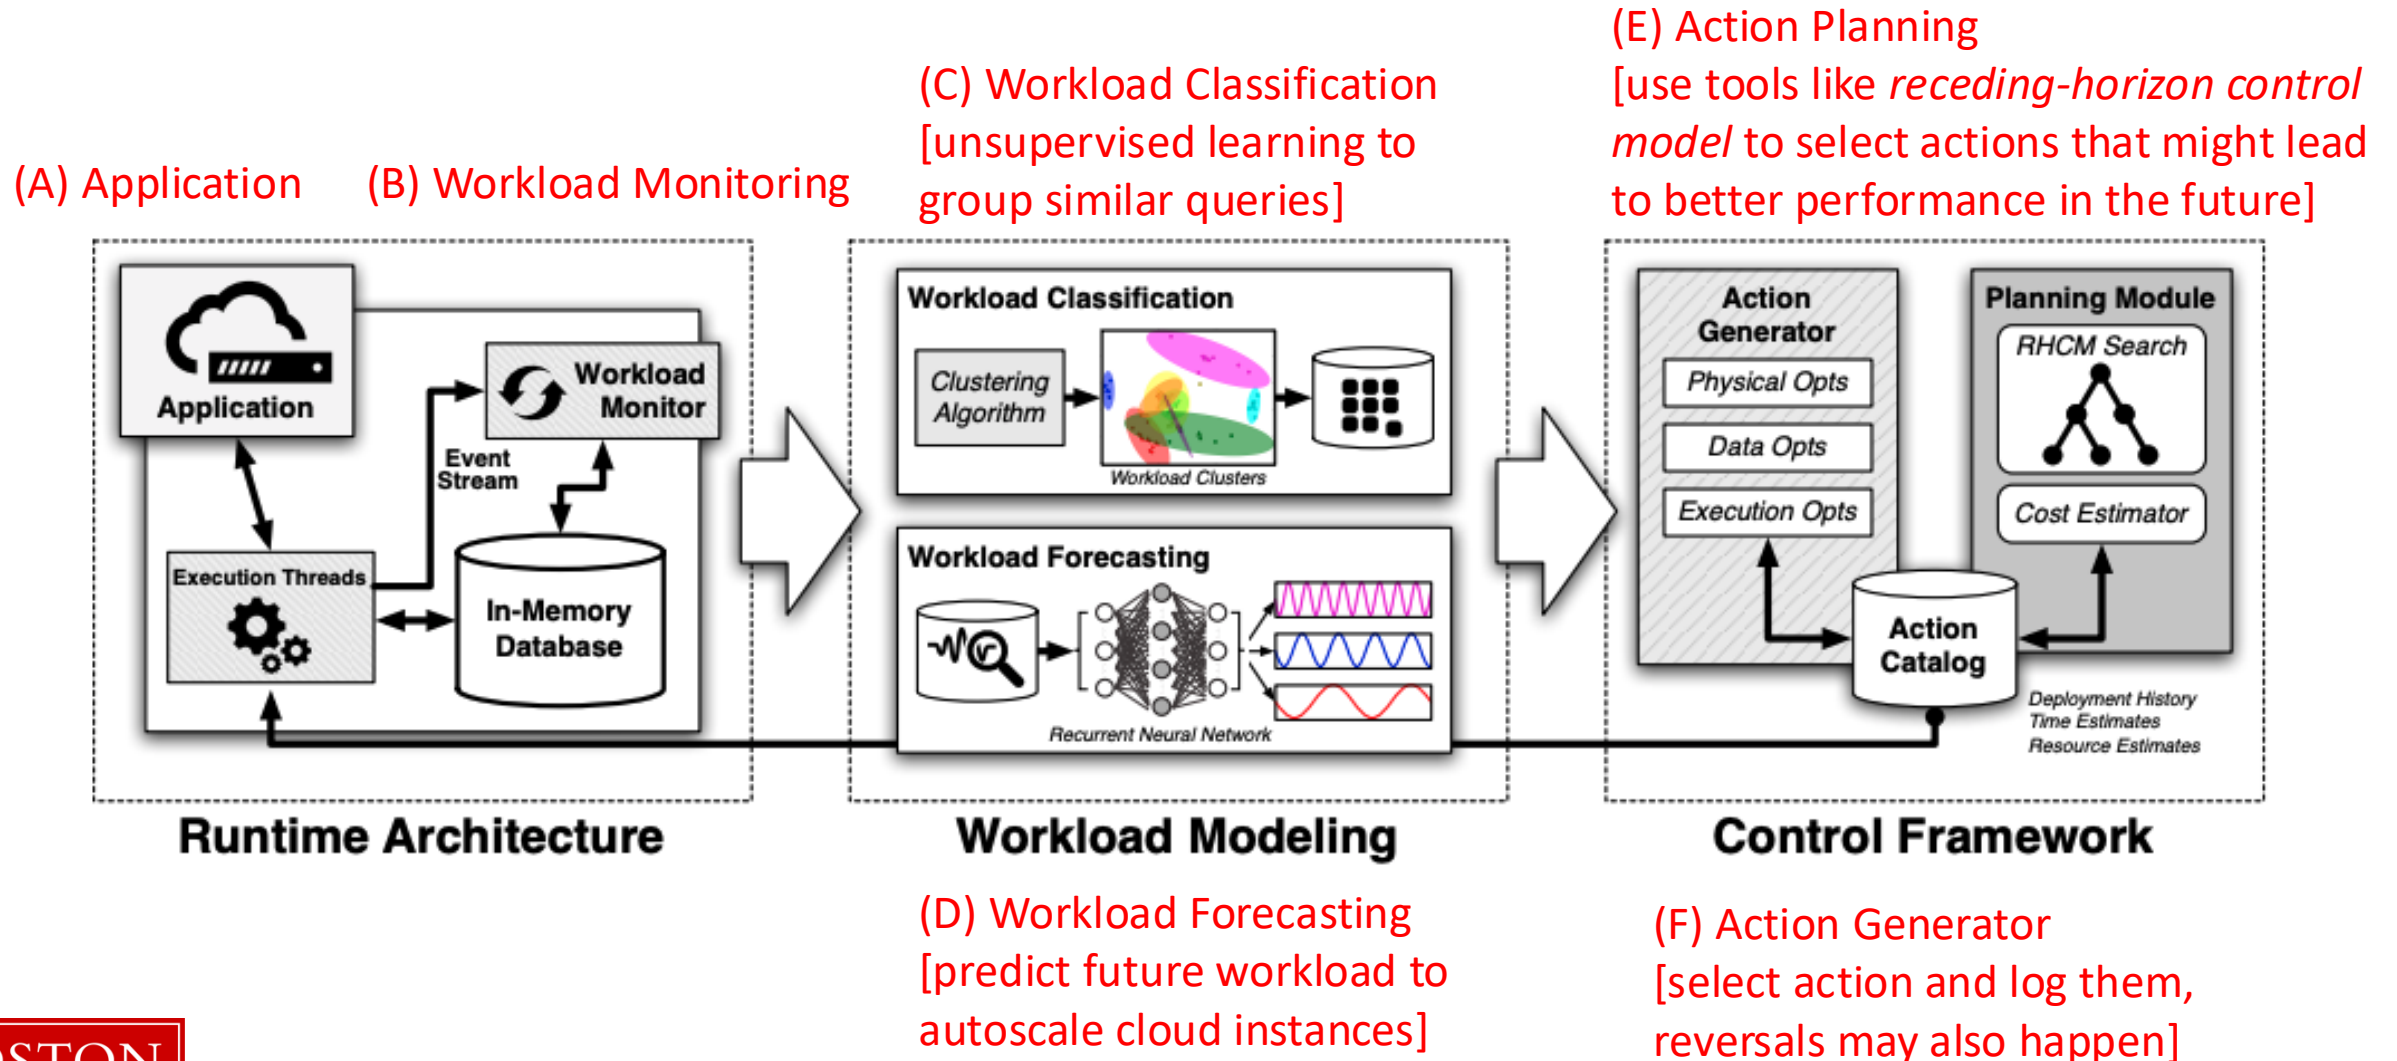
Use ML models to replace the cost-models of the database ***Tuner***

Self-driving Data systems

Types of actions that
a self-driving system
needs to take
automatically


	Types	Actions
PHYSICAL	Indexes	AddIndex, DropIndex, Rebuild, Convert
	Materialized Views	AddMatView, DropMatView
	Storage Layout	Row→Columnar, Columnar→Row, Compress
DATA	Location	MoveUpTier, MoveDownTier, Migrate
	Partitioning	RepartitionTable, ReplicateTable
RUNTIME	Resources	AddNode, RemoveNode
	Configuration Tuning	IncrementKnob, DecrementKnob, SetKnob
	Query Optimizations	CostModelTune, Compilation, Prefetch

Use-case: Peloton Self-Driving Architecture




Workload forecasting

Using Recurrent
Neural Networks (RNN)
the model learns patterns
and adapts to changes




(a) RNN Forecast Model (24-Hour Horizon)



(b) RNN Forecast Model (7-Day Horizon)

Action example: adapting the storage layout




Columns are better for OLAP

Rows are better for OLTP


Hybrid matches the best when workload alternates

Why automatic tuning is hard? (1/2)



(a) Dependencies


Complex interdependencies between different tuning knobs!




(b) Continuous Settings

Continuous domain ("too many" knob options) with irregular benefits

Why automatic tuning is hard? (2/2)

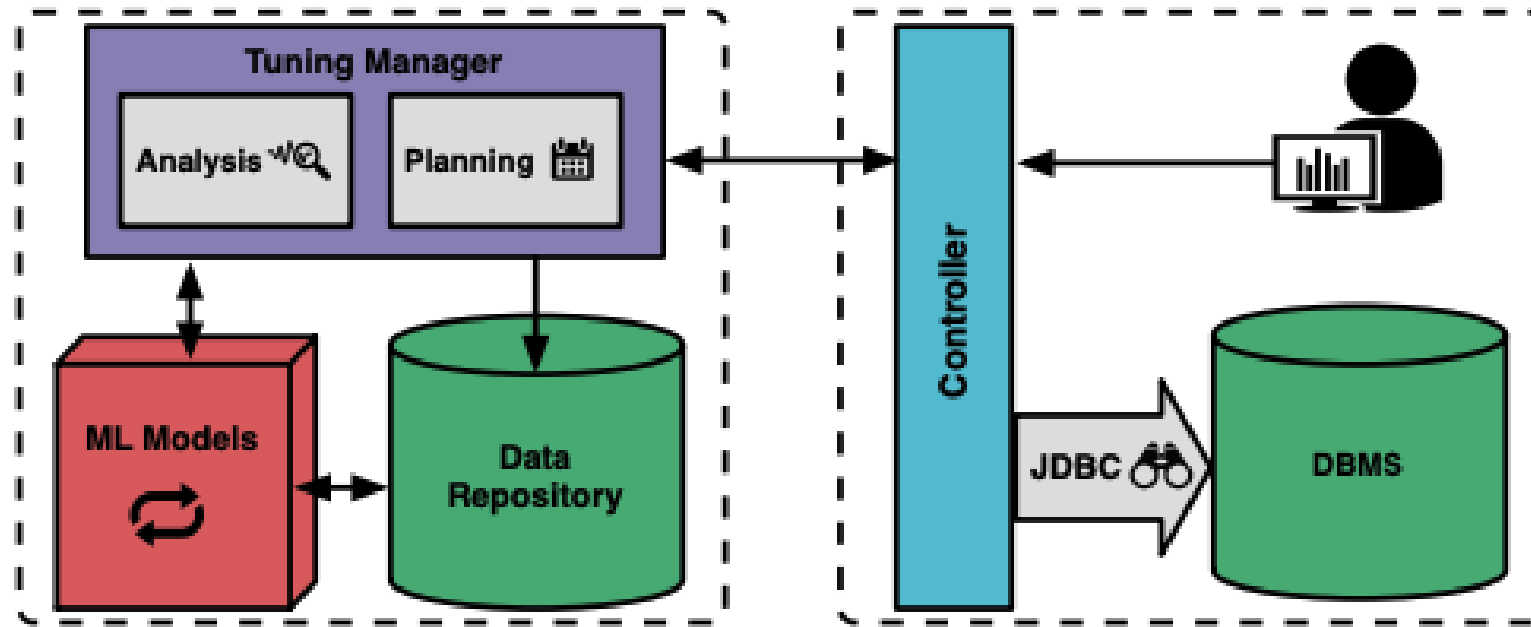


Non-reusable configurations!



Increasing tuning complexity

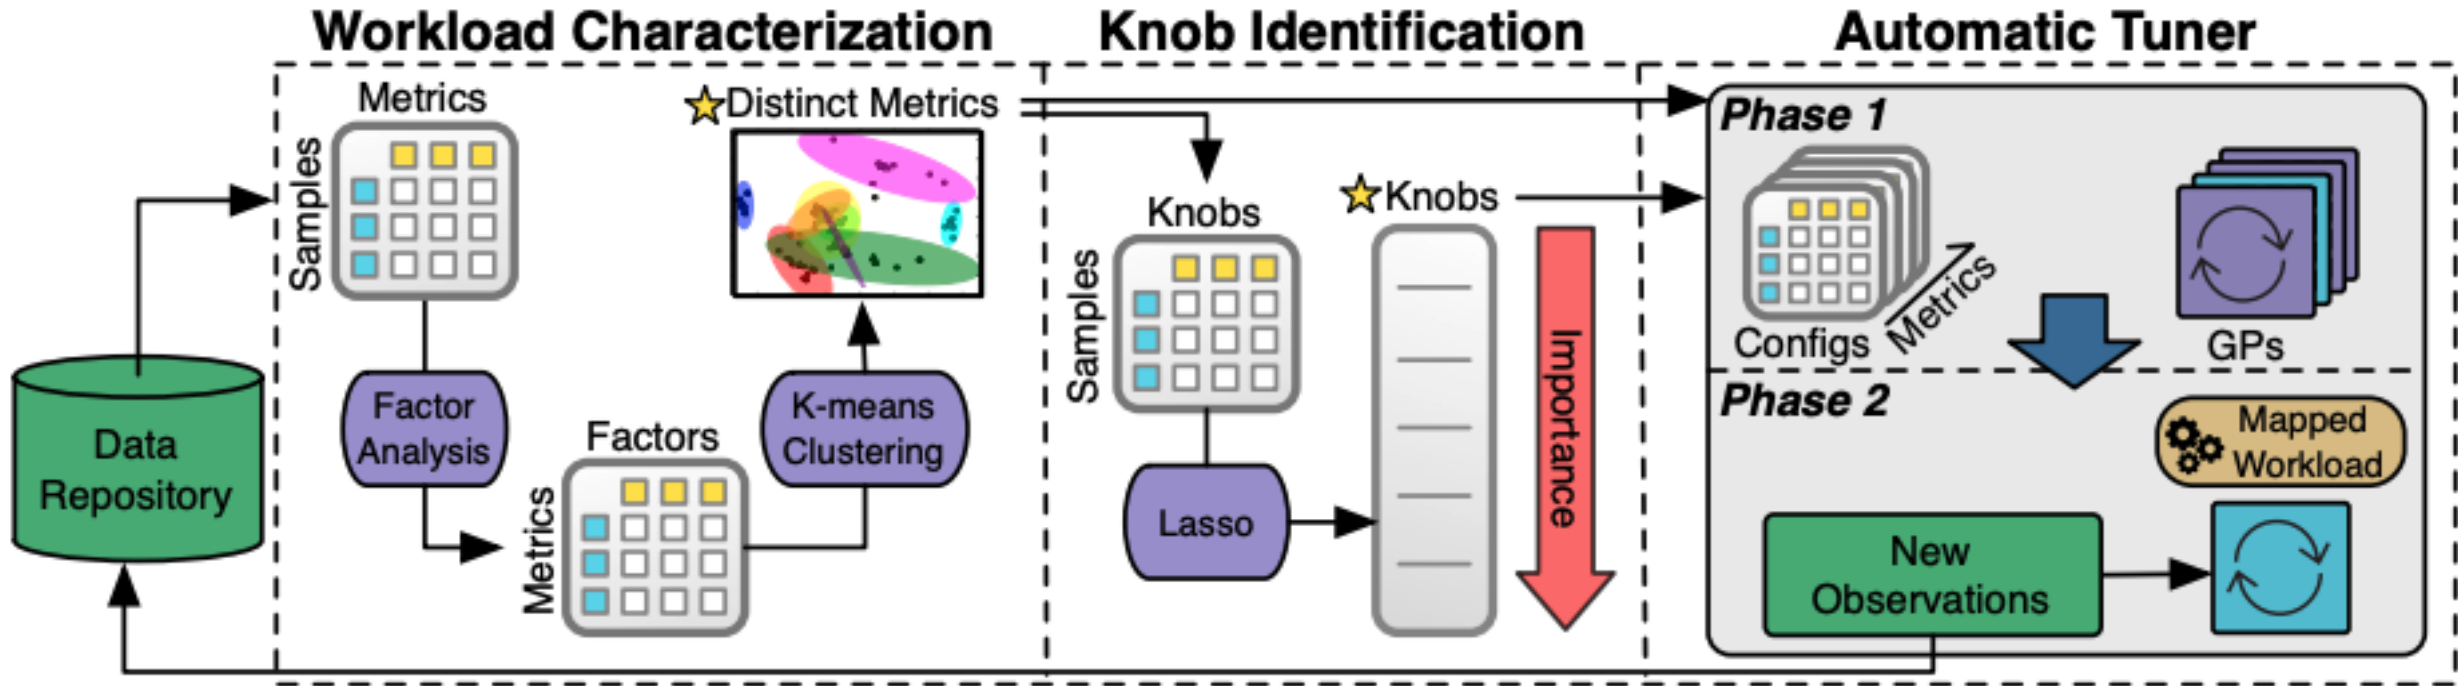
Use case: Ottertune



Two distinct components: the tuning manager **does not have access to data**, only to **performance metrics** and the values of the **tuning knobs**

All performance data are organized per system and per major version to ensure that no wrong, deprecated, or non-existing knobs are tuned.

OtterTune Machine Learning Pipeline



How to classify/characterize a workload? .

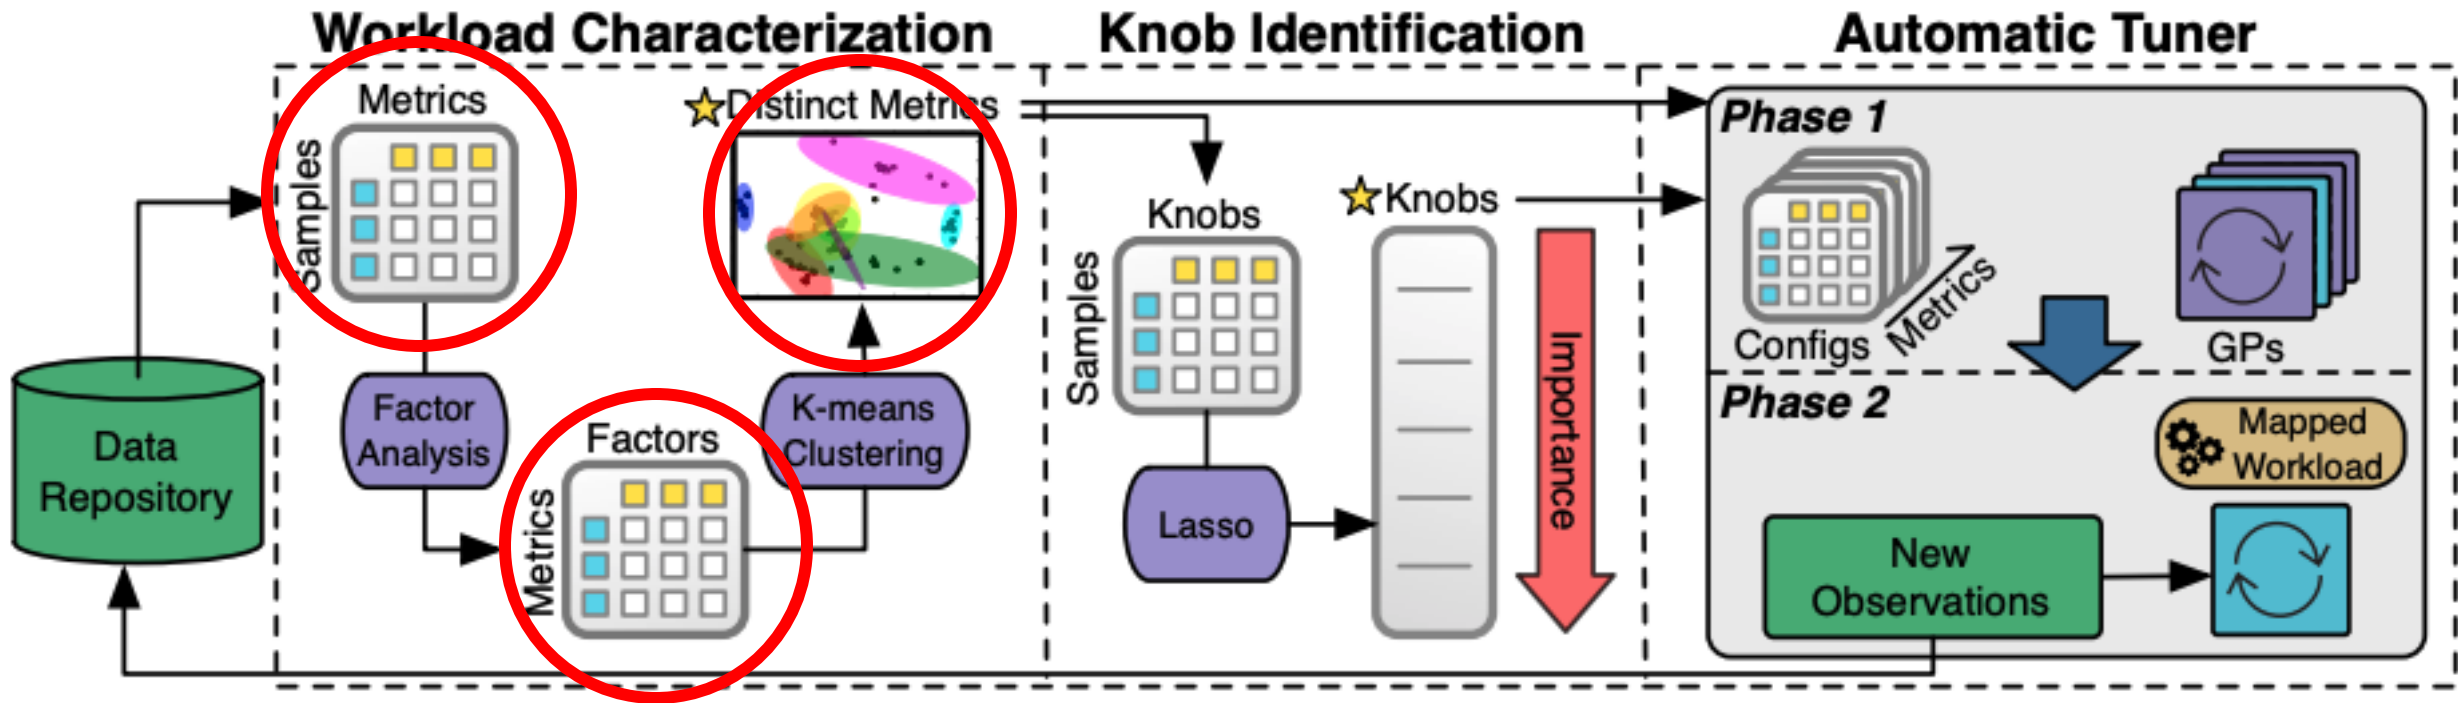


What are possible challenges of this approach? .



A workload is characterized based on the system metrics when it is executed (e.g., #pages reads/writes, cache utilization, locking overhead)

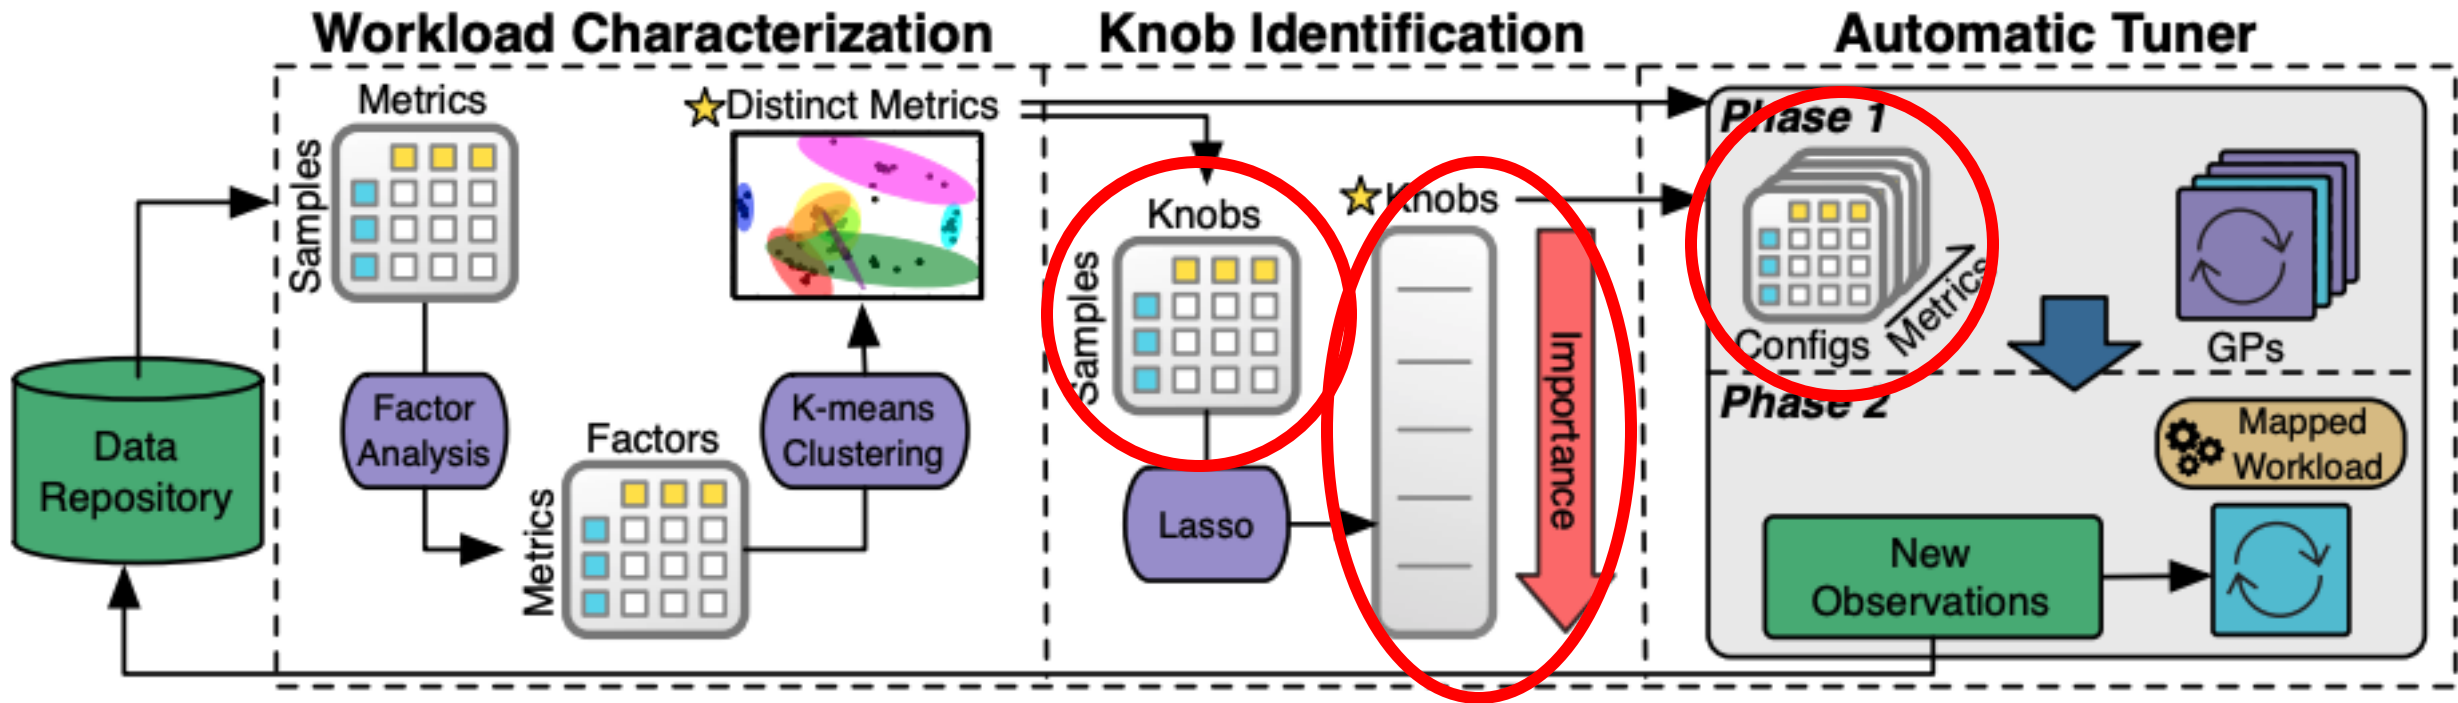
OtterTune Machine Learning Pipeline



Collect statistics at the global level (system-wide), per table proves to be challenging for various systems

Prune redundant metrics (e.g., data read and pages read are directly linked) via factor analysis and k-means clustering

OtterTune Machine Learning Pipeline



Identify important knobs

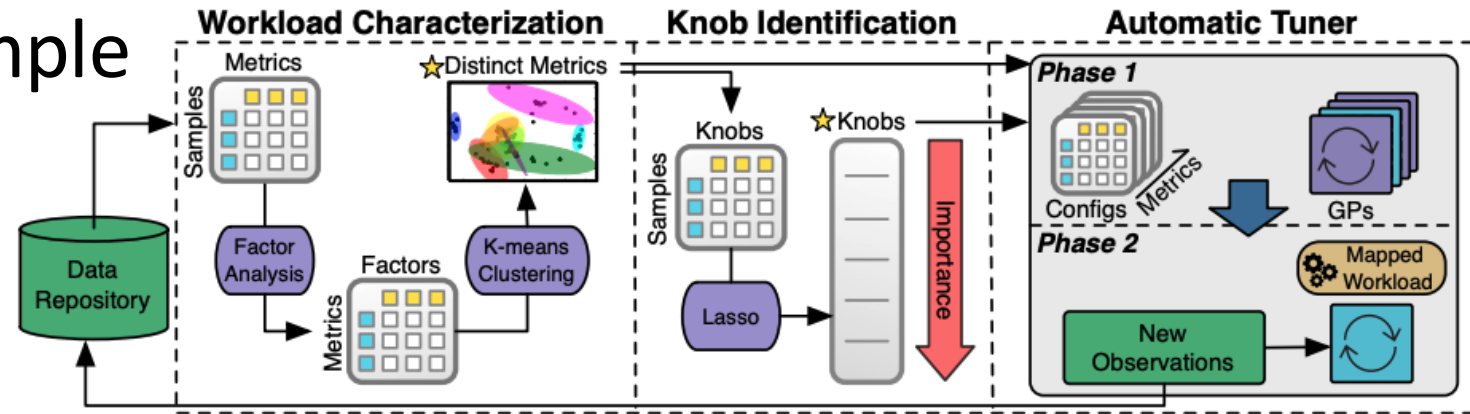
Order the knobs based on their significance on the system's performance (and identify knobs interdependencies)

Store in a repository observations

OtterTune Machine Learning Pipeline

Automated Tuning: an Example

Use the systems metrics to identify (classify) the workload




Iterative configuration recommendation balancing **exploration** vs. **exploitation**

Exploration: try out a configuration for which there is not enough data in the repository
this is done when (i) there is not enough data for this workload (so more data are needed), or
(ii) the system decides to try out new configurations that help collect more data in general

Exploitation: the systems uses small variations of a configuration that is close to optimal using the existing data

OtterTune in Action

Start by sweeping values of knobs to collect “training data”




The optimal number of knobs varies per *DBMS* and *workload*!

Increasing the number of knobs gradually is the best approach, because it balances complexity and performance.

OtterTune tunes MySQL and Postgres that have few impactful knobs, and Action Vector that requires more knobs to be tuned in order to achieve good performance.


OtterTune vs iTunes on TPCC

iTuned uses an initial set of 10 DBMS configurations at the beginning of the tuning session.



OtterTune is trained with more data, so it can achieve a better end result!

OtterTune vs iTunes on TPC-H





Action Vector allows fewer “bad” options, so the training is easier.


“A tuning knob is a database engineer not knowing what do”

take this with a grain of salt!

OtterTune Efficacy Comparison




(a) TPC-C (Throughput)




(b) TPC-C (99%-tile Latency)

MySQL



(a) TPC-C (Throughput)

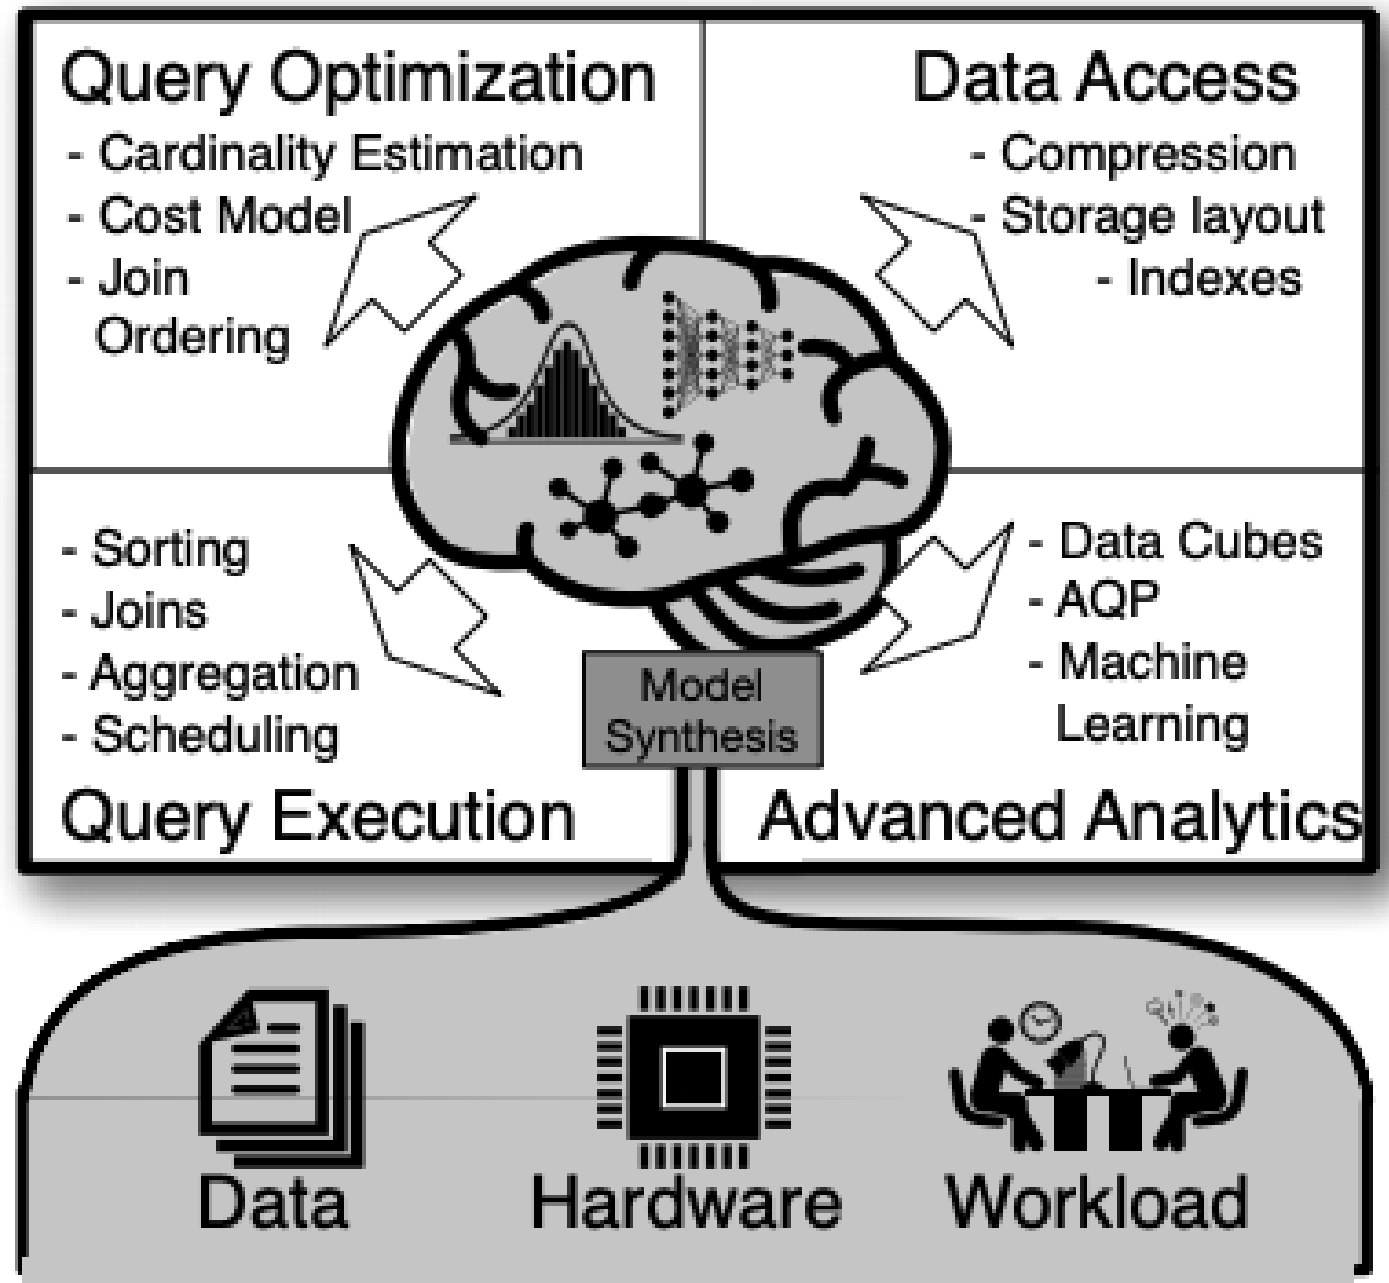



(b) TPC-C (99%-tile Latency)

PostgreSQL

It is hard (but not impossible) to beat an expert DBA!

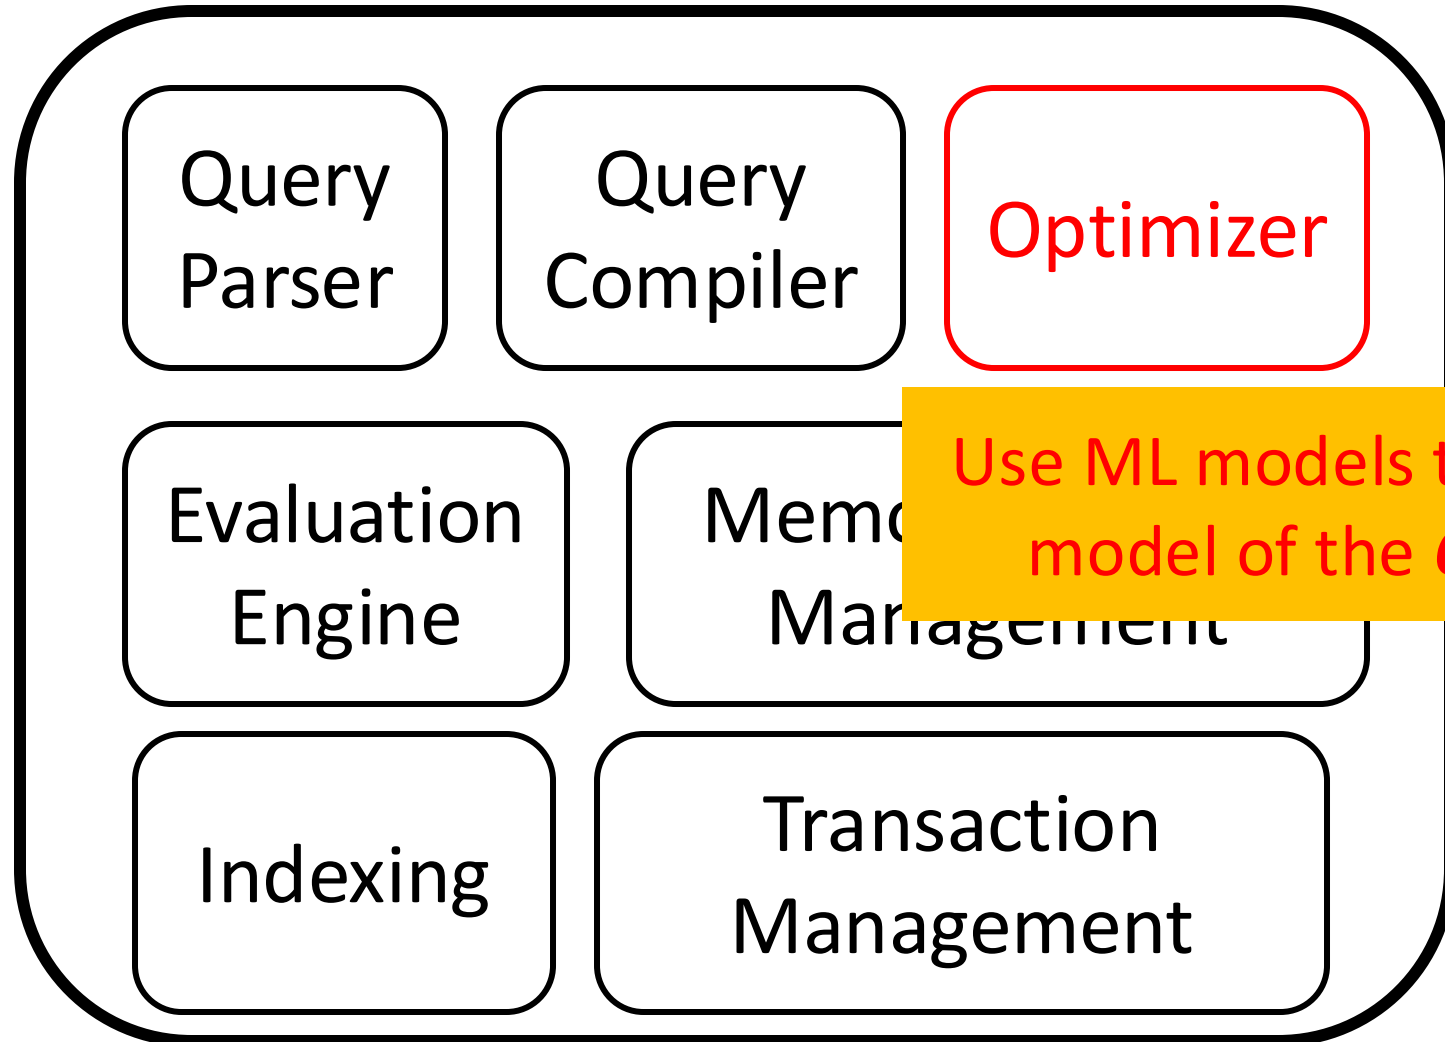
A Learned Database System






*application/SQL
access patterns
complex queries*

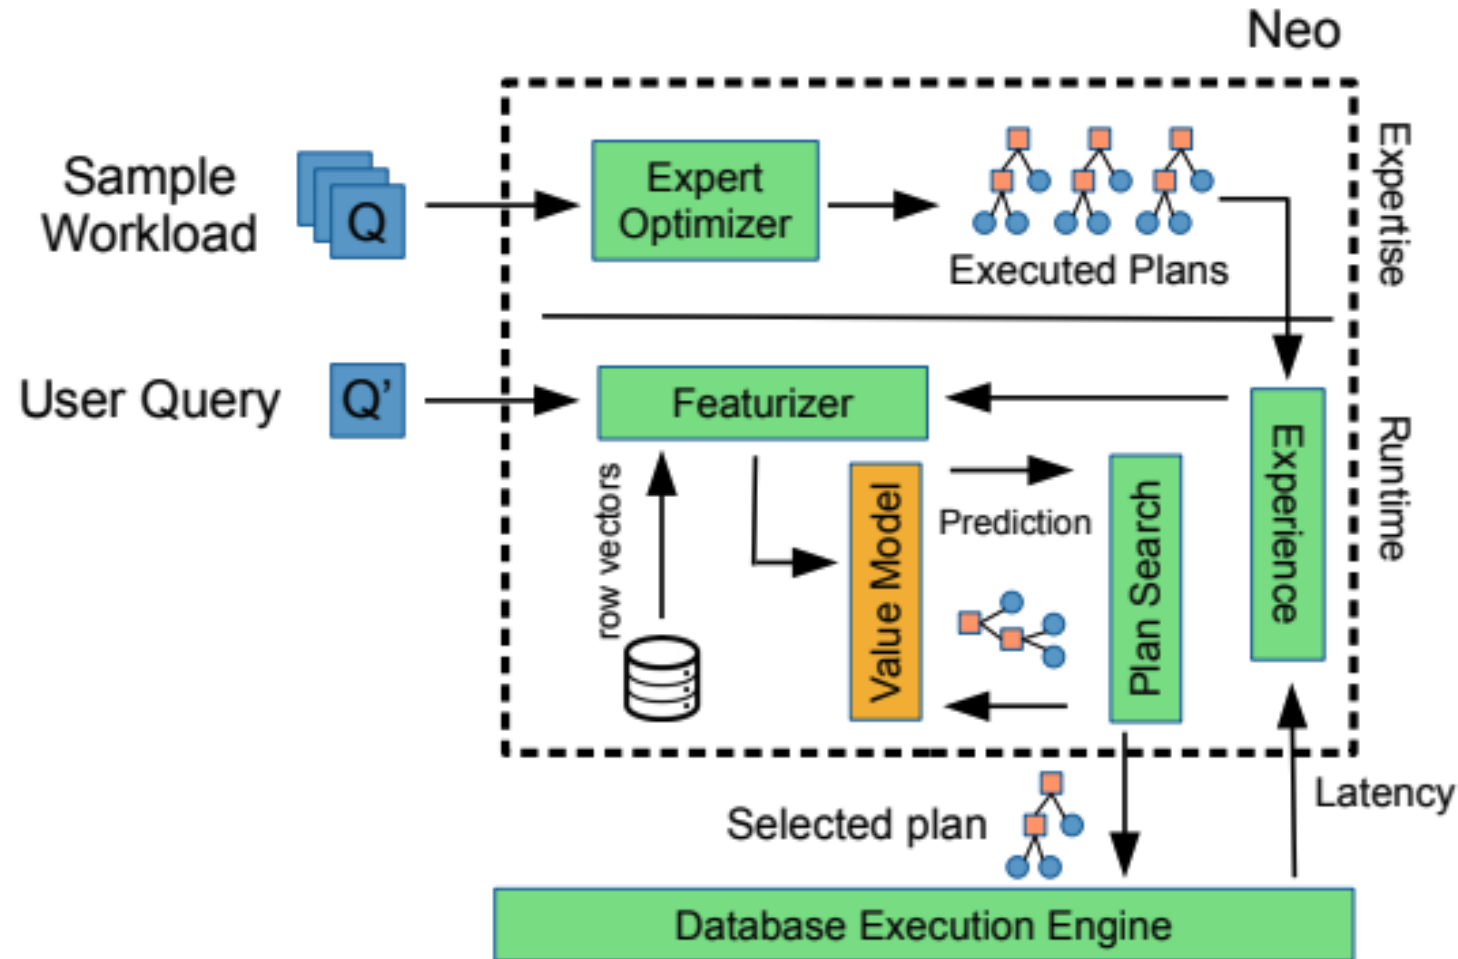
modules




Use ML models to replace the cost-model of the **Query Optimizer**



Learned Query Optimization



Learned Query Optimization







PostgreSQL

SQLite





MS SQL Server

Oracle





JOB




TPC-H




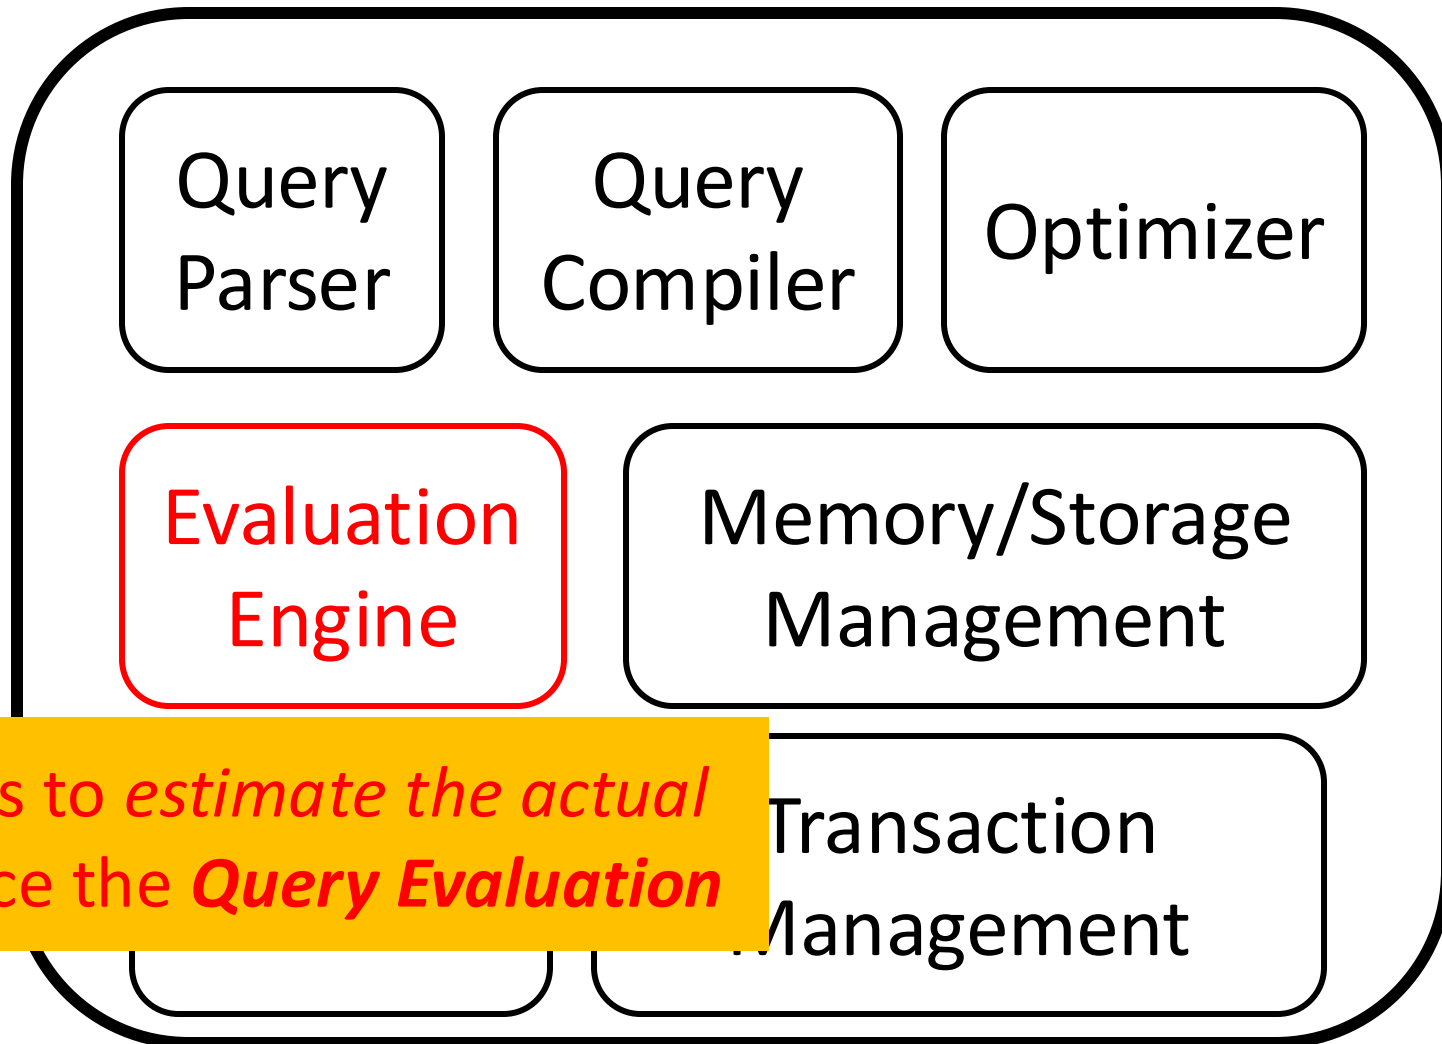
Corp



modules




*application/SQL
access patterns
complex queries*



Use ML models to *estimate the actual data* and replace the **Query Evaluation**

Motivation

In the era of big data, exact analytical query processing is too “expensive”.



Agarwal, Sameer, et al. "BlinkDB: queries with bounded errors and bounded response times on very large data." *Proceedings of the 8th ACM European Conference on Computer Systems*. ACM, 2013.

Motivation

In the era of big data, exact analytical query processing is too “expensive”.

A large class of analytical queries takes the form:

```
SELECT AF(y) FROM table  
WHERE x BETWEEN lb AND ub  
[GROUP BY z]
```

Such queries are very popular on emerging datasets/workloads: IoT, sensors, scientific, etc.

Approximate Query Processing

Targeting ***Analytical*** Queries – **why?**

Goal: fast data analytics over large volumes of data

Tradeoff: accuracy vs. latency – **why?**

Is an accurate response always necessary?

exploratory analytics, business intelligence, analytics for ML

Basic tool: sampling

Current Solutions

- Online Aggregations
- Data Sketches
- Sample-based Approaches (the dominating approach)



Uniform Sampling

Limited supported aggregate functions

Stratified Sampling

Still, very time-consuming

Hash Sampling

Space Overhead – samples can be very large

Support for join (multi-way)

Support for nesting

Query-time sampling

Queries ***explicitly specify*** sample operations

Sample then execute query

Uniform sampling: may miss small groups

Distinct sampler: online sampling of distinct values

With joins: want to sample ***before*** joins not after – **why?**

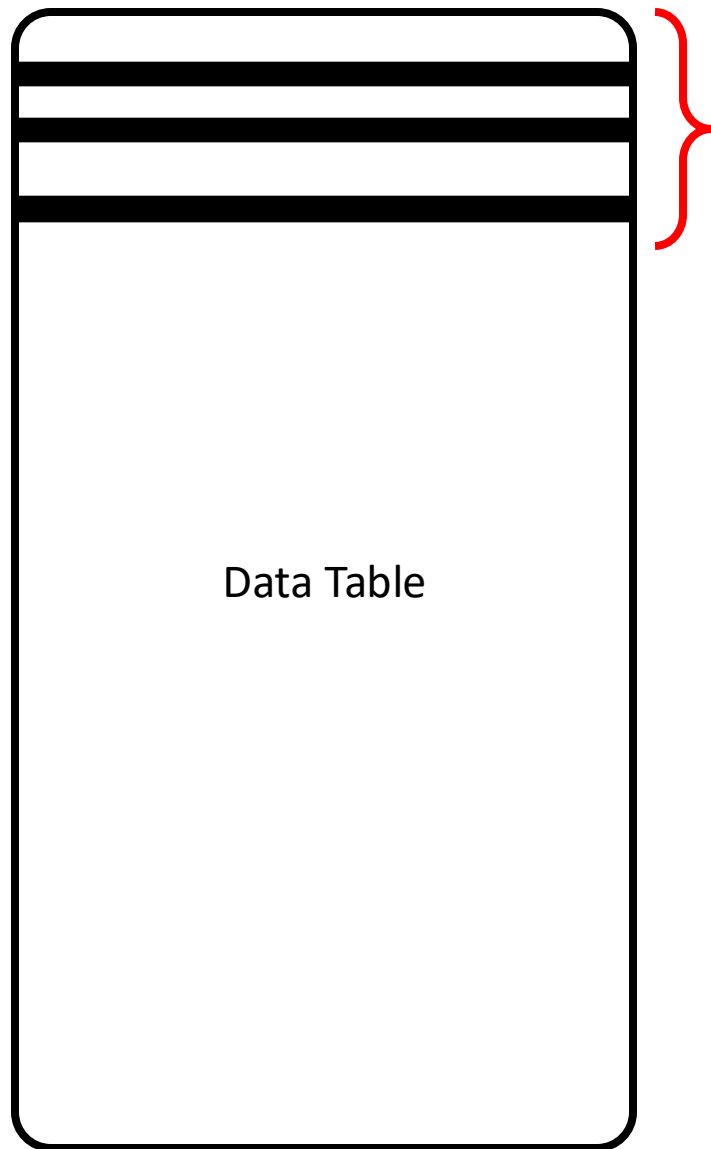
Online aggregation

Execute query on growing random samples

Preliminary outputs are constantly updated – **which?**

Query result

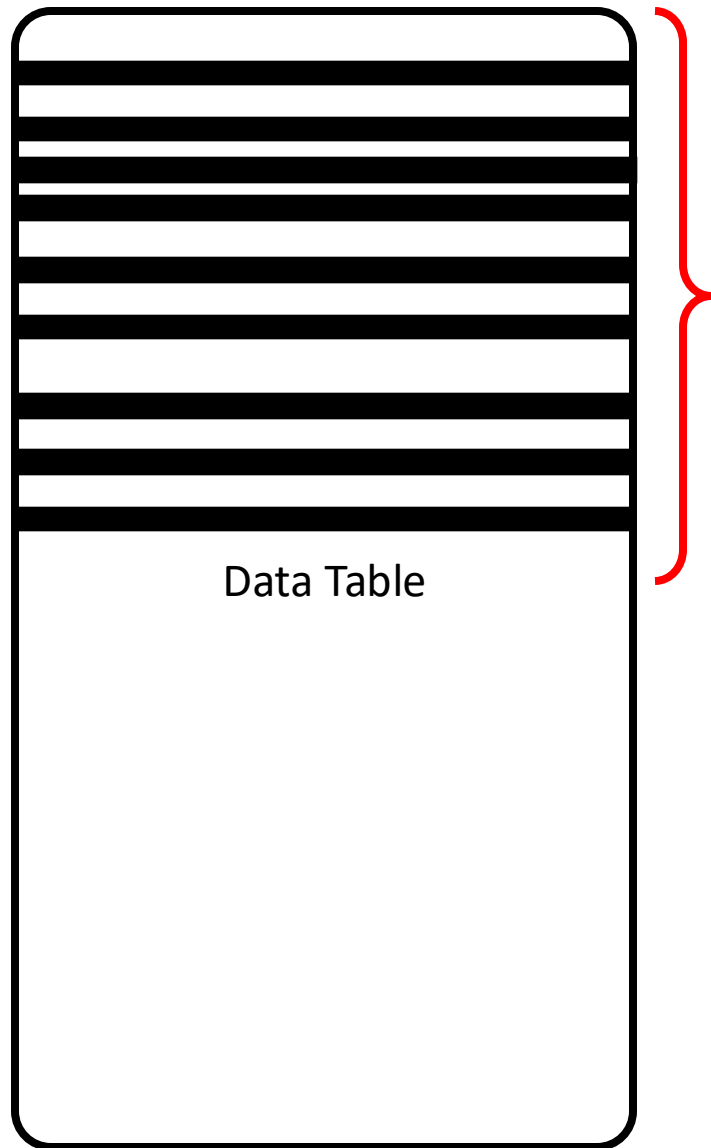
Estimated error



The diagram shows a vertical rectangle representing a data table. The top portion of the rectangle is divided into five horizontal bands: a thin white band, a thick black band, a thin white band, a thick black band, and a thin white band. A red bracket on the right side of these bands indicates a specific section of the table. The rest of the rectangle is a single large white area.

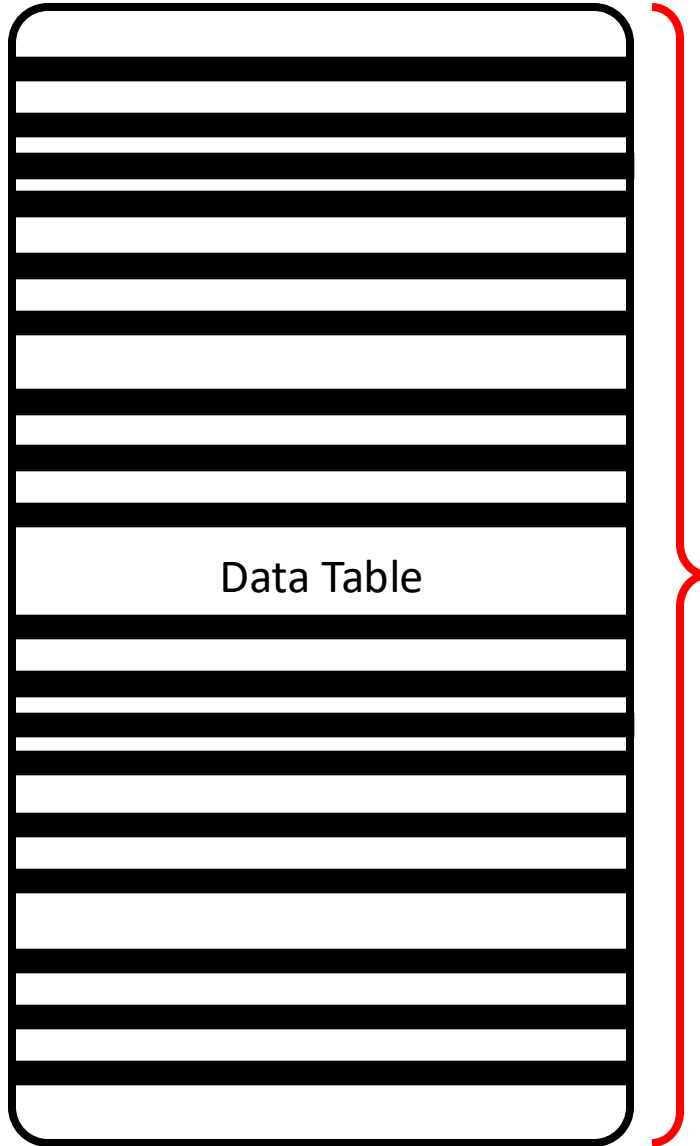
Data Table

expected mean: 1003
[990, 1020] with confidence 95%



Data Table

expected mean: 1002
[995, 1007] with confidence 96%



Data Table

expected mean: 1001
[1001, 1001] with confidence 100%

Online aggregation

Execute query on growing random samples

Preliminary outputs are constantly updated – **which?**

- Query result

- Estimated error

Hard to execute efficiently – **why?**

- Random sample → Random access

- Random samples might contain few rows that join

- Can be improved using join indices

Queries on Pre-Computed Samples

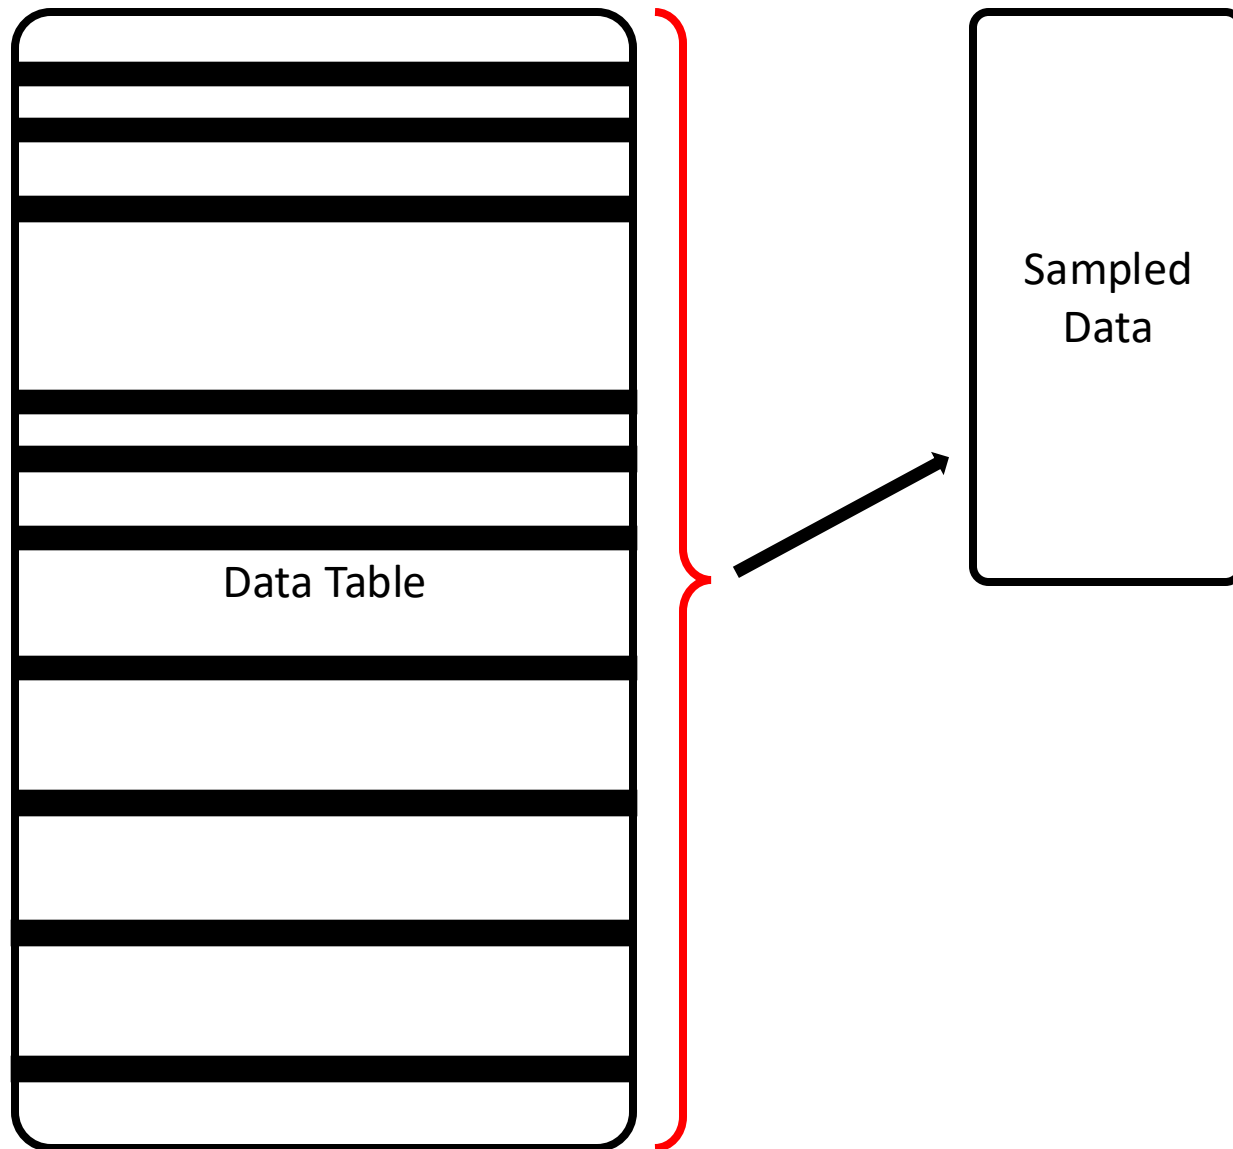
Low latency because **sampling cost** is assumed **offline**
operate **only on the sample**

Additional space (to keep sample)

Cannot provide fixed error bounds

Error bounds are data dependent (high variance = large error)

They can be arbitrarily large



SQL additions

Aggregate is computed on a group

Group is defined based on certain columns

Extend specification with bounds

Error-bound query

```
SELECT count(*)  
FROM Sessions  
WHERE Genre=`western`  
GROUP BY OS  
ERROR WITHIN 10% AT CONFIDENCE 95%
```

Time-bound query

```
SELECT count(*)  
FROM Sessions  
WHERE Genre=`western`  
GROUP BY OS  
WITHIN 5 SECONDS
```

Offline vs online sampling

	Offline	Online
Assumption:	(partially) known workload	No assumption
Speedup:	High	Low

Offline vs online sampling

	Offline	Online
Assumption:	(partially) known workload	No assumption
Speedup:	High	Low

Offline vs online sampling

	Offline	Online
Assumption:	(partially) known workload	No assumption
Speedup:	High	Low

Both are helpful:

- offline sampling is used for (partially) predictable workloads,
- online sampling is for the rest.

DBEst: transparent AQP

Very small query execution times (e.g., ms),
With **small state** (memory/storage footprint) (e.g., KBs), and
High accuracy (e.g., a few % relative error)
Regardless of data size?

YES! (for a large class of analytical queries)
rests on simple SML models
Built over samples of tables

DBEst Contributions

DBEst shows that

- Models can be built over small samples

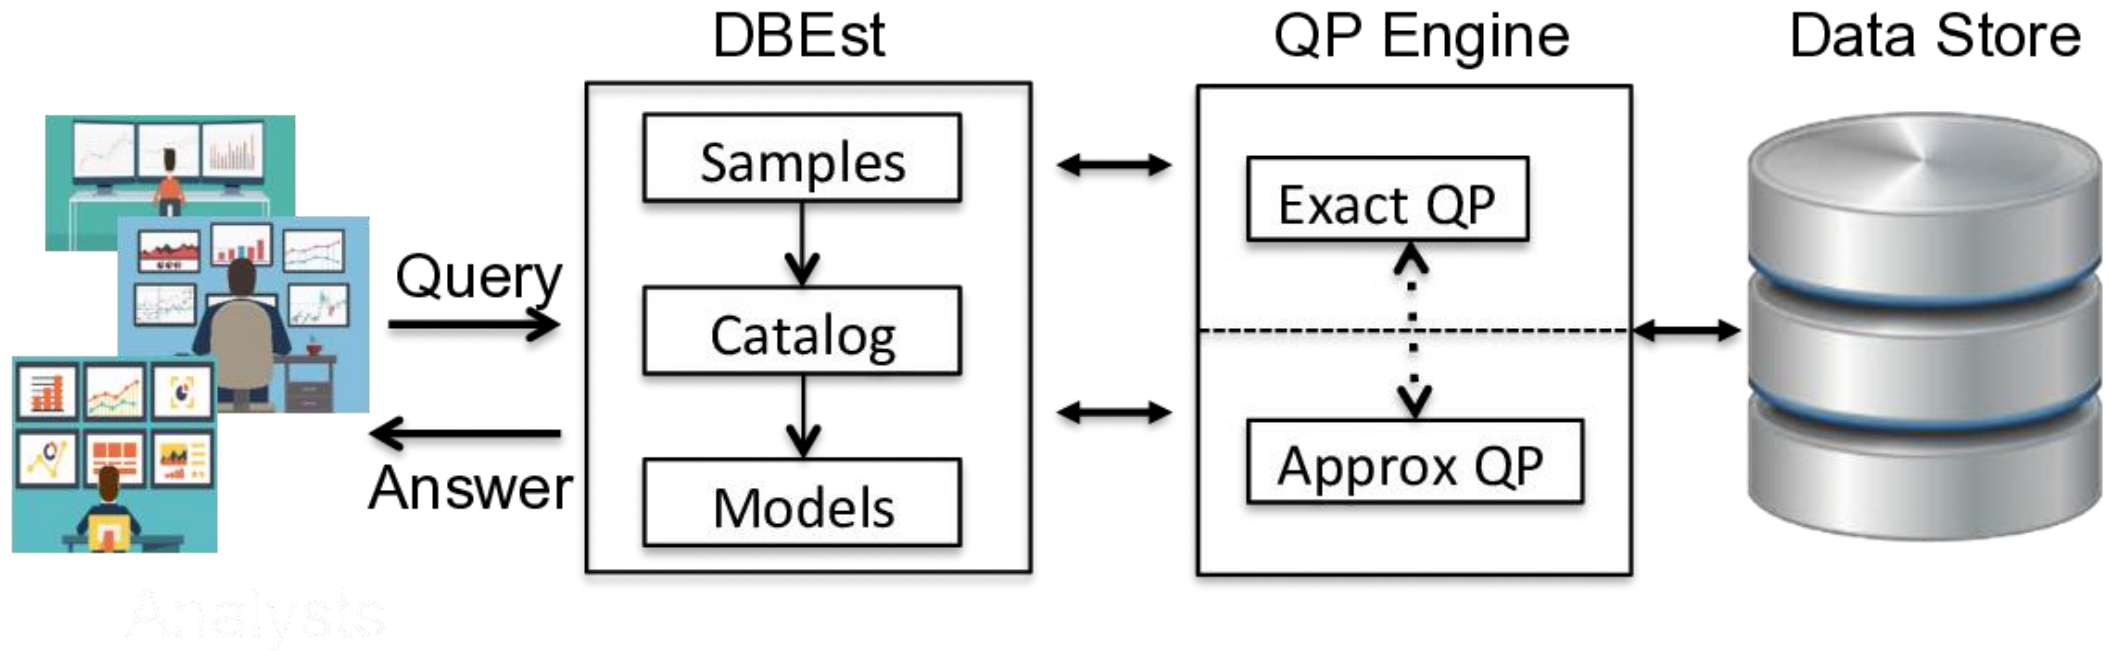
- Can generalize nicely, ensuring accuracy

- Model state is small (KBs)

- AQP over models is much faster than over samples***

- Model training overhead is acceptable – inline with sample generation.

DBEst Architecture



DBEst and ML models

*which **aggregate functions** are very **hard** to
answer via **approximate** query processing?*

- Problem SQL query

SELECT AF(**y**) from table

WHERE **x** between *low* and *high*

[GROUP BY **z**]

which are easy?

- What models?


Regression $y=R(x)$

- LR, PR...
- **XGBoost**, GBoost...

Density Estimator
 $D(x)$

- **Kernel Density**
- Nearest neighbor method
- Orthogonal series estimator

Density Estimator




Histograms is the simplest form of **density estimator**

DBEst is **gradually learning** a function
that **approximates** the **actual density** function of the data

e.g., “how many values exist between *low* and *hi*?”

Regression Model



A **regression model** describes the **relationship between two variables**
 $y = F(x)$

DBEst uses a regression model to capture “matches” from selection

e.g., “which values of y exist for x between low and hi?”

How to use regression and density estimation to answer queries?

```
SELECT count(*)  
FROM Table  
WHERE x between lb and ub
```

$$COUNT(y) \approx N \cdot \int_{lb}^{ub} D(x) dx$$

fraction of values in [lb,ub]

```
SELECT avg(y)  
FROM Table  
WHERE x between lb and ub
```

$$\begin{aligned} AVG(y) &= \mathbb{E}[y] \\ &\approx \mathbb{E}[R(x)] \\ &= \frac{\int_{lb}^{ub} D(x) R(x) dx}{\int_{lb}^{ub} D(x) dx} \end{aligned}$$

relationship of x values with y values

fraction of values in [lb,ub]

```
SELECT sum(y)  
FROM Table  
WHERE x between lb and ub
```

$$\begin{aligned} SUM(y) &= COUNT(y) \cdot AVG(y) \\ &\approx COUNT(y) \cdot \mathbb{E}[R(x)] \\ &= N \cdot \int_{lb}^{ub} D(x) dx \cdot \frac{\int_{lb}^{ub} D(x) R(x) dx}{\int_{lb}^{ub} D(x) dx} \\ &= N \cdot \int_{lb}^{ub} D(x) R(x) dx \end{aligned}$$

How to use regression and density estimation to answer queries?

```
SELECT variance(y)
FROM Table
WHERE x between lb and ub
```

$$\begin{aligned} \text{VARIANCE}_y(y) &= \mathbb{E}[y^2] - [\mathbb{E}[y]]^2 \\ &\approx \mathbb{E}[R^2(x)] - [\mathbb{E}[R(x)]]^2 \\ &= \frac{\int_{lb}^{ub} R^2(x)D(x)dx}{\int_{lb}^{ub} D(x)dx} - \left[\frac{\int_{lb}^{ub} R(x)D(x)dx}{\int_{lb}^{ub} D(x)dx} \right]^2 \end{aligned}$$

PERCENTILE.

If the reverse of the CDF, $F^{-1}(p)$, could be obtained, then the p^{th} percentile for Column x is

```
SELECT percentile(x,p)
FROM Table
```

$$\alpha = F^{-1}(p) \tag{5}$$

Note that $F^{-1}(p)$ is derived using $F(p) = \int_{-inf}^p D(x)dx$

More support on SQL

```
SELECT avg(y)
FROM Table
WHERE x1 between lb1 and ub1
      AND x2 between lb2 and ub2
```

$$\begin{aligned} AVG(y) &= \mathbb{E}[y] \\ &\approx \mathbb{E}[R(x_1, x_2)] \\ &= \frac{\int_{lb_1}^{ub_1} \int_{lb_2}^{ub_2} D(x_1, x_2) R(x_1, x_2) dx_2 dx_1}{\int_{lb_1}^{ub_1} \int_{lb_2}^{ub_2} D(x_1, x_2) dx_2 dx_1} \end{aligned}$$

Supporting GROUP BY

- build models for each group by value,
- create **model bundles**:
 - E.g., each bundle stores ~500 groups
 - Store bundles in, say, an SSD (~100 ms to deserialize and compute AF on bundle).

Supporting join

- Join table is flattened -> make samples -> build models.

Limitations

- Group By Support ->too many groups
 - Model Training time ↑, Query Response time ↑, space overhead ↑.
- No error guarantee

DBEst Summary

- DBEst: a model-based AQP engine, using simple SML models:
 - Much smaller query response times
 - High(er) accuracy
 - Much smaller space-time overheads
 - Scalability
- Ensuring high accuracy, efficiency, scalability with low money investments -- resource (cpu, memory/storage/ network) usage.
- Future work: more efficient support for
 - Joins
 - Categorical attributes
 - Improved parallel/distributed DBEst

A perspective on ML in Database Systems

from: ML-In-Databases: Assessment and Prognosis, IEEE Data Engineering Bulletin

New *Forces*

(1) End-user want to

- democratize data (all business units to have access to all data)
- make data-driven decisions (often in real time)

(2) New applications

- structured query processing (SQL) + natural language processing (NLP) + Complex Analytics (exploratory + predictive ML)

New *Forces*

(3) Data integration

diverse and inconsistent datasets are combined in common data repositories (data lakes)

(2) New hardware + the move to the cloud

moving from full ownership to pay-as-you-go

self-tuning systems *en masse* in the cloud (as we discussed today)

Consequences and New Directions

Storage ***hierarchy*** is still relevant, but the layers are elastic (in the cloud)

ML models can be deployed at-will as “functions”

New push for ***serverless computing***

use only services and not rent an entire server

CS 561: Data Systems Architectures

class 22

Machine Learning & Data Systems

Prof. Manos Athanassoulis

<https://bu-disc.github.io/CS561/>

DBEst experiments

Evaluation

systematically showing sensitivities on

- range predicate selectivity + sample sizes + AFs

Performance under Group By and Joins

Comparisons against

- State of the art AQP (VerdictDB and BlinkDB)
- State of the art columnar DB (MonetDB)

Using data from TPC-DS and 3 different UCI-ML repo datasets.

Experimental Setup

Ubuntu 18.04 with Xenon X5650 12-core CPU, 64 GB RAM And 4TB SSD

Datasets: TPC-DS, Combined Cycle Power Plant (CCPP), Beijing PM2.5

Query types:

- Synthetic queries: 0.1%, 1%, to 10% query range
- Number of queries: vary between 30 to 1000 queries.
- Complex TPC-DS queries: Query 5, 7, and 77.

Compared against VerdictDB, BlinkDB and MonetDB, for error

- VerdictDB uses 12 cores while DBEst runs on 1 core. (Multi-threaded DBEst is also evaluated)


Report execution times + system throughput for the parallel version

Report performance of joins and group by

Performance – Sensitivity Analysis

Query range effect

Dataset: TPC-DS
Sample size: 100k
540 synthetic queries
Column pair:
[ss_list_price, ss_wholesale_cost]




Influence of query range on relative error


Performance – Sensitivity Analysis

Sample size effect

Dataset: TPC-DS
Query range: 1%
1200 synthetic queries
Column pair:
[ss_list_price, ss_wholesale_cost]



Influence of sample size on relative error




Influence of sample size on space overhead


Performance Comparison

TPC-DS dataset

Query range: 0.1%, 1%, 10%
~100 queries, involving 16
column pairs.
Sample size: 10k, 100k



Relative Error: DBEst vs VerdictDB




Query Response Time: DBEst vs VerdictDB


Performance Comparison

CCPP dataset

2.6 billion records, 1.4TB
Query range: 0.1%, 0.5%, 1.0%
108 queries, involving 3 column pairs.
Sample size: 10k, 100k



Relative error (10k sample)




Relative error (100k sample)


Performance Comparison Group By

```
SELECT AF(ss_list_price)
FROM store_sales
WHERE ss_wholesale_cost_sk ...
GROUP BY ss_store_sk
```

- 90 queries, 57 groups
- Sample size: 10k



Relative error for group by queries





Accuracy histogram for SUM


Performance Comparison Join

```
SELECT AF(ss_wholesale_cost), AF(ss_net_profit)
FROM store_sales, store
WHERE ss_store_sk=s_store_sk
AND s_number_of_employees BETWEEN ...
```

- 42 queries.




Parallel Query Execution



Group by query response time reduction (TPC-DS)

1 core versus 12 cores



Throughput of parallel execution (CCPP)