



# Cloud Data Lakes

..Or “LakeHouse”  
..Or “Open Data Lake”

Andrew Lamb | Staff Engineer, InfluxData

April 17, 2025, Boston University Guest Lecture, [CS-561](#)





**influxdata**<sup>®</sup>



---

**Andrew Lamb**

Staff Engineer  
InfluxData

> 22 🤓 years in enterprise software development

Oracle: Database (2 years)

DataPower: XSLT compiler (2 years)

Vertica: DB / Query Optimizer (6 years)

Nutonian/DataRobot: ML Startups (7 years)

InfluxData: InfluxDB 3.0, Arrow, DataFusion (5 years)

MIT VI-2 2002, MEng 2003

# Outline

- Why this topic is important
- Database Architectures through the Ages
- Trends driving the move “to the cloud”
- Disaggregated Databases, and common architecture features

# Context:



## Lakehouse: A New Generation of Open Platforms that Unify Data Warehousing and Advanced Analytics

Michael Armbrust<sup>1</sup>, Ali Ghodsi<sup>1,2</sup>, Reynold Xin<sup>1</sup>, Matei Zaharia<sup>1,3</sup>

<sup>1</sup>Databricks, <sup>2</sup>UC Berkeley, <sup>3</sup>Stanford University

### Abstract

This paper argues that the data warehouse architecture as we know it today will wane in the coming years and be replaced by a new architectural pattern, the Lakehouse, which will (i) be based on open direct-access data formats, such as Apache Parquet, (ii) have first-class support for machine learning and data science, and (iii) offer state-of-the-art performance. Lakehouses can help address several major challenges with data warehouses, including data staleness, reliability, total cost of ownership, data lock-in, and limited use-case support. We discuss how the industry is already moving toward Lakehouses and how this shift may affect work in data management. We also report results from a Lakehouse system using Parquet that is competitive with popular cloud data warehouses on TPC-DS.

### 1 Introduction

This paper argues that the data warehouse architecture as we know it today will wane in the coming years and be replaced by a new architectural pattern, which we refer to as the Lakehouse, characterized by (i) open direct-access data formats, such as Apache Parquet and ORC, (ii) first-class support for machine learning and data science workloads, and (iii) state-of-the-art performance.

The history of data warehousing started with helping business leaders get analytical insights by collecting data from operational databases into centralized warehouses, which then could be used for decision support and business intelligence (BI). Data in these warehouses would be written with schema-on-write, which ensured that the data model was optimized for downstream BI consumption. We refer to this as the first generation data analytics platforms.

A decade ago, the first generation systems started to face several challenges. First, they typically coupled compute and storage into an on-premises appliance. This forced enterprises to provision and pay for the peak of user load and data under management, which became very costly as datasets grew. Second, not only were datasets growing rapidly, but more and more datasets were completely unstructured, e.g., video, audio, and text documents, which data warehouses could not store and query at all.

To solve these problems, the second generation data analytics platforms started offloading all the raw data into *data lakes*: low-cost storage systems with a API that hold data in generic and usually open file formats, such as Apache Parquet and ORC [8, 9]. This approach started with the Apache Hadoop movement [5], using the Hadoop File System (HDFS) for cheap storage. The data lake was a schema-on-read architecture that enabled the agility of storing any data at low cost but on the other hand, punted the problem of data

This article is published under the Creative Commons Attribution License (<http://creativecommons.org/licenses/by/3.0/>). 11th Annual Conference on Innovative Data Systems Research (CIDR '21), January 11–15, 2021, Online.

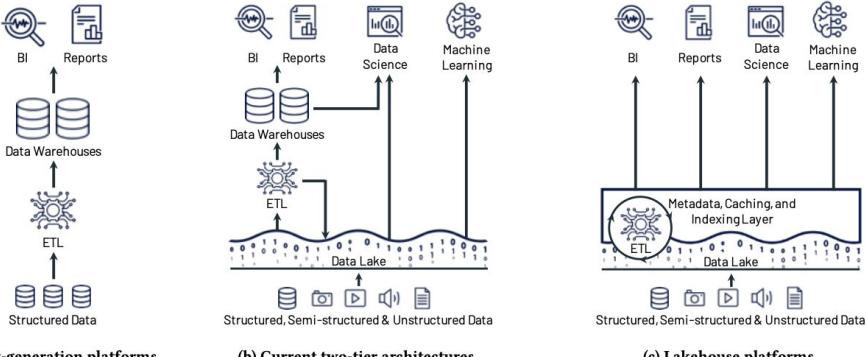


Figure 1: Evolution of data platform architectures to today's two-tier model (a-b) and the new Lakehouse model (c).

This article is published under the Creative Commons Attribution License (<http://creativecommons.org/licenses/by/3.0/>). 11th Annual Conference on Innovative Data Systems Research (CIDR '21), January 11–15, 2021, Online.

# Context:



[databricks.com/blog/2021/11/02/databricks-sets-official-data-warehousing-performance-record.html](https://databricks.com/blog/2021/11/02/databricks-sets-official-data-warehousing-performance-record.html)

[snowflake.com/en/blog/industry-benchmarks-and-databricks-set-official-data-warehousing-performance-record.html](https://snowflake.com/en/blog/industry-benchmarks-and-databricks-set-official-data-warehousing-performance-record.html)

## Databricks Sets Official Data Warehousing Performance Record

Published: November 2, 2021 | Data Warehousing | 10 min read | By Reynold Xin and Mostafa Mokhtar

Today, we are proud to announce that [Databricks SQL](#) has set a [new world record](#) in [100TB TPC-DS](#), the gold standard performance benchmark for data warehousing. [Databricks SQL outperformed the previous record by 2.2x](#). Unlike most other benchmark news, this result has been formally audited and reviewed by the TPC council.

These results were corroborated by research from Barcelona Supercomputing Center, which frequently runs benchmarks that are derivative of TPC-DS on popular data warehouses. Their latest research [benchmarked Databricks and Snowflake, and found that Databricks was 2.7x faster and 12x better in terms of price performance](#). This result validated the thesis that data warehouses such as Snowflake become prohibitively expensive as data size increases in production.

Databricks has been rapidly developing full blown data warehousing capabilities directly on data lakes, bringing the best of both worlds in one data architecture dubbed [the data lakehouse](#). We announced our full suite of data warehousing capabilities as Databricks SQL in November 2020. The open question since then has been whether an open architecture based on a lakehouse can provide the

snowflake.com/en/blog/industry-benchmarks-and-databricks-set-official-data-warehousing-performance-record.html

[snowflake](#) Product Solutions Why

BLOG CATEGORY

STRATEGY & INSIGHTS NOV 12, 2021

## Industry Benchmark Integrity

Databricks perform and outperf

Price  
Snowflake  
Databricks

When we founded Snowflake, we set out to build an account what had worked well and what hadn't in prior or leverage the cloud to rethink the limits of what was possible system that "just worked." We knew there were many opportunities to lead on performance and scale, simplicity and innovation to lead on performance and scale, simplicity and innovation.

In the same way that we had clarity about many things we didn't want to do. One such thing was engaging in benchmarks divorced from real-world experiences. This practice customers first.

Twenty years ago, the game of leapfrogging benchmark results in industry and both of us were on the front line fighting the new world records were being set on a regular basis. Most special settings, and very specific optimizations that would have little or even negative impact on customers' day-to-day development teams are distracted from focusing on what underserved with more complex technology. Anyone who saw the reality that the benchmark race became a distraction reason why all the relevant players in the database industry workloads, have largely stopped publishing new results.

Since founding Snowflake, we have focused on our customers and their workloads, and not on synthetic

MPP architecture

There inherent latency trade off



Why Databricks Product Solutions Resources About

DATA + AI SUMMIT

Blog / Data Warehousing / Article

## Snowflake Claims Similar Price/Performance to Databricks, but Not So Fast!

Published: November 15, 2021

Data Warehousing

6 min read

By Mostafa Mokhtar, Reynold Xin and Matei Zaharia

On Nov 2, 2021, we announced that [we set the official world record](#) for the fastest data warehouse with our Databricks SQL lakehouse platform. These results were audited and reported by the official Transaction Processing Performance Council (TPC) in a 37-page document [available online](#) at tpc.org. We also shared a third-party benchmark by the Barcelona Supercomputing Center (BSC) outlining that Databricks SQL is significantly faster and more cost effective than Snowflake.

A lot has happened since then: many congratulations, some questions, and some sour grapes. We take this opportunity to reiterate that [we stand by our blog post and the results: Databricks SQL provides superior performance and price performance over Snowflake, even on data warehousing workloads \(TPC-DS\)](#).

Snowflake's response: "lacking integrity"?

Snowflake responded 10 days after our publication (last Friday) claiming that our results were "lacking integrity." They then presented their own benchmarks



Blog

Snowflake Claims Similar Price/Performance to Databricks, but Not So Fast!

Keep up with

Recommendations



Blog

Eliminating for Databr



# Table Format wars 🗡 (preview)

## Delta Lake: High-Performance ACID Table Storage over Cloud Object Stores

Michael Armbrust, Tathagata Das, Liwen Sun, Burak Yavuz, Shixiong Zhu, Mukul Murthy, Joseph Torres, Herman van Hovell, Adrian Ionescu, Alicja Luszczak, Michal Switakowski, Michał Szafrański, Xiao Li, Takuya Ueshin, Mostafa Mokhtar, Peter Boncz, Ali Ghodsi\*, Sameer Paranjape, Pieter Senster, Reynold Xin, Matei Zaharia\*

Databricks, CWI, UC Berkeley, University of Amsterdam, Intel, Oracle, Alibaba Cloud, Microsoft, AWS, Google, Facebook, LinkedIn, and others.

### ABSTRACT

Cloud object stores such as Amazon S3 are some of the largest and most cost-effective storage systems on the planet, making them an attractive fit to store large datasets and data lakes. Understanding their limitations as key-value stores makes it difficult to achieve ACID transactions and high performance: metadata operations such as listing objects are expensive, and consistency guarantees such as atomic writes are missing. A high-performance ACID table storage layer over cloud object stores initially developed at Databricks, Delta Lake uses a transaction log that is compact and efficient, supports ACID transactions with low time travel, and significantly faster metadata operations for large tabular datasets (e.g., the ability to quickly search billions of table partitions for a specific column value). Delta Lake also attempts to provide high-level features such as automatic data layout optimization, upserts, caching, and audit logs. Delta Lake tables can be queried via Spark, Presto, Redshift, and other systems. Delta Lake is deployed at the scale of Databricks customers that process exabytes of data per day, with the largest instances managing exabyte-scale datasets and billions of objects.

**TYLDR Reference Form**  
Armbrust et al., Delta Lake: High-Performance ACID Table Storage over Cloud Object Stores. *PVLDB*, 13(12), 3411-3424, 2020.  
DOI: <https://doi.org/10.14778/3478.341560>

### 1. INTRODUCTION

Cloud object stores such as Amazon S3 [6] and Azure Blob Storage [17] have become some of the largest and most widely used storage systems on the planet, holding exabytes of data for millions of customers. Applications that add value on top of these services, such as real-as-you-go billing, customer support, and export management [15], cloud object stores are especially attractive because they allow users to scale computing and storage resources separately: for example, a query can process a petabyte of data but only run a cluster for a few seconds, query over it for a few hours,

As a result, many organizations now use cloud object stores to manage large structured datasets in data warehouses and data lakes.

This work is licensed under the Creative Commons Attribution Non-Commercial NoDerivatives 4.0 International License. To view a copy of this license, visit <http://creativecommons.org/licenses/by-nd/4.0/>. For any use beyond those contained in the license, permission is requested by emailing [info@valley.org](mailto:info@valley.org). Copyright is held by the owner(s). Publication rights licensed to ACM. VLDB Endowment, Vol. 13, No. 12, ISSN 2150-8097,  
DOI: <https://doi.org/10.14778/3415478.3415560>.

3411

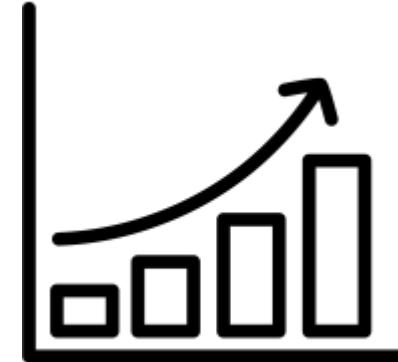


Delta Lake from  
databricks

\* First mover,  
better support

\* Arguably  
technically  
superior

VS



Apache  
**ICEBERG**

\* much faster / wider  
adoption

\* More Neutral  
governance

\* among other  
things championed  
by snowflake

"Delta Lake: High-Performance ACID Table Storage over Cloud Object Stores"

Databricks reportedly paid \$2 billion in Tabular acquisition

## Databricks Agrees to Acquire Tabular, the Company Founded by the Original Creators of Apache Iceberg

June 4, 2024

Share this post



Databricks and Tabular will work together towards a joint vision of the open lakehouse

SAN FRANCISCO — June 4, 2024 — Databricks, the Data and AI company, today announced it has agreed to acquire Tabular, a data management company founded by Ryan Blue, Daniel Weeks, and Jason Reid. By bringing together the original creators of Apache Iceberg™ and Lirux Foundation Delta Lake, the two leading open source lakehouse formats, Databricks will lead the way with data compatibility so that organizations are no longer limited by which of these formats their data is in. Databricks intends to work closely with the Delta Lake and Iceberg communities to bring format compatibility to the lakehouse; in the short term, inside Delta Lake UniForm and in the long term, by evolving toward a single, open, and common standard of interoperability. Databricks and Tabular will work together towards a joint vision of the open lakehouse.

[source](#)

The screenshot shows a news article from TechCrunch. At the top, there's a large image of a digital interface with various data visualizations and text boxes containing "AI". Below the image, the title "IN BRIEF" is visible. The main text discusses the acquisition of Tabular by Databricks, mentioning the joint vision for the open lakehouse and the intention to work closely with the Delta Lake and Iceberg communities. A quote from Maxwell Zeff is included: "Databricks reportedly paid \$2 billion in Tabular acquisition".

Analytics and AI giant Databricks reportedly paid nearly \$2 billion when it [acquired Tabular](#) in June, a startup that was only doing \$1 million in annual recurring revenue, according to [Bloomberg](#). That's a pretty outrageous exit multiple, and it was purportedly fueled by a battle between Databricks and Snowflake.

Tabular had over \$30 million in funding, backed by Altimeter Capital, Andreessen Horowitz and Zetta Venture Partners, when it was acquired

[source](#)

What if you didn't have to choose a format?



ON DEMAND

### Beyond Lakehouse Table Formats

The original creators of Delta Lake and Apache Iceberg™ take on interoperability



### Access on demand

\* First Name

\* Last Name

\* Company Email

\* Company Name

\* Job Title

\* Phone Number

\* Country  United States

### The path to table format interoperability

Choosing the best unified platform for data, analytics and AI is easy — it's lakehouse. Choosing the right open table format for your lakehouse? Not as easy.

For most organizations, it's a daunting decision that delays lakehouse adoption. The holdup hurts your ability to capitalize on analytics and AI.

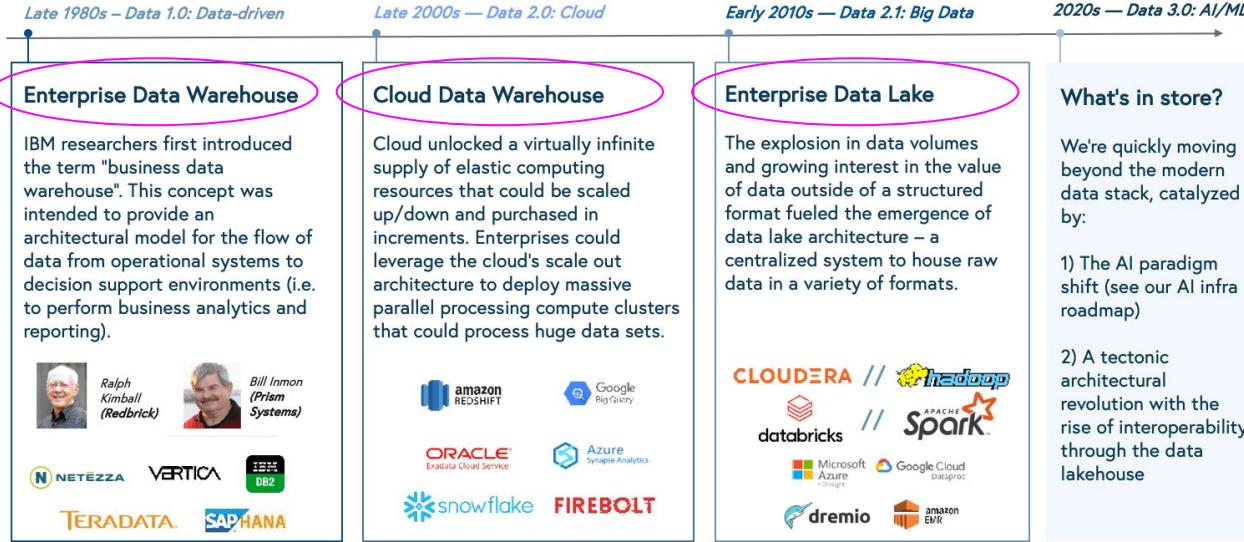
What if you didn't have to choose a format?

Join us for a conversation about interoperability with Michael Armbrust, original creator of Delta Lake, and Ryan Blue, an original creator of Apache Iceberg. They'll discuss the state of open table formats and how Databricks is solving interoperability.

[source](#)

# Roadmap: Data 3.0 in the Lakehouse Era

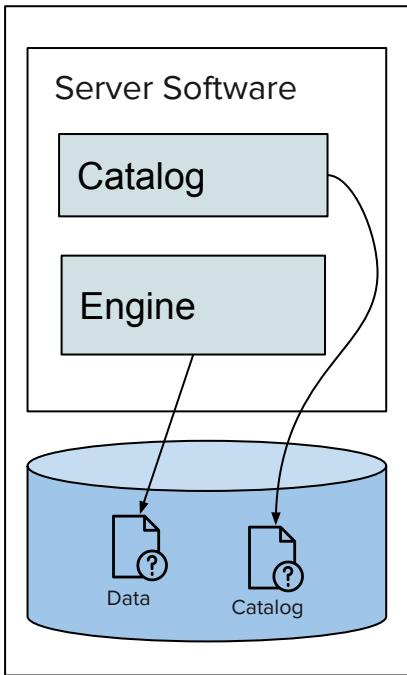
## Enterprise data architecture is constantly evolving



1

[Roadmap: Data 3.0 in the Lakehouse Era - Bessemer Venture Partners \(3/25/2025\)](#)

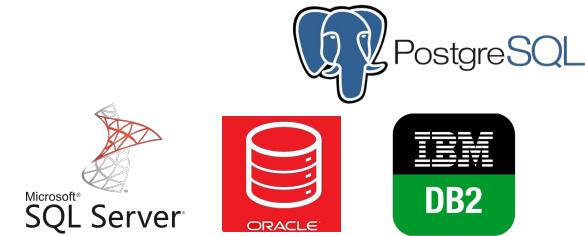
# Single Node Servers (1990s)



Tightly integrated  
**engine, storage**  
and **catalog**

**CPU + MEM +  
STORAGE**

*Data + Catalog  
stored in  
proprietary  
formats on local  
file systems*

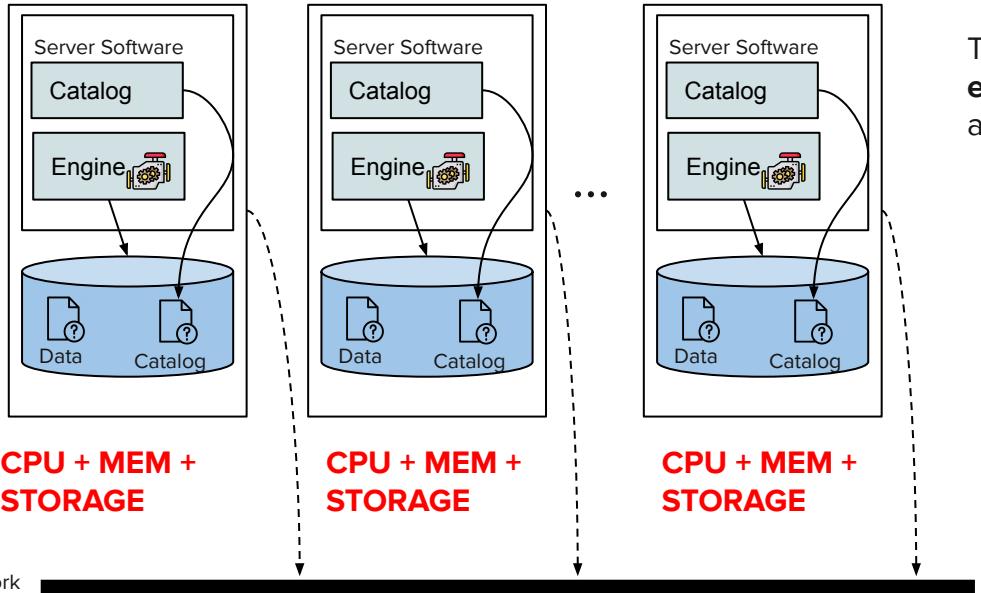


Exemplars

**Driven by**

- Minicomputer → Servers
- Local hard drive capacity

# MPP / Shared Nothing (2000s)



Tightly integrated  
**engine, storage**  
and **catalog**

Data stored on  
each node,  
**next to**  
compute

Proprietary Catalog  
and Data format

Nodes  
*communicate over  
commodity  
network*

**PARACCEL**



**Greenplum**



**VERTICA**

**Exemplars**



**Driven by**

- Fast local networks
- “Inexpensive” (10x cheaper) commodity Linux servers

# Arrival of the “Cloud”

# Object Stores: What

**What:** “Infinite FTP server in the sky”

**What:** Distributed Key/Value stores

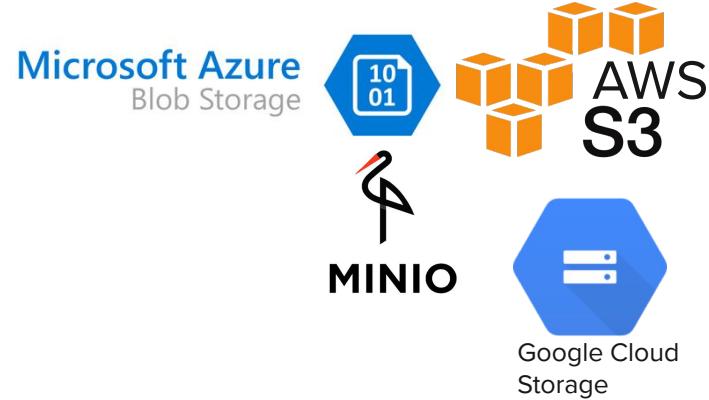
Basic CRUD interface:

**GET** <URL> → Bytes

**PUT** <URL> Bytes

**LIST** <PREFIX> (lists keys)

**DELETE** <URL>



# Object Stores: Why

Why give up nice File system APIs

E.g. can't append or modify parts of objects

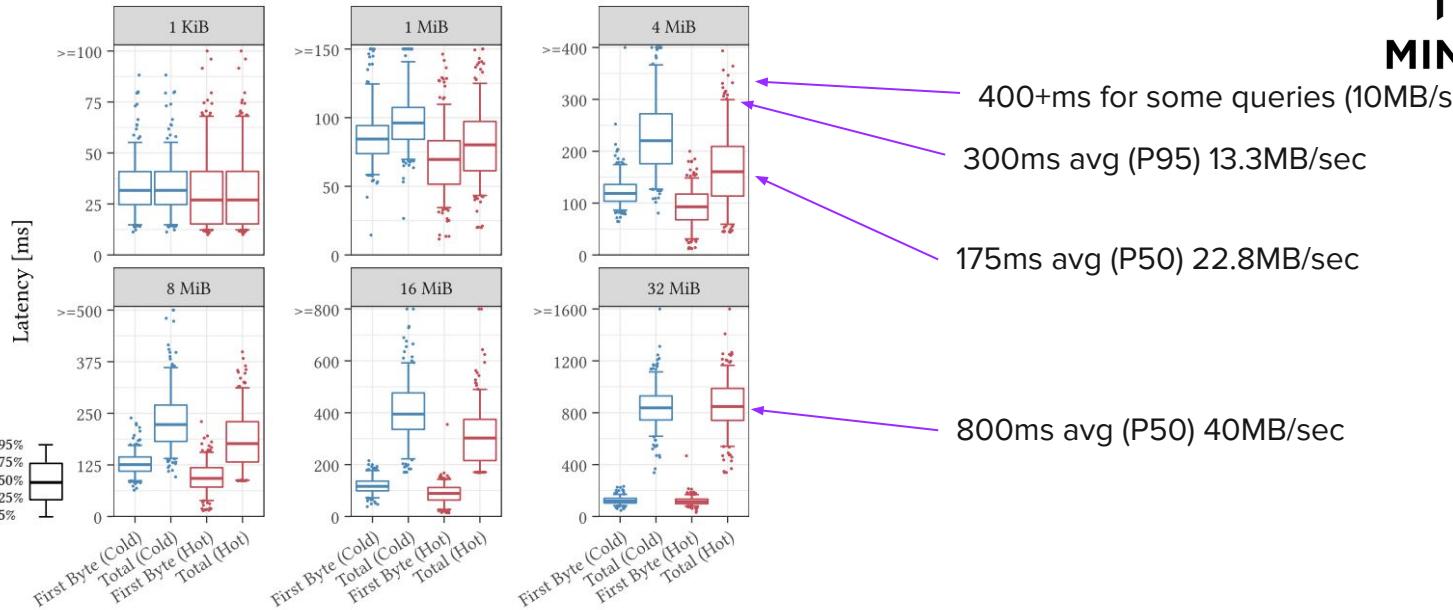


- Durability (3x replication, cross AZ, handled transparently)
- “Infinite” scale + capacity
- cheap (\$23/TB/month\*)
- Pay per access (not per byte): \$0.40/million requests

⇒ Compelling to outsource persistent storage to Object Stores

# Object Stores: Ugly

Significant latency / latency unpredictability



Microsoft Azure  
Blob Storage



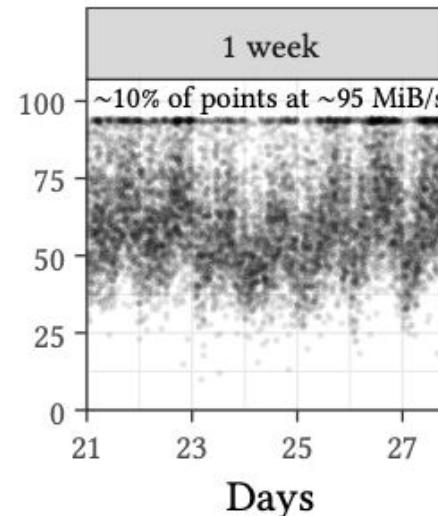
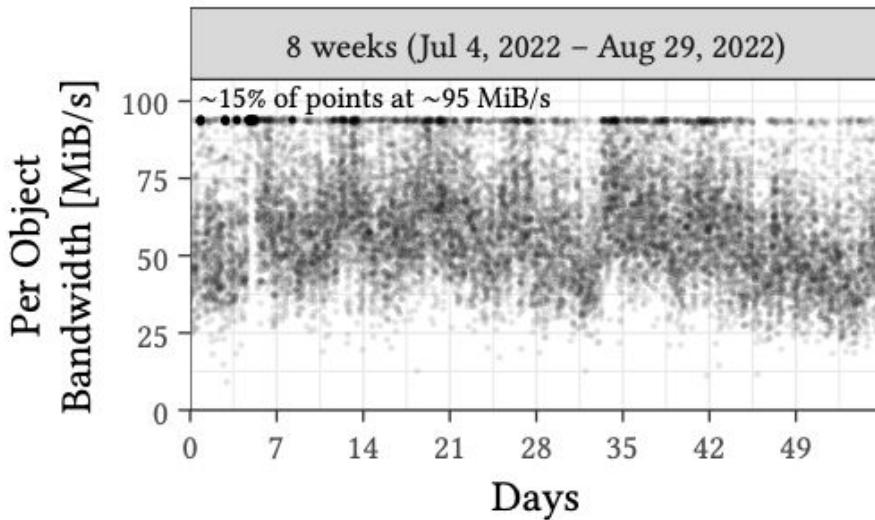
MINIO



Google Cloud  
Storage

# Object Stores: Ugly

Significant latency / latency unpredictability



Microsoft Azure  
Blob Storage



IO



Google Cloud  
Storage

Not isolated from  
other workloads:  
observed bandwidth  
varies cyclically

Source: [Exploiting Cloud Object Storage for High-Performance Analytics](#)

# Object Stores: Ugly

Isn't as cheap as it turns out, more expensive over time

Amazon S3 Historical Prices (2008-2025)



Microsoft Azure  
Blob Storage



Google Cloud  
Storage

🤔 no price change  
in the last 10 years...

# Elastic Compute: What



GCP Compute  
Engine



AzureVM



Amazon EC2

**What:** “Rent VM’s by the day, hour or minute (now)”

Pricing: <https://aws.amazon.com/ec2/pricing/on-demand/>

**Example:** t2.xlarge: 4 vCPU 16GB RAM @ \$0.1856/hour)

0:00:00      **Start** VM (start billing)

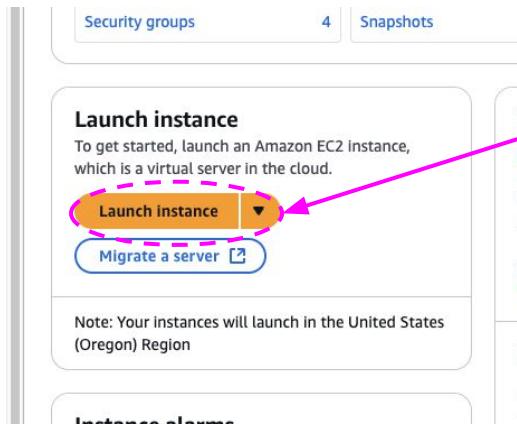
3:25:24      **Stop** VM (stop billing)

$$3 \cdot 60 \cdot 60 + 25 \cdot 60 + 24 = 12324 \text{ seconds}$$

$$12324 \text{ seconds} / 3600 \text{ seconds/hour} * \$0.1856/\text{hour} \Rightarrow \$0.64$$

# Elastic Compute: Why

No upfront capital investment  
⇒ Much more efficiently use hardware



The screenshot shows the AWS Lambda service dashboard. At the top, there are tabs for 'Security groups' (highlighted in blue), '4' (likely indicating the number of functions), and 'Schemas'. Below these, there's a search bar and a 'Create function' button. The main area is titled 'Lambda functions' and lists three functions: 'lambda-tutorial', 'lambda-tutorial-1', and 'lambda-tutorial-2'. Each function entry includes a 'Edit' button, a 'Logs' link, and a 'Test' button. A pink arrow points from the text 'VS' to the 'Launch instance' button in the screenshot below.

VS



Personal Anecdote: budgeting ~ \$250K 6 months in advance for server clusters to test on

HP DL380s (popular midrange server in late 2000s)

[Ebay “Lot of 21 HP Proliant DL380”](#)

# Elastic Compute: The Ugly

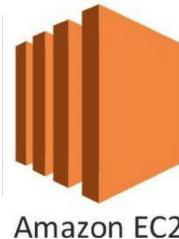
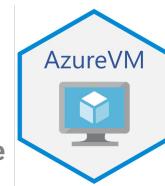
**Overabundance leads to waste:** 💰

Easy to spend

Developers often leave machines running by accident 😷

Rise of cost optimization software

Personal Anecdote: \$1m/month AWS bills



# Elastic Compute: The Ugly

**Kinda crappy compared to your own machine**

Sequential Write Throughput:

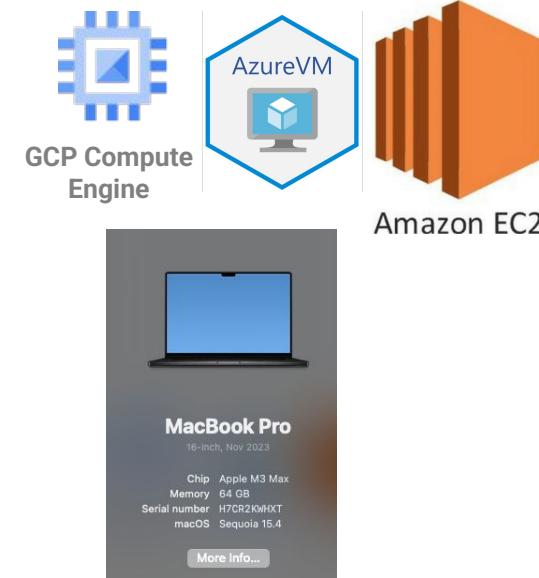
- My Macbook Pro 1TB SSD: > **1GB/s**
- GCP VM\* 4 SSD @ 1.5TB RAID 0: **815 MB/s**

\* c3-standard-22-lssd (22 vCPUs, 88 GB Memory)

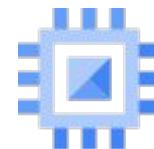
4 “local”SSDs, RAID0, Intel Sapphire Rapids x86\_64

```
# Test the IO throughput using `dd`  
dd if=/dev/zero of=/data/test1.img bs=1G count=10 oflag=dsync  
# 10737418240 bytes (11 GB, 10 GiB) copied, 13.179 s, 815 MB/s
```

Read more: [tpchgen-rs World's fastest open source TPC-H data generator, written in Rust - Apache DataFusion Blog](#)



VS

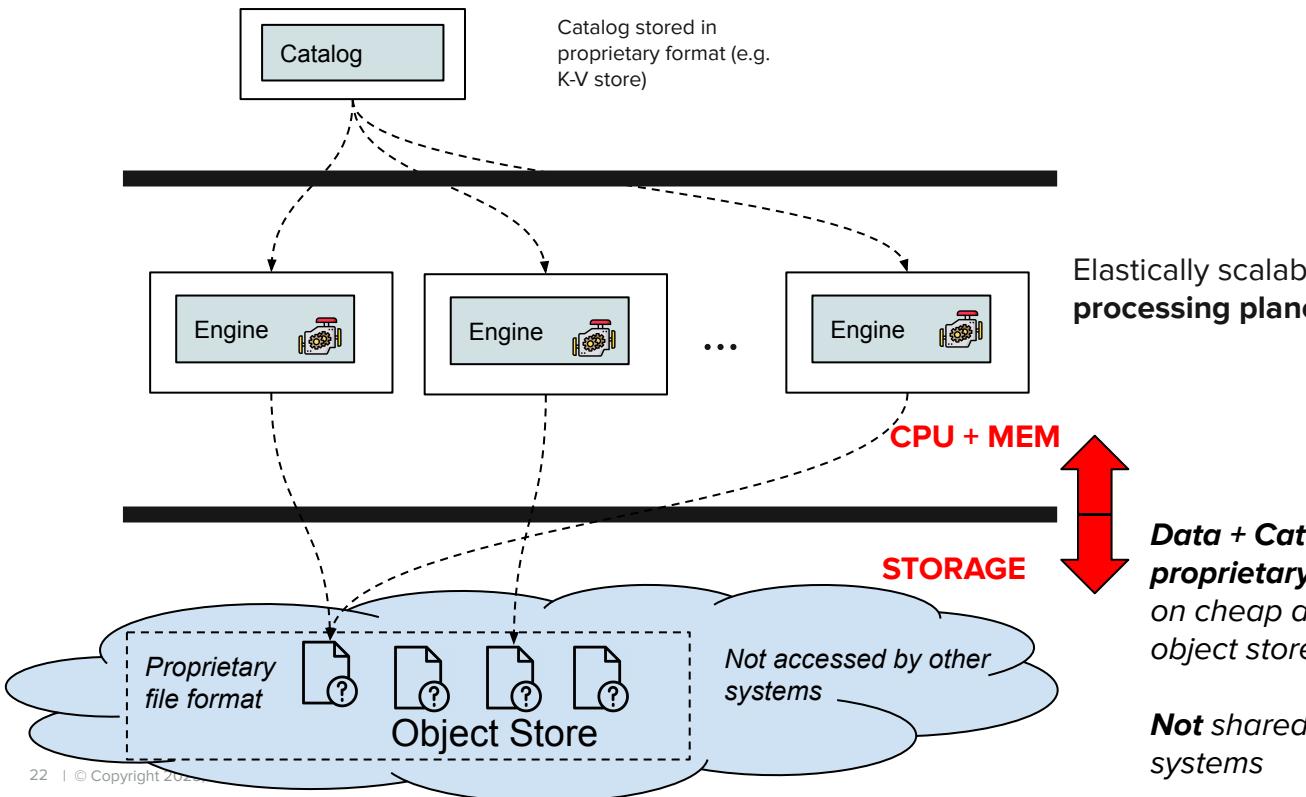


GCP Compute Engine

# Cloud Database Architecture “Disaggregated Storage Design”

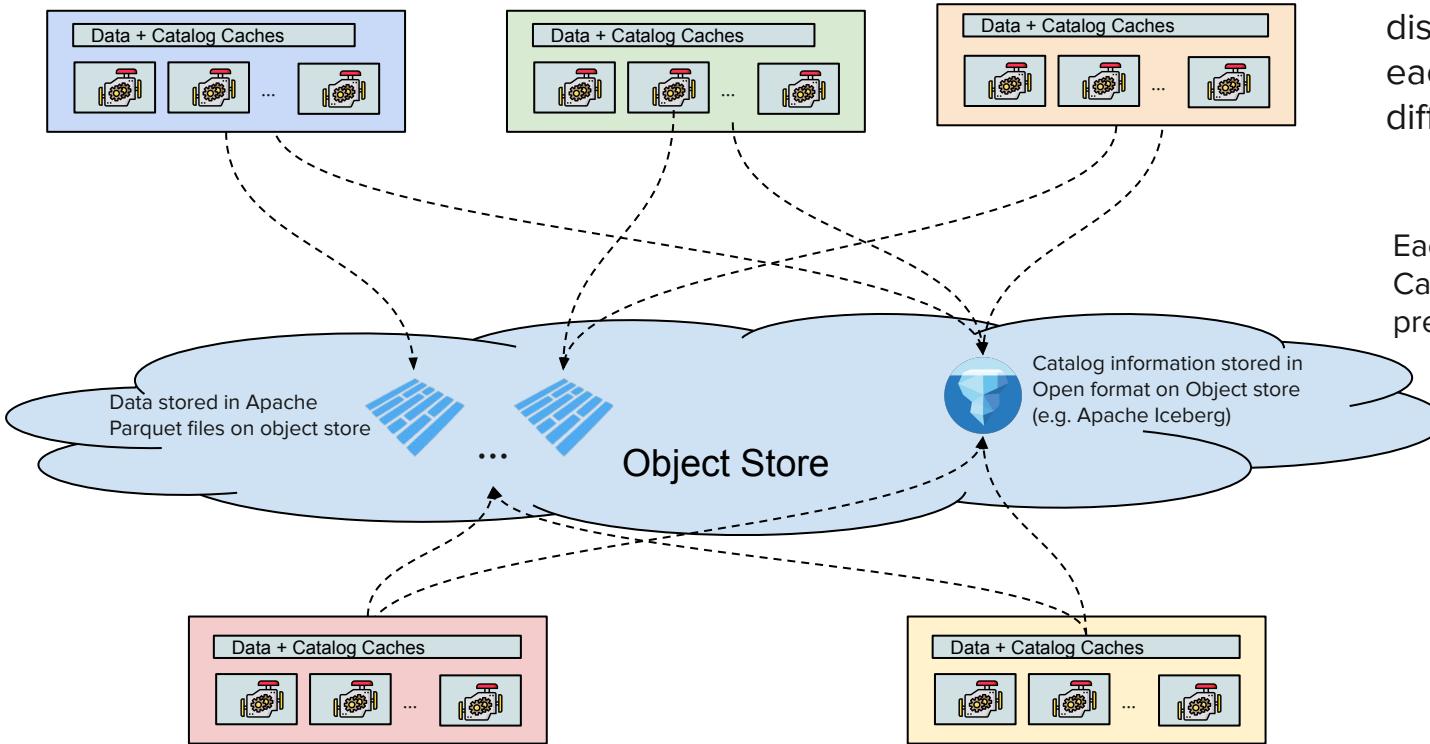
# Disaggregated Architectures (2010s)

## “Cloud Data Warehouse”



- Driven by**
- Hourly VM rental
  - 10x cheaper storage (AWS S3)

# Cloud Data Lakes (2020s)



**Constellation** of disaggregated systems, each focused on different use case

Each includes use case specific Catalog + Data Caching / precomputation / engines

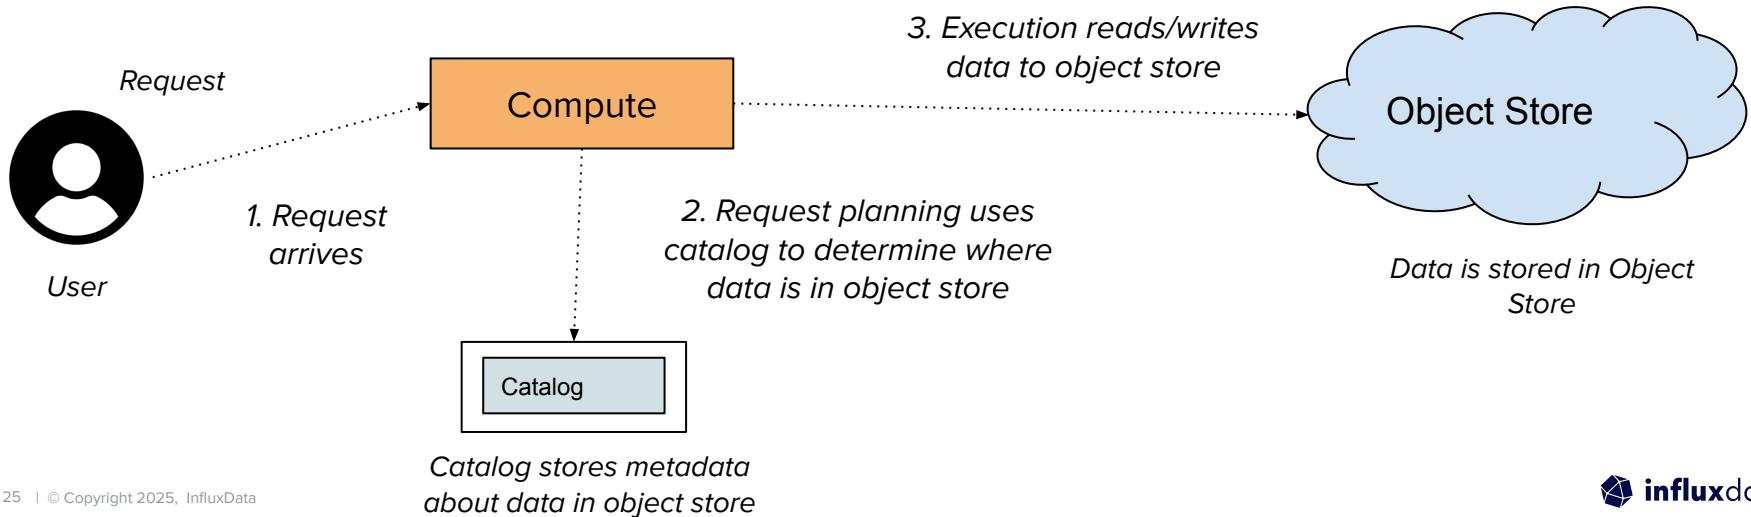
# Common Features of Disaggregated Databases

# Metadata Store (“Catalogs”)

Object storage latency (100s of ms) is too high for planning for many workloads (both read and write)

No multi-object transactions

⇒ metadata ‘catalog’ describes data layout in object store



# Metadata Store (“Catalogs”)

Popular choices:

- Key Value store (FoundationDB)
- Traditional transactional SQL systems (postgres)



Typical Contents

- **Schema:** tables, columns, types, etc.
- **Partitioning:** partitions, partition values, etc.
- **File Locations:** paths on object store
- **Pruning:** per-column min/maxes (Small Materialized Aggregates / Zone Maps), Bloom Filters, etc.

Reference: [How FoundationDB Powers Snowflake Metadata Forward](#)

# Separate Scalable Operations

Separate major responsibilities into separately scalable sets of VMs

## Typical Components:

- Write / Ingestion
- Query
- Reorganization (compaction, garbage collection, etc)

**Why:** scale capacity along with demand (e.g scale writers up to handle bursts)

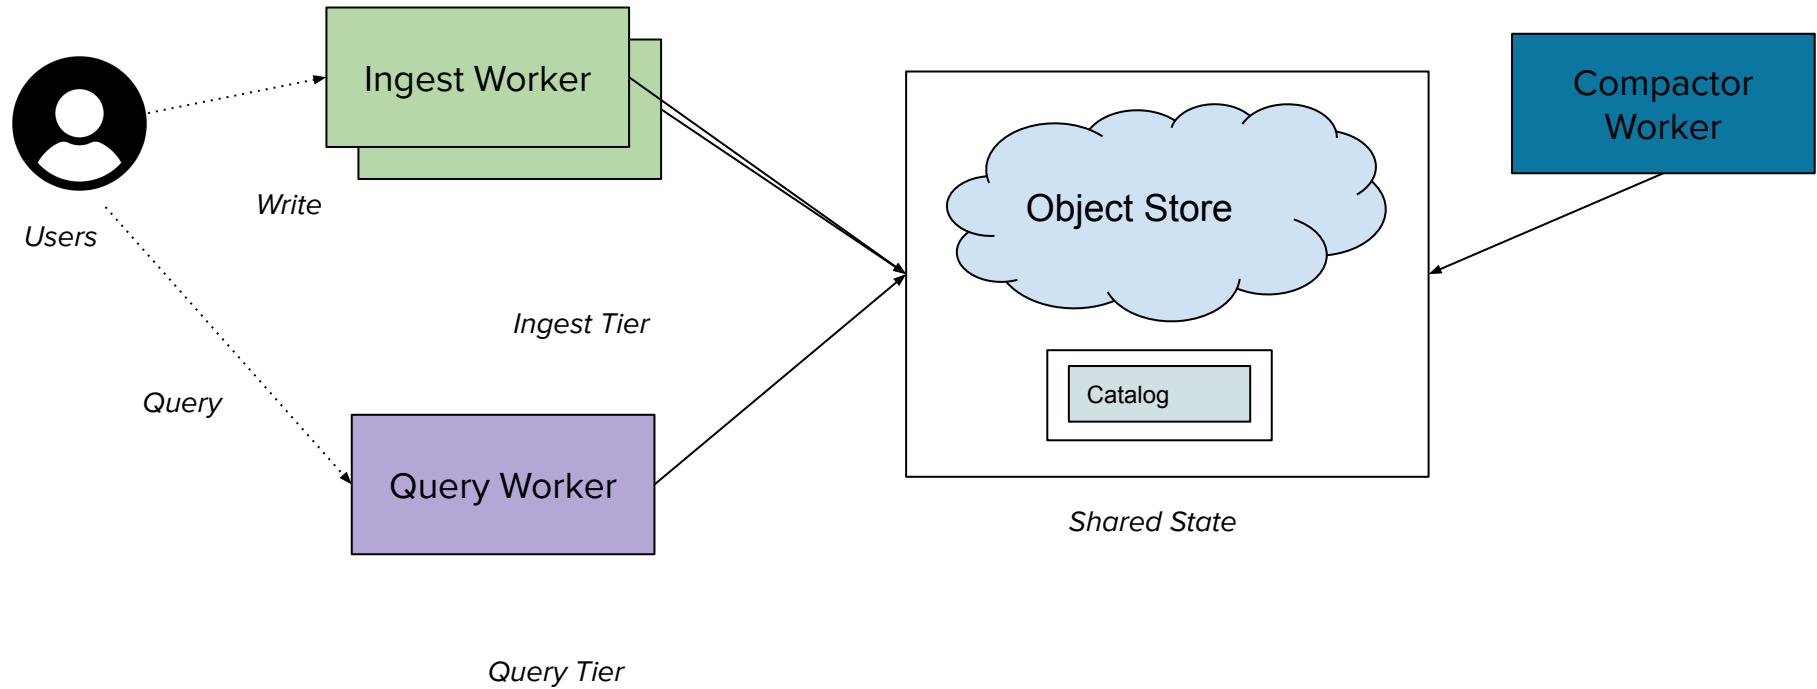
## Industrial Examples:

<https://docs.snowflake.com/en/user-guide/intro-key-concepts> (virtual warehouses)

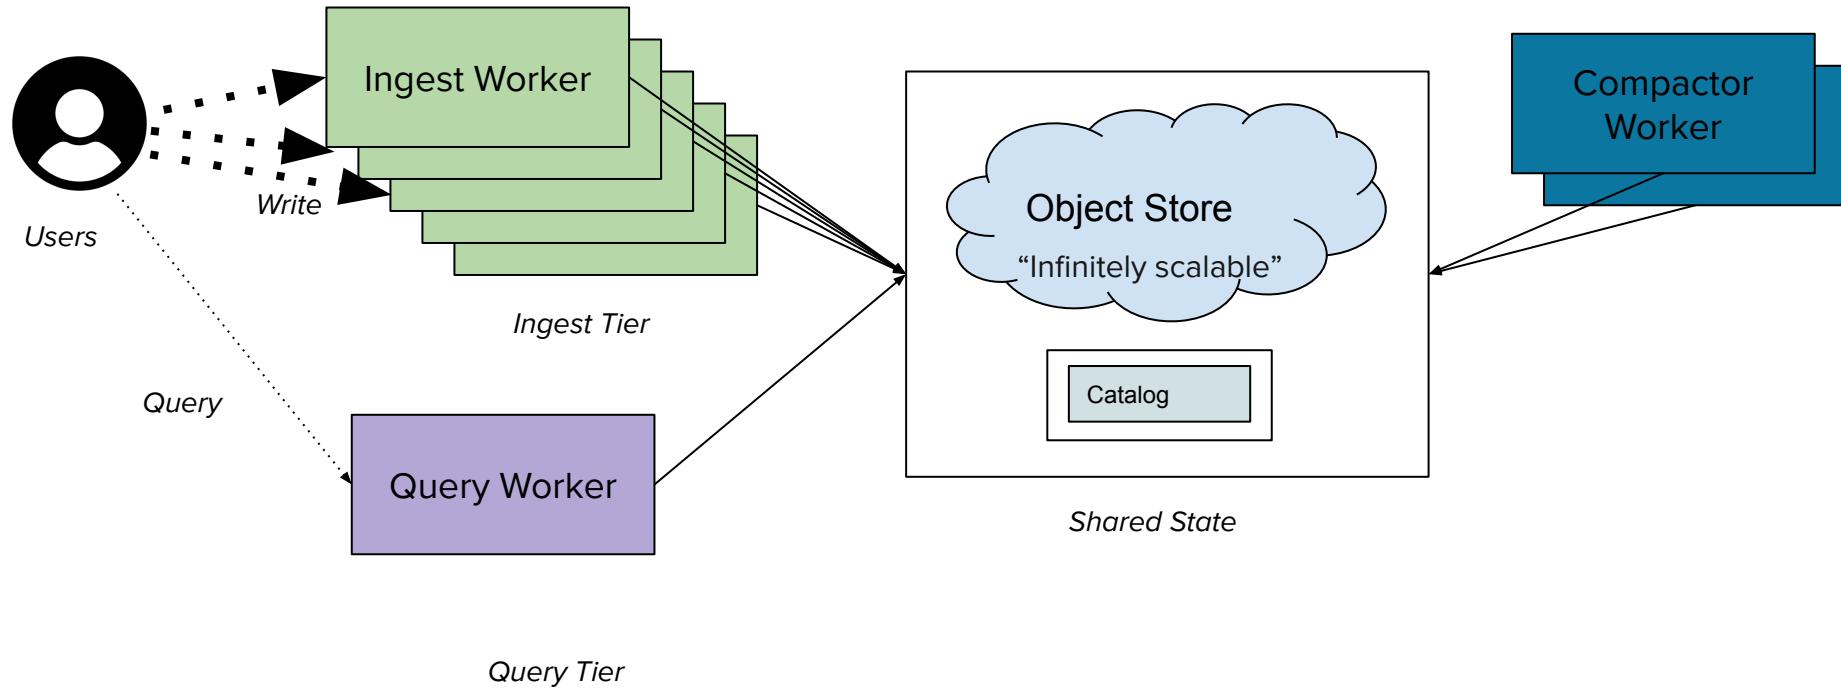
<https://www.datadoghq.com/blog/engineering/introducing-husky/>

<https://www.influxdata.com/blog/influxdb-3-0-system-architecture/>

# Separate Scalable Operations: Example



# Separate Scalable Operations: Example



Increase write workload leads to more ingest and compactor workers, no need to increase query tier

# Write Buffering

Large per-request overhead to object store (\$\$ and latency)

⇒ Buffer in ram + locally to amortize cost across many requests

1. Writes are collected in memory buffer

Write

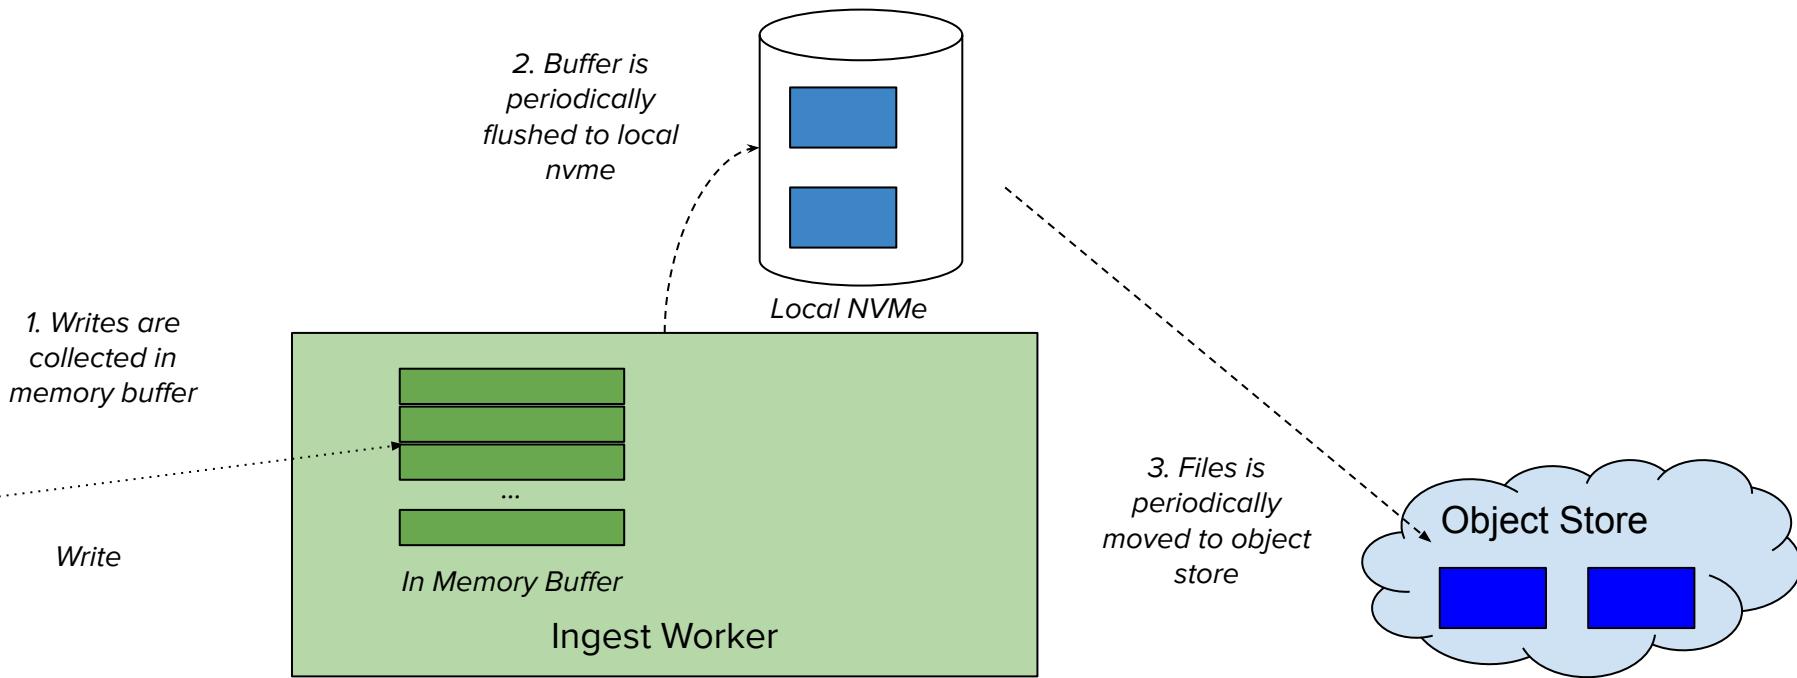
In Memory Buffer

Ingest Worker

2. Buffer is periodically flushed to object store

Object Store

# Write Buffering / Local Storage



# Write Buffering

## Challenges:

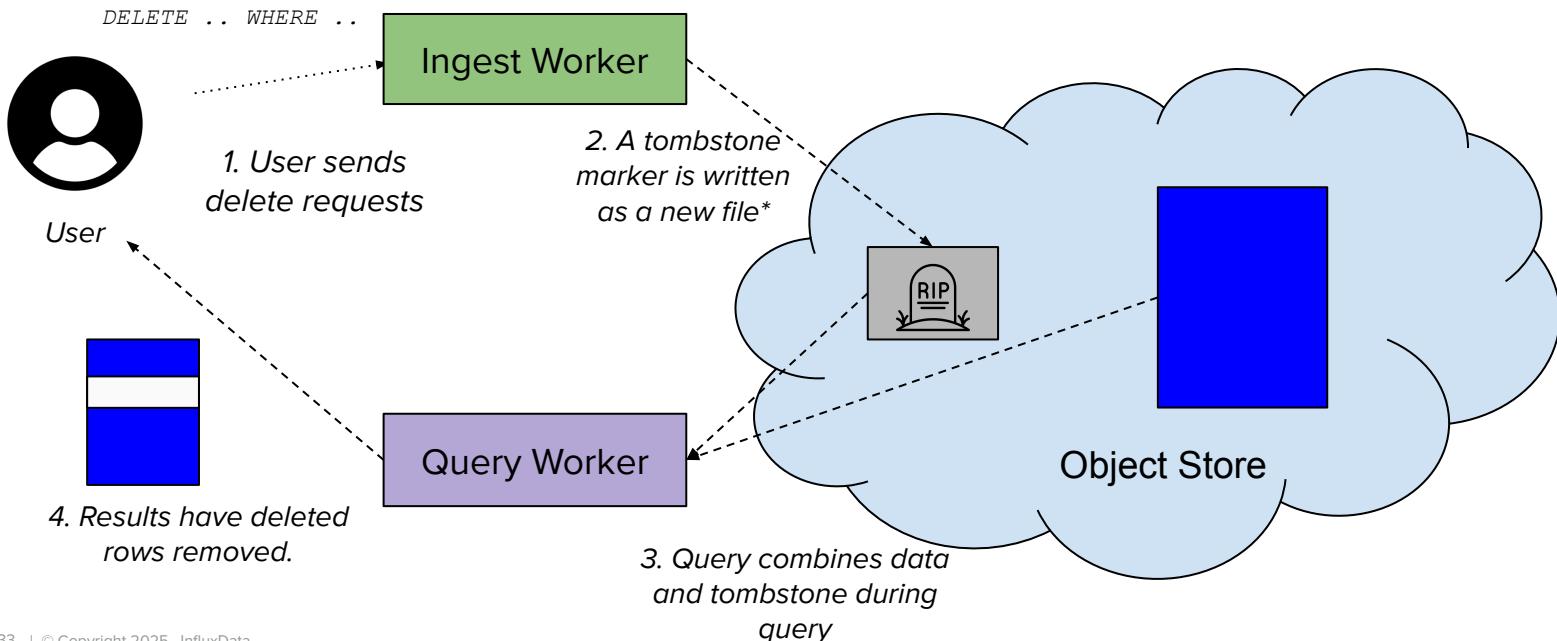
- Durability of data before it is written to object store
- Time to become readable (is memory in buffer readable?)

## Examples:

- [Monarch: Google's Planet-Scale In-Memory Time Series Database](#)
- [Architecture | WarpStream](#)
- [Architecture | SlateDB](#)

# Deletes (+ Updates) via Tombstones

Write once (no updates) storage ⇒ Delete / Update writes new things



# Deletes (+ Updates) via Tombstones

## Variations:

- **Delete Vectors** (row ids / offsets deleted)
  - Often stored in objects
  - Time consuming to create / Faster query execution
- **Stored Predicates**
  - Often stored in meta store
  - Fast to create / Potentially slower query execution
  - Tied to predicate expressions
  - Slower as number of deletes increase

Offset
21
31
67
104

*Delete Vector with resolved row ids*

## Challenges

- Sequencing deletes with inserts
- Performance
- Eventually reclaiming Storage

`user_id IN (123, 456)`

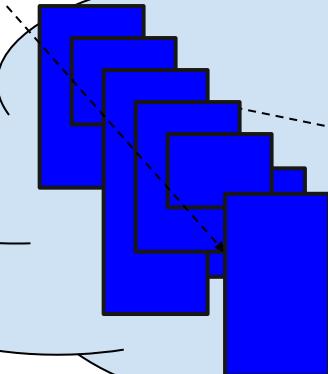
*Delete Predicate*

# Data Layout optimization

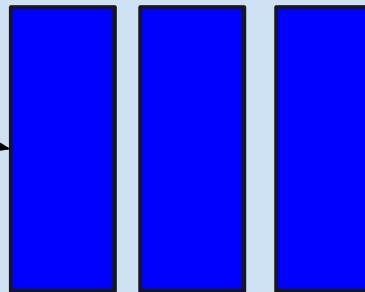
Object Storage is write once: write new objects, but not modify existing  
⇒ Rewrite objects overtime (better organized, garbage collect, etc)

1. New data arrives in  
new object

2. objects are compacted,  
garbage collected



3. Old files removed



Object Store

# “Table Formats”

Metadata catalogs / stores are proprietary, add operational overhead.

Use object store to store metadata (cost of increased planning latency)

⇒ Standardize describing what files make up a table

Examples:

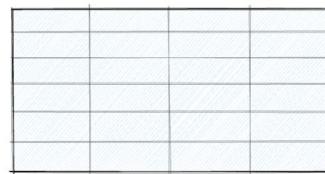


# “Table Formats”

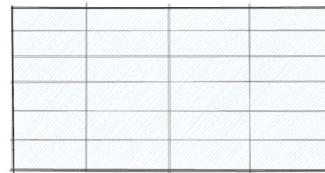
Separate Data from Metadata



DATA



METADATA



# “Table Formats”

Classic case of solve the problem with a layer of indirection

1. Current list of files stored in one object, used by current queries

001  
002

3. List file is updated atomically, updating the pointers.



4. Objects not used in future queries

2. New file is written

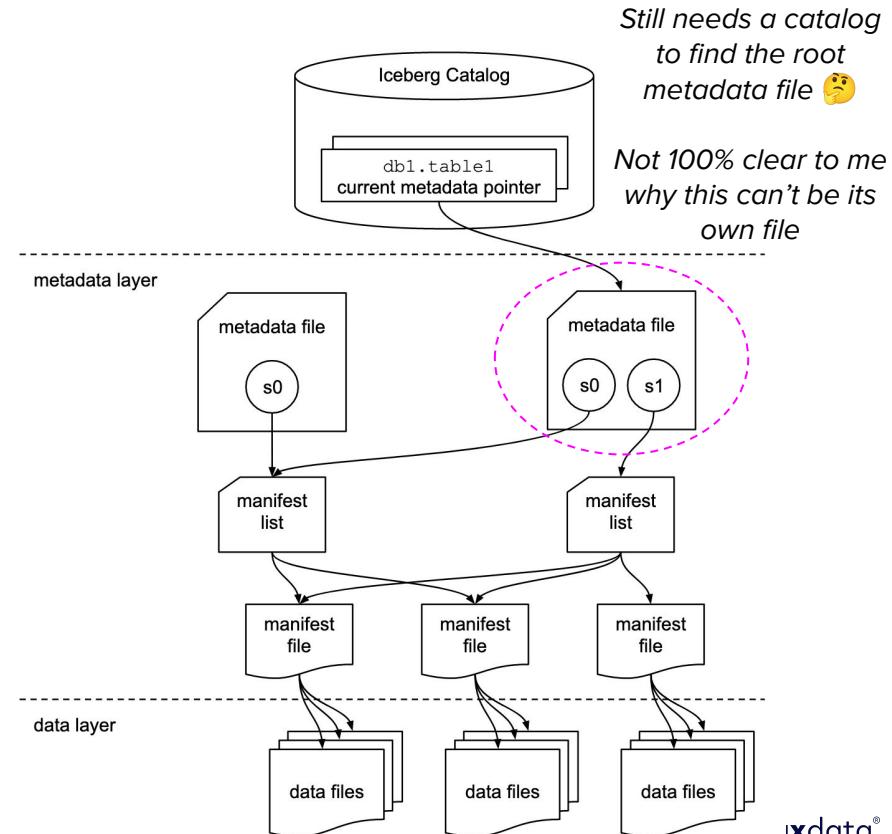
# Table Formats

Apache

**ICEBERG**



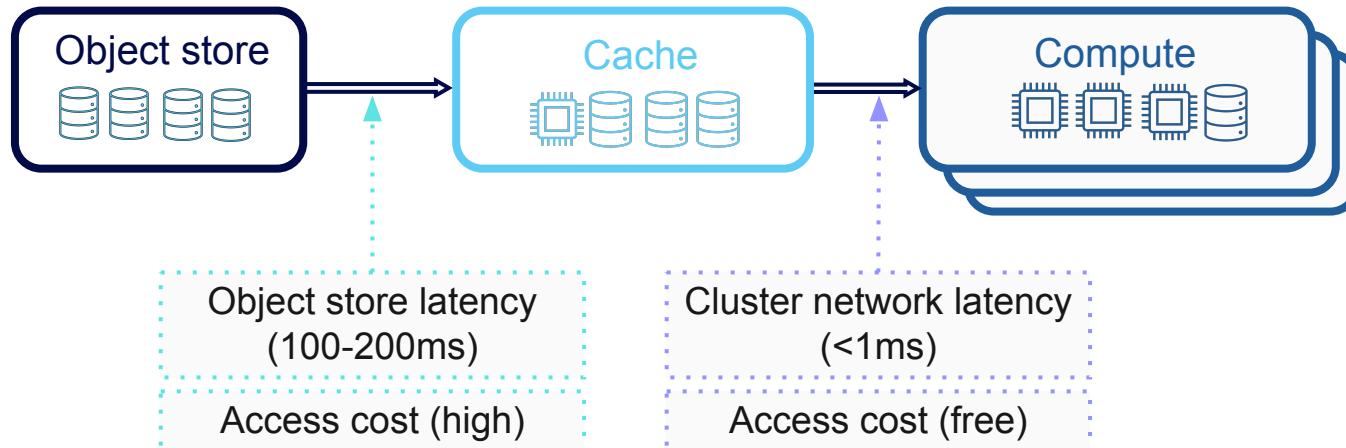
Adds extra layer(s) of indirection



# Object cache

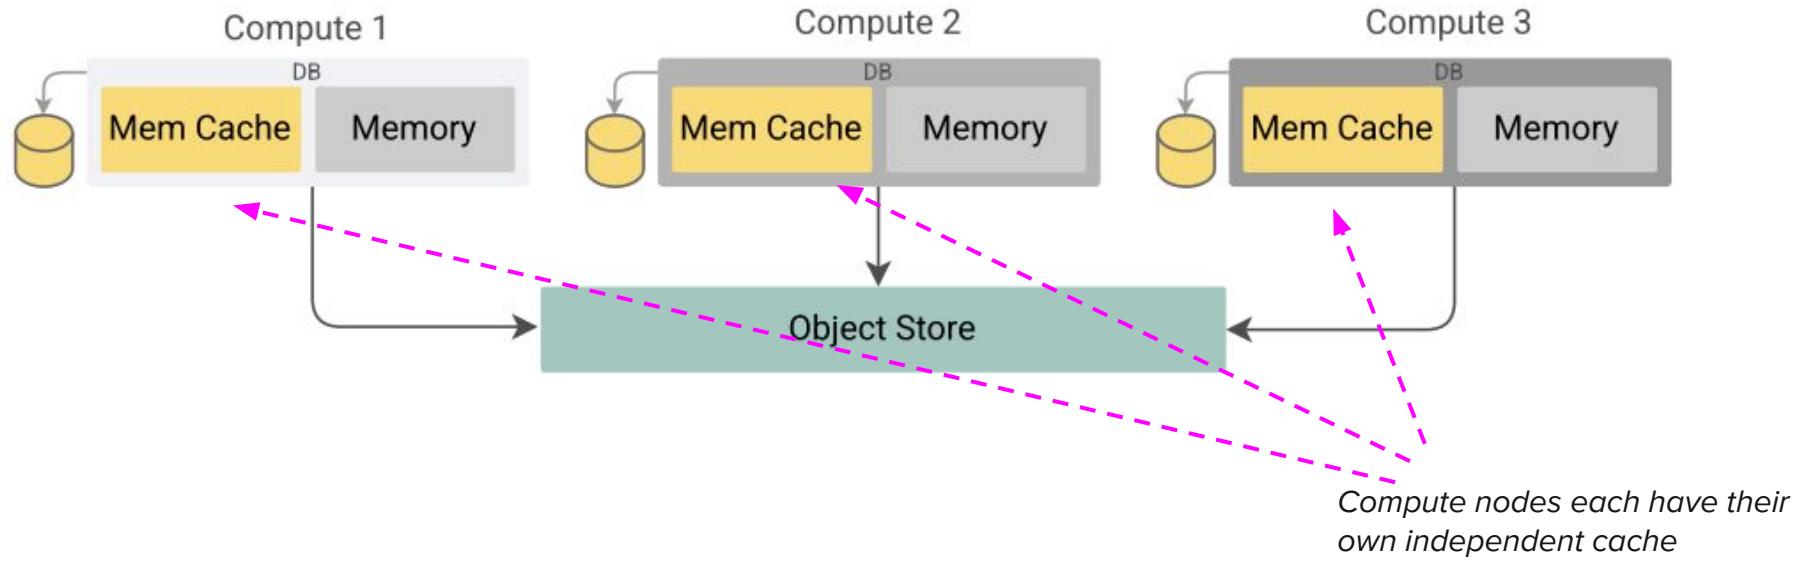
Object storage latency, unpredictability, and cost per access

⇒ Reduce via in cluster caches



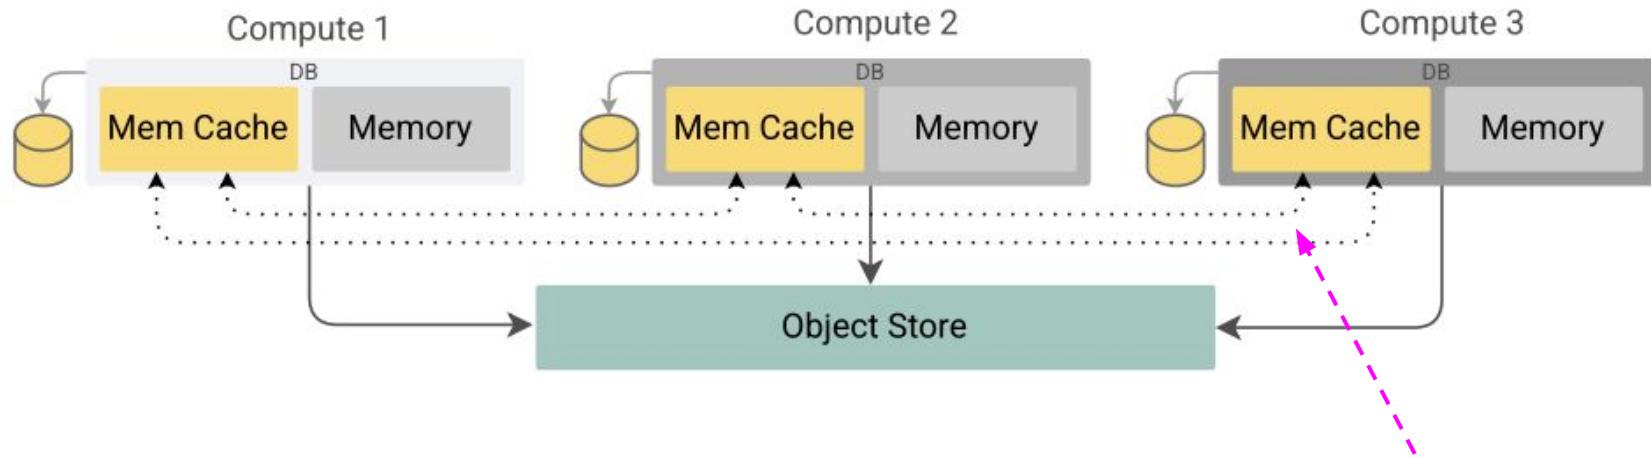
Credit: [Xiangpeng Hao](#), UW Madison

# Object cache: Common Topologies



CIDR 2025: [The Five-Minute Rule for the Cloud: Caching in Analytics Systems](#)

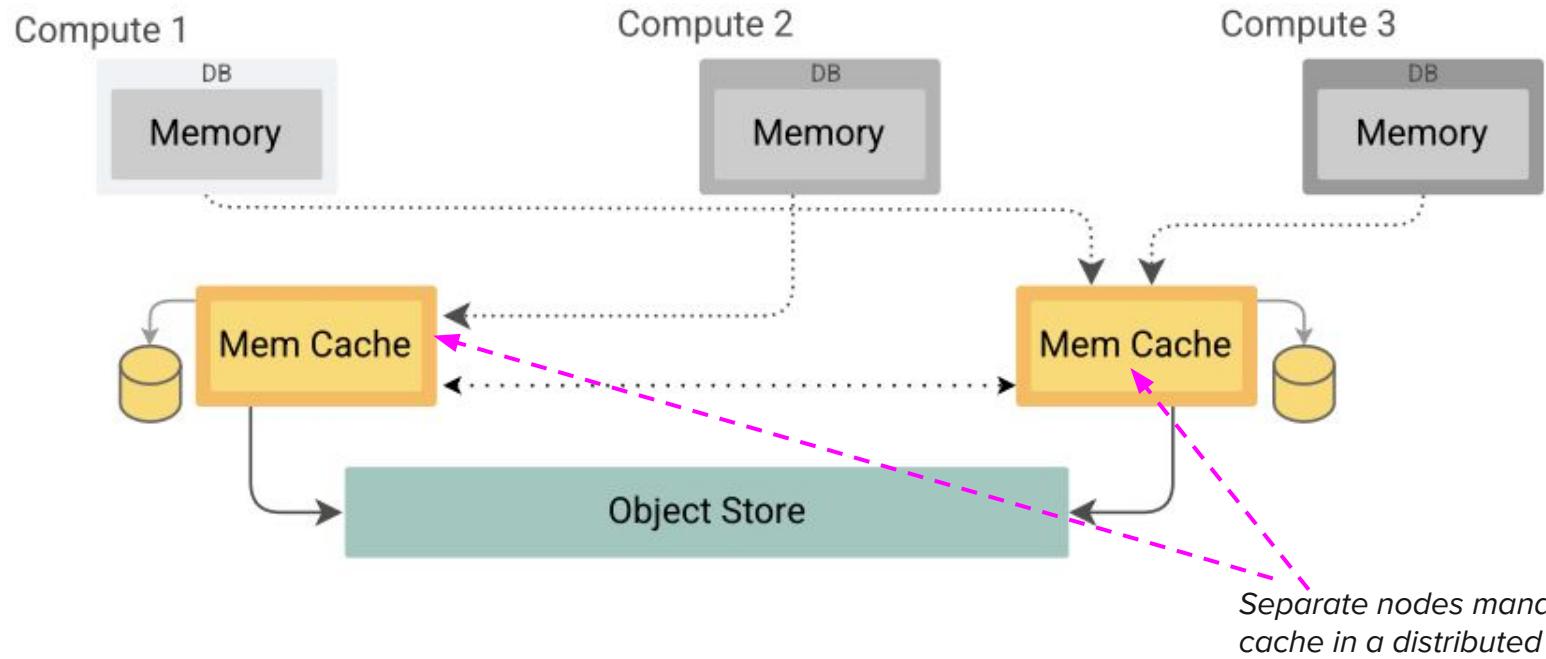
# Object cache: Common Topologies



*Compute nodes cooperatively manage a distributed cache*

CIDR 2025: [The Five-Minute Rule for the Cloud: Caching in Analytics Systems](#)

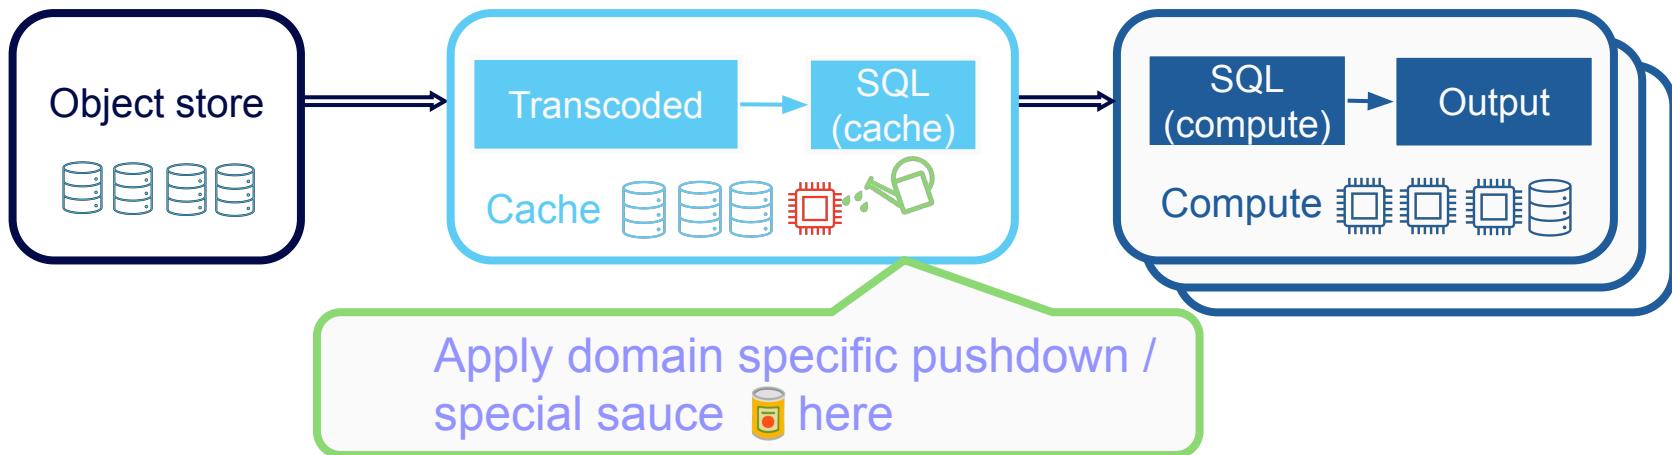
# Object cache: Common Topologies



CIDR 2025: [The Five-Minute Rule for the Cloud: Caching in Analytics Systems](#)

# Object cache++

If you have a cache anyways, opportunities for transcoding

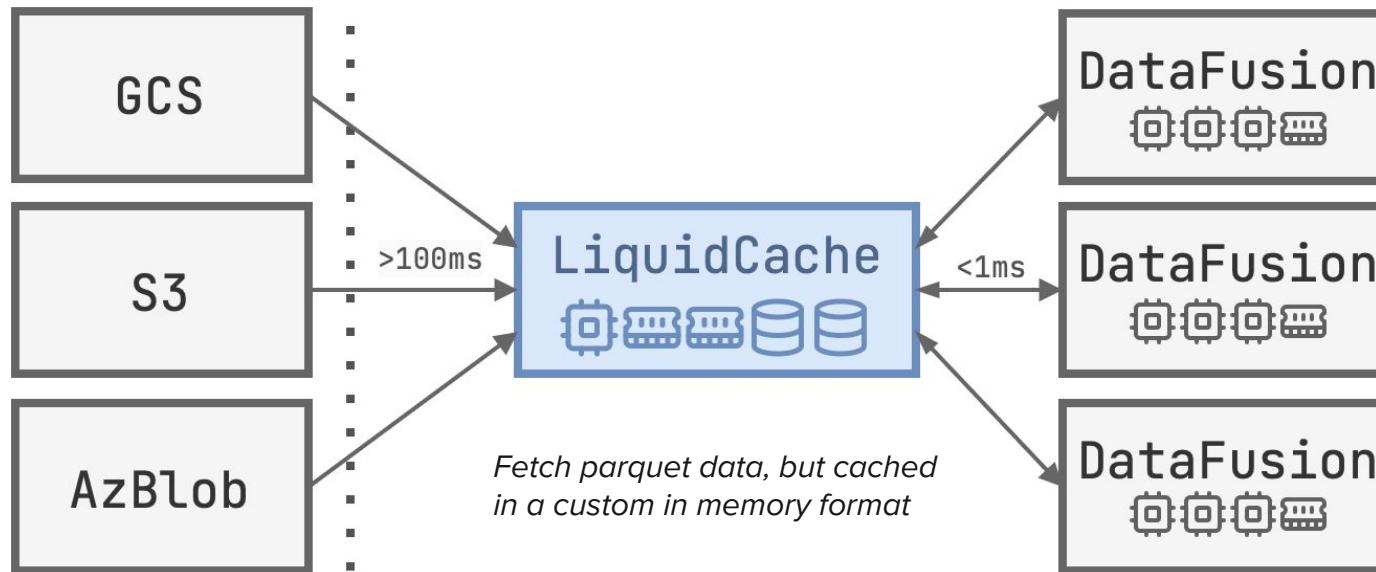


Credit: [Xiangpeng Hao](#), UW Madison

# Object cache++

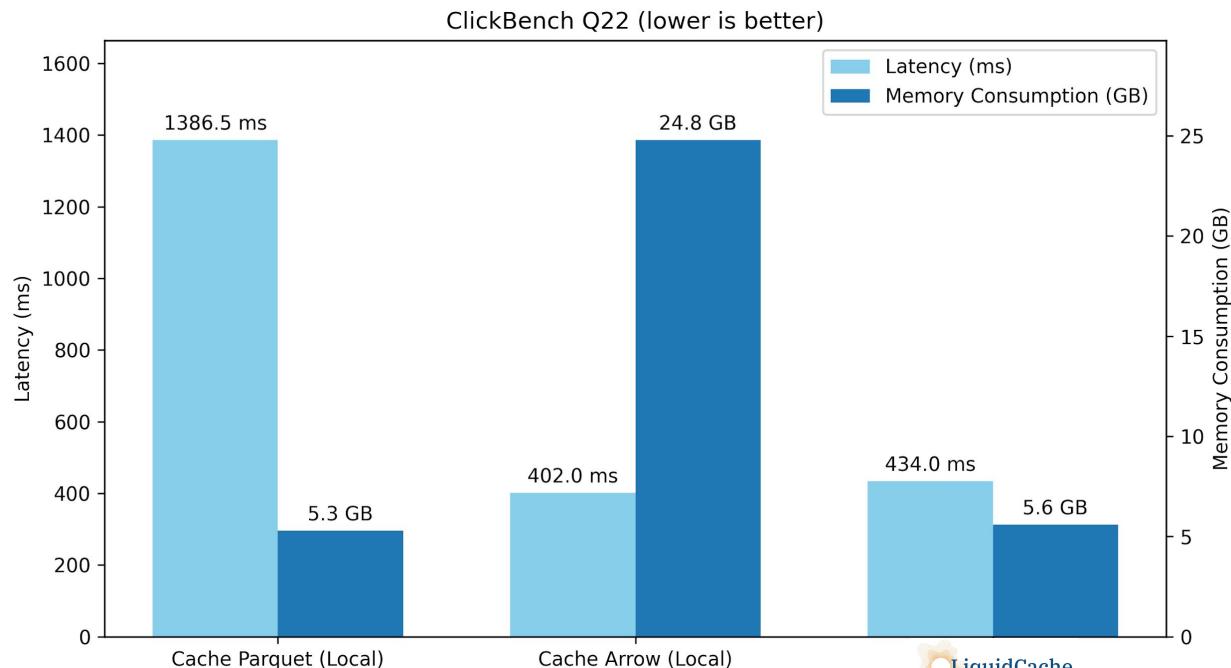


Example: <https://github.com/XiangpengHao/liquid-cache>



*Disclosure: I am an advisor to this project*

# Liquid Cache – ClickBench Q22

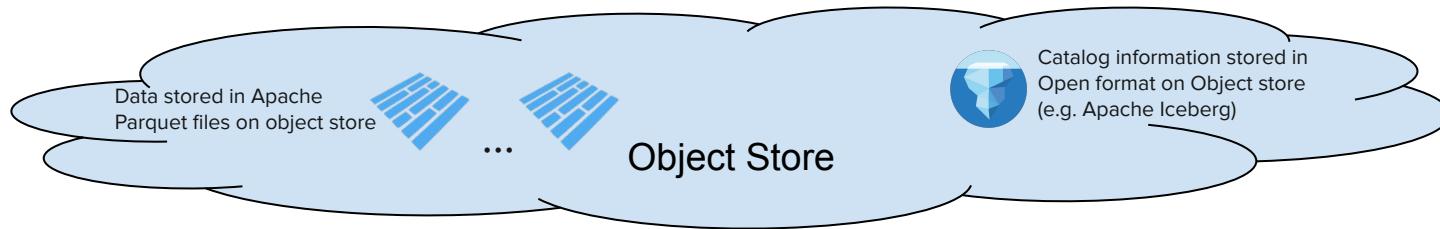


```
SELECT
  "SearchPhrase",
  MIN("URL"),
  MIN("Title"),
  COUNT(*) AS c,
  COUNT(DISTINCT "UserID")
FROM hits
WHERE
  "Title" LIKE '%Google%'
  AND
  "URL" NOT LIKE '%.google.%' AND
  "SearchPhrase" ◇ ''
GROUP BY "SearchPhrase"
ORDER BY c DESC
LIMIT 10;
```

# Future Directions / Predictions

# Increased adoption and interest in Open Formats

Specifically: Apache Parquet and Apache Iceberg



**Implication:** the classic business model of being the data platform that has huge data gravity (hard to move) may be changing

⇒ (Startup) Opportunities for many new specialized engines, etc.

# Further disaggregation

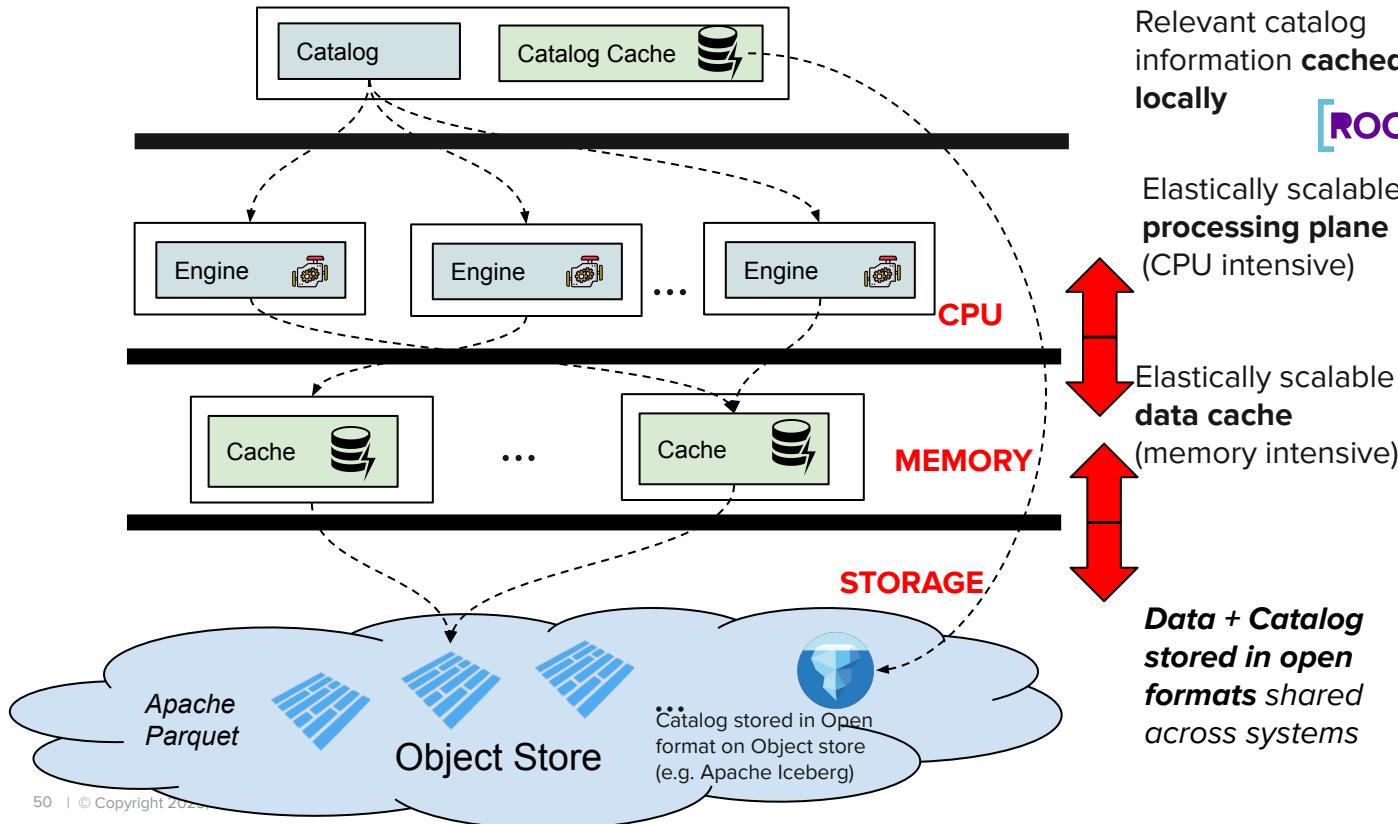
Currently have disaggregated storage:

- Storage
- Memory + compute

Predict further disaggregation of memory from compute:

- Storage
- Memory (cache)
- Compute

# Disaggregated Memory/Cache (2030s) Conjecture



**Early signs**

Driven by

- Serverless (stateless) processing,
- High bandwidth interconnect
- Demand for interoperability



Thank you!  
Questions?